

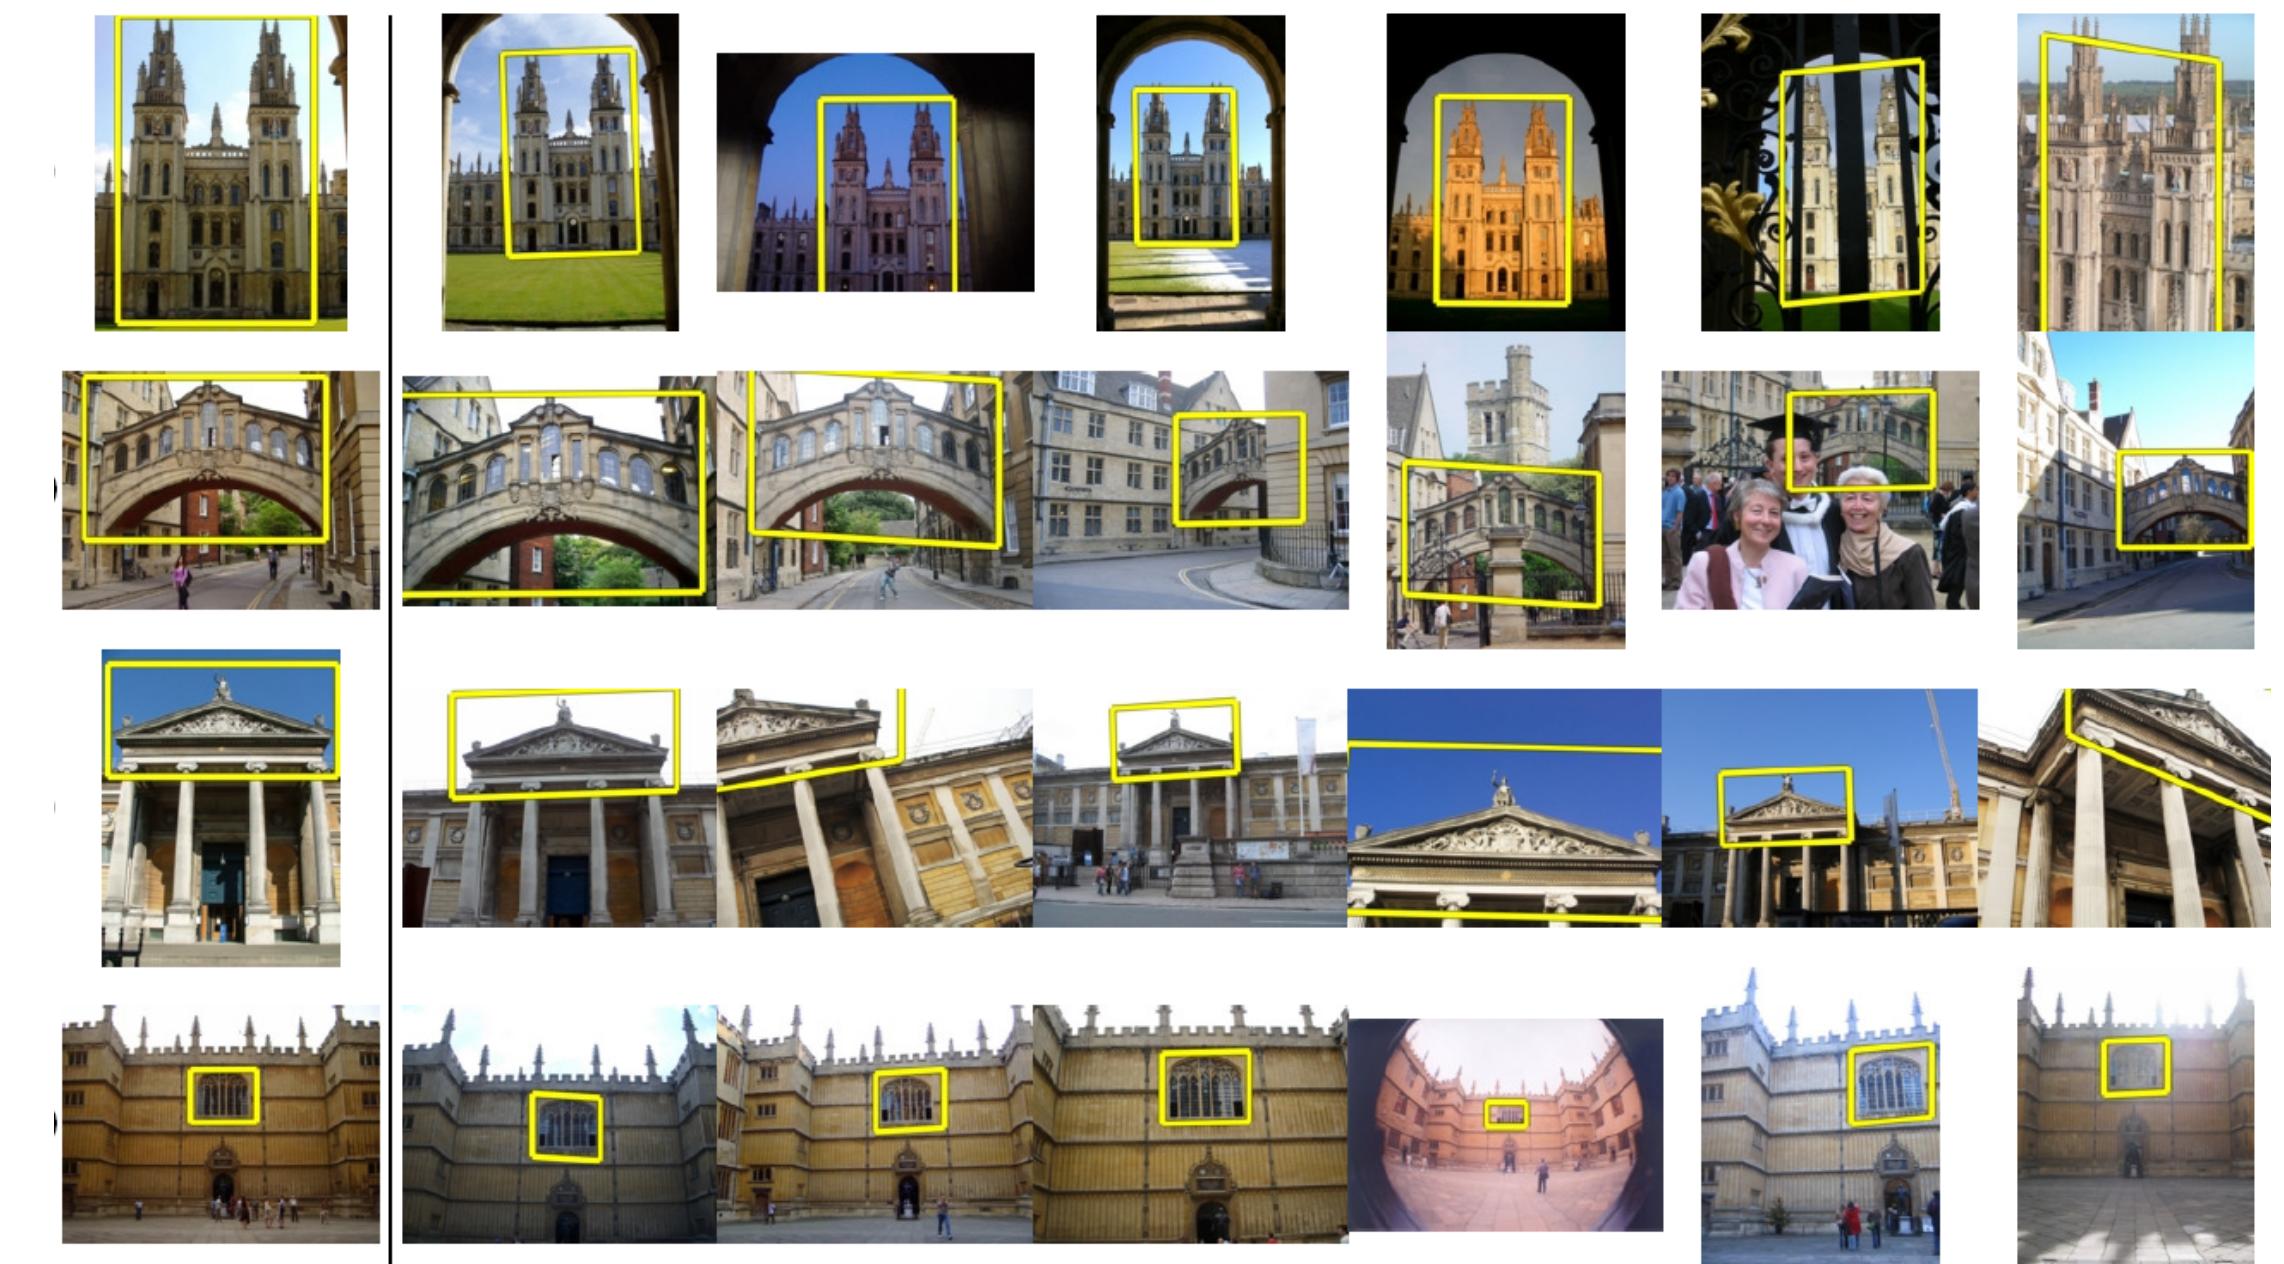


# Transformations and Camera Geometry

CS 650: Computer Vision

# Spatial Transformations

- Many things in computer vision involve analyzing things in different spatial configurations:
  - Different 2D positions (translation)
  - Different 2D orientations (rotation)
  - Different 2D sizes (scaling)
  - Other 2D deformations (e.g., affine)
  - Different 3-D to 2-D projections
  - 3-D to 3-D transformations



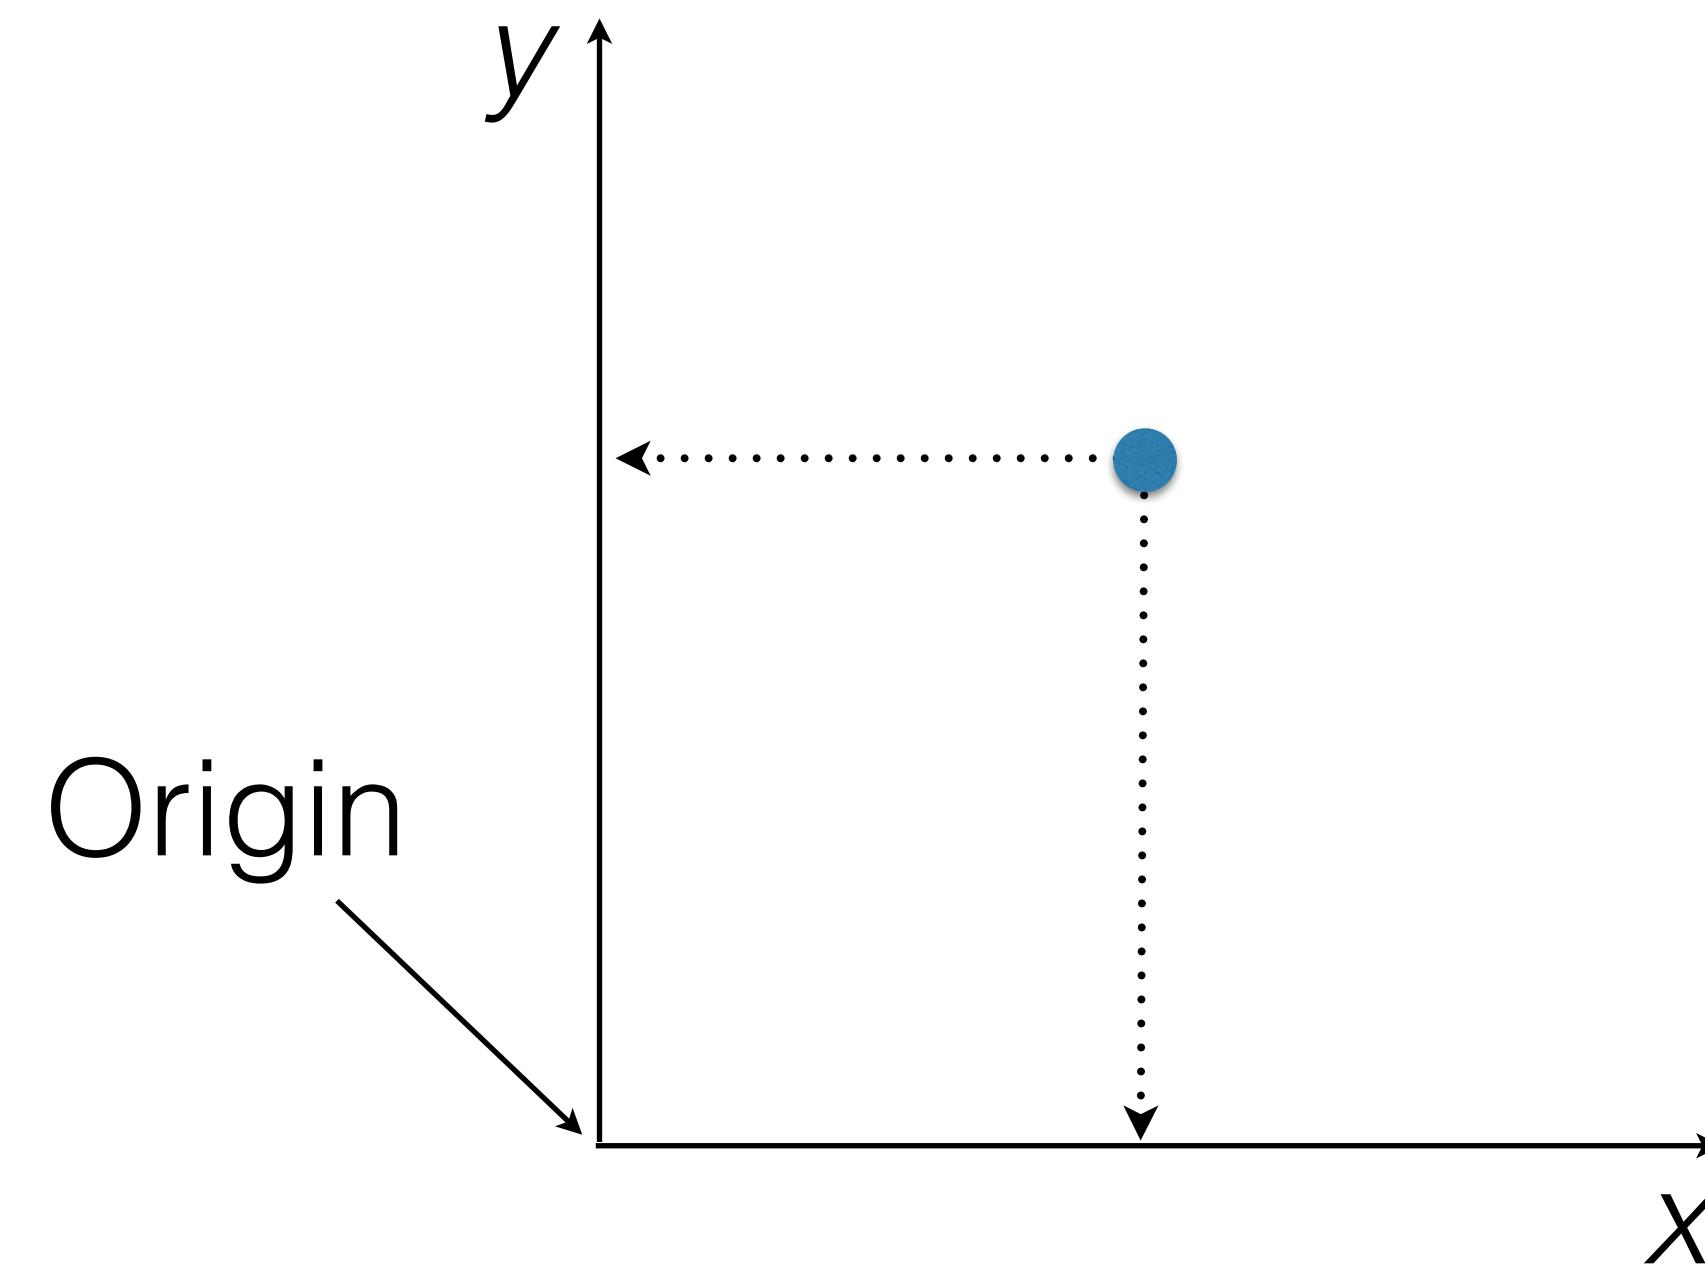
**Quick detour:**  
**Describing spatial things numerically...**

# Describing Points



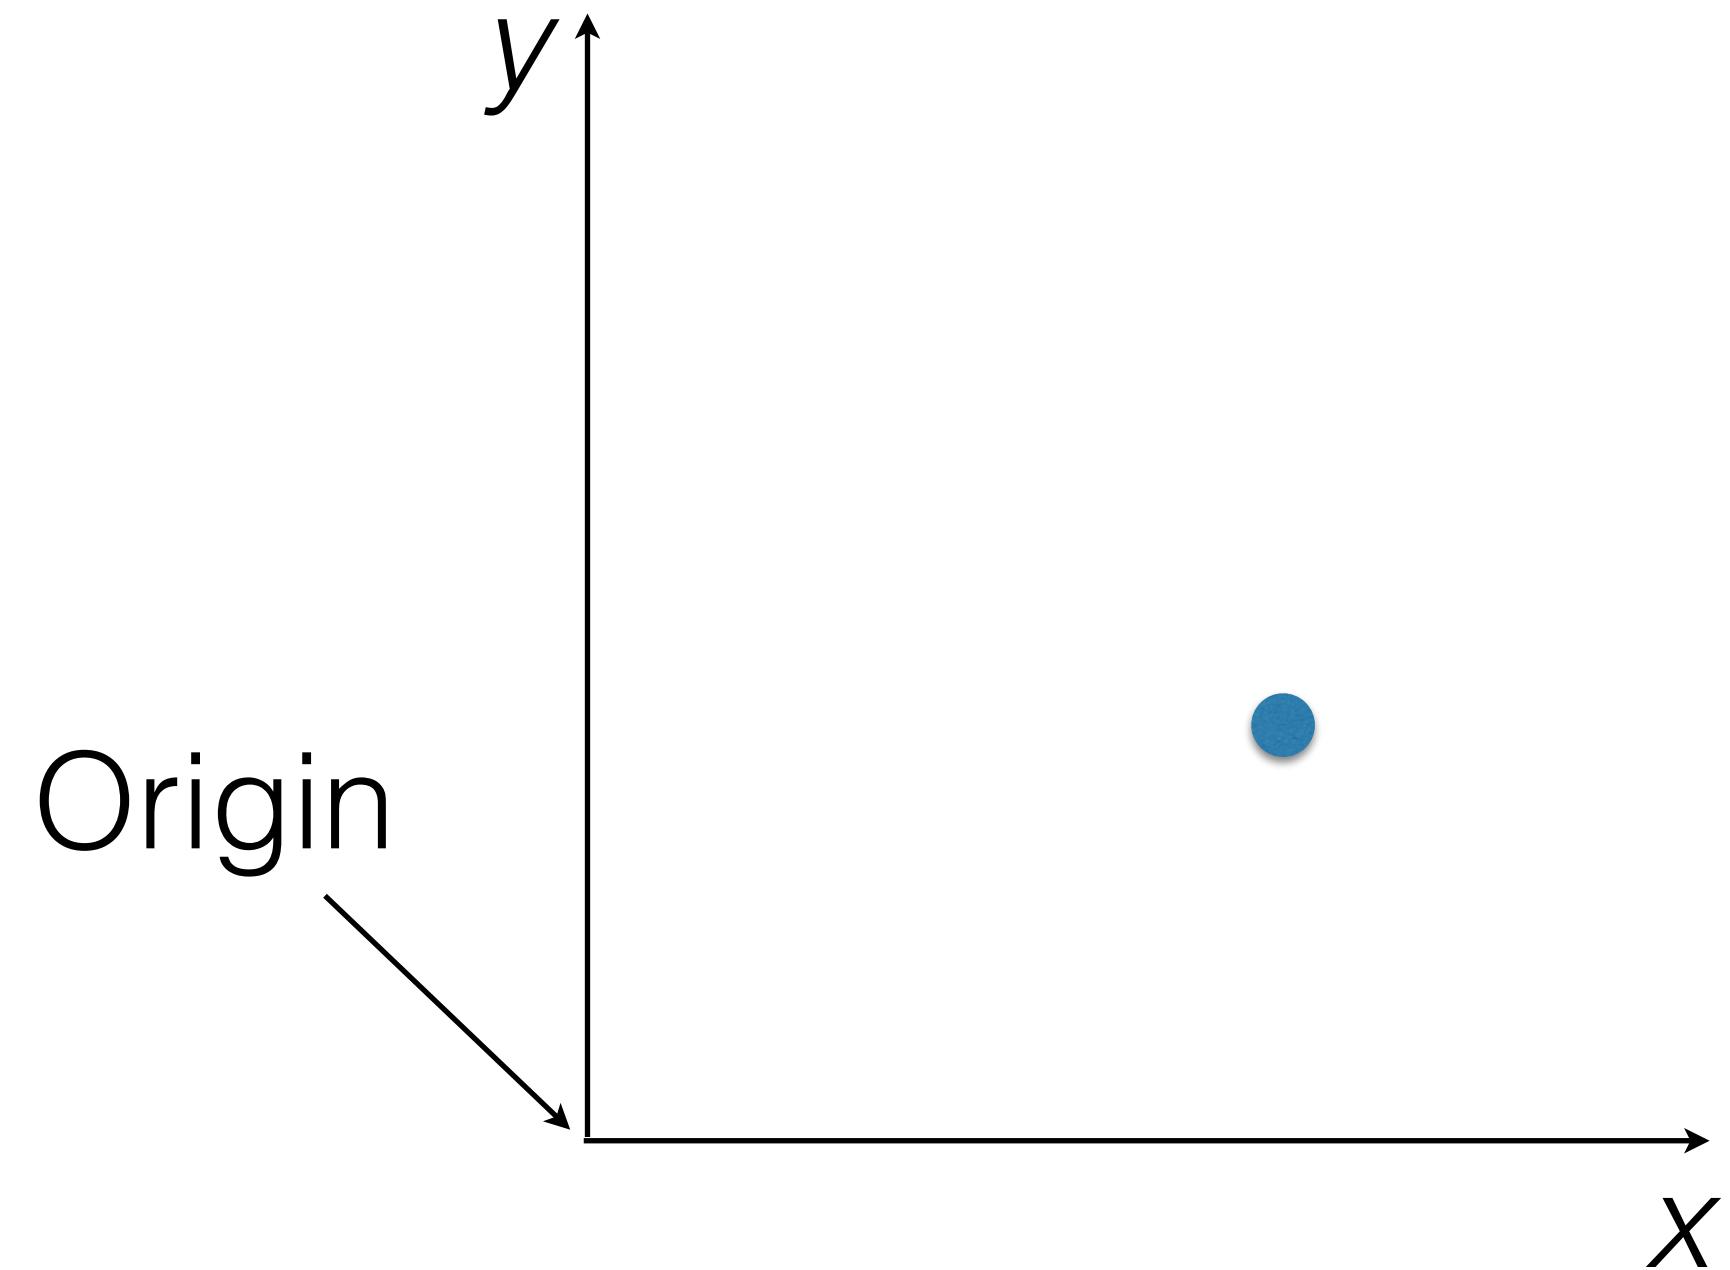
How do you describe this point numerically?

# Coordinate Systems



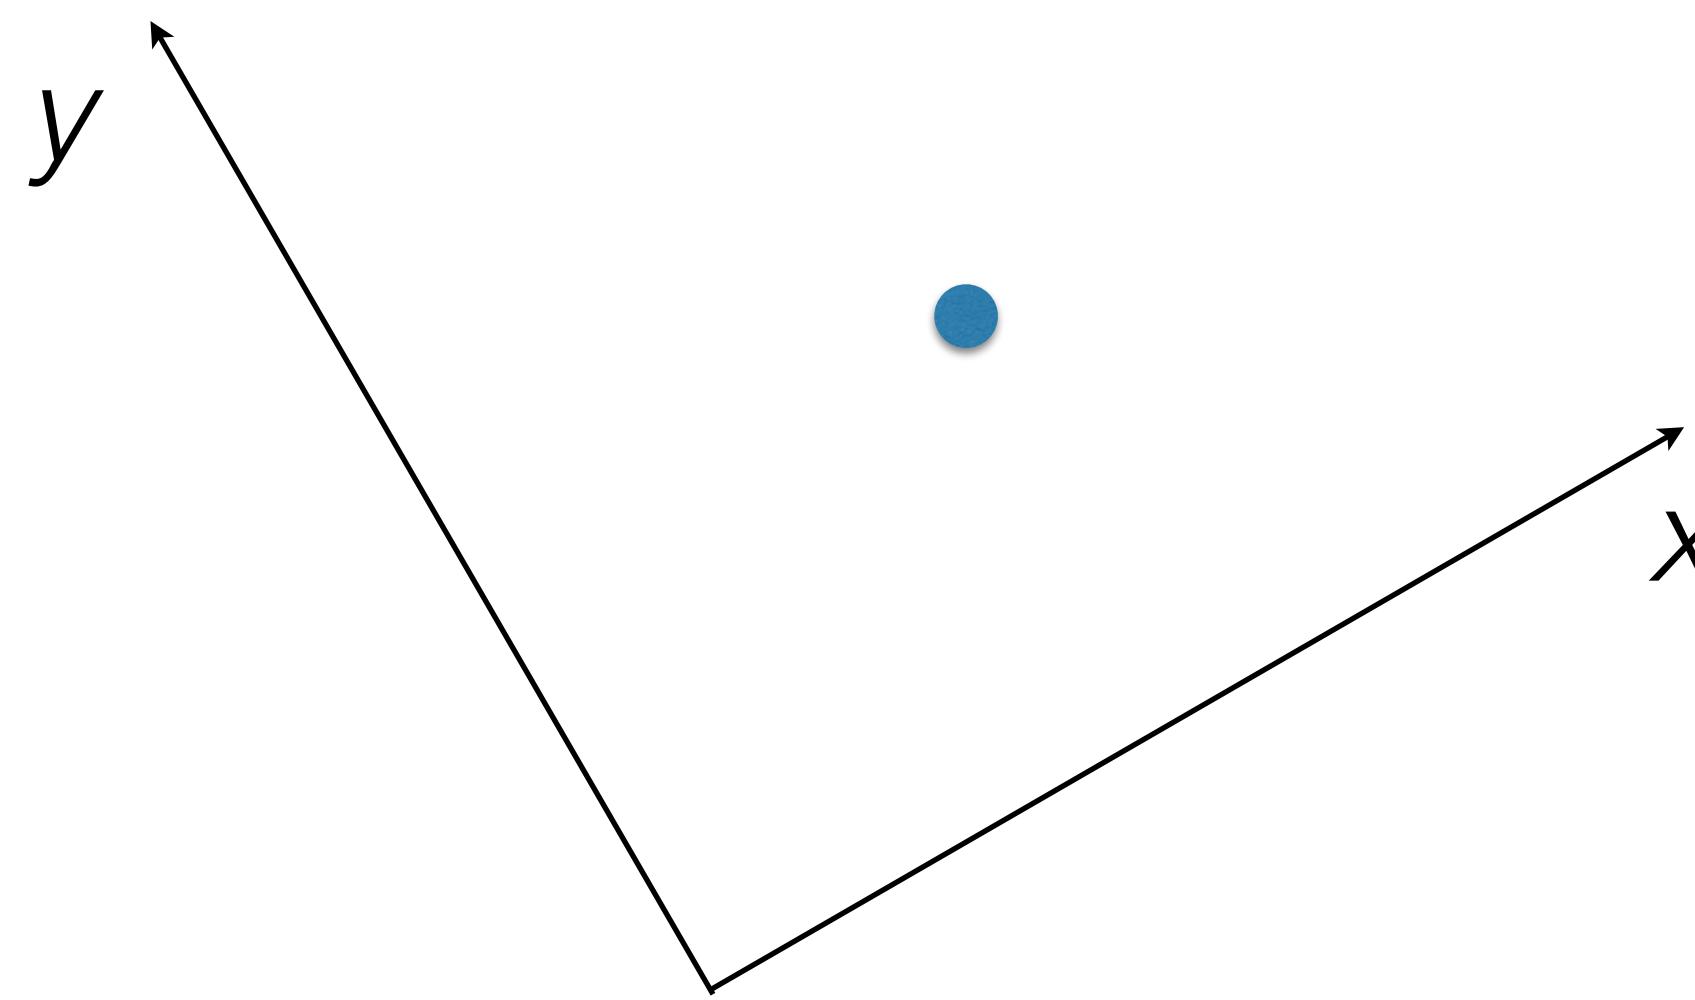
How do you describe this point numerically?

# Coordinate Systems



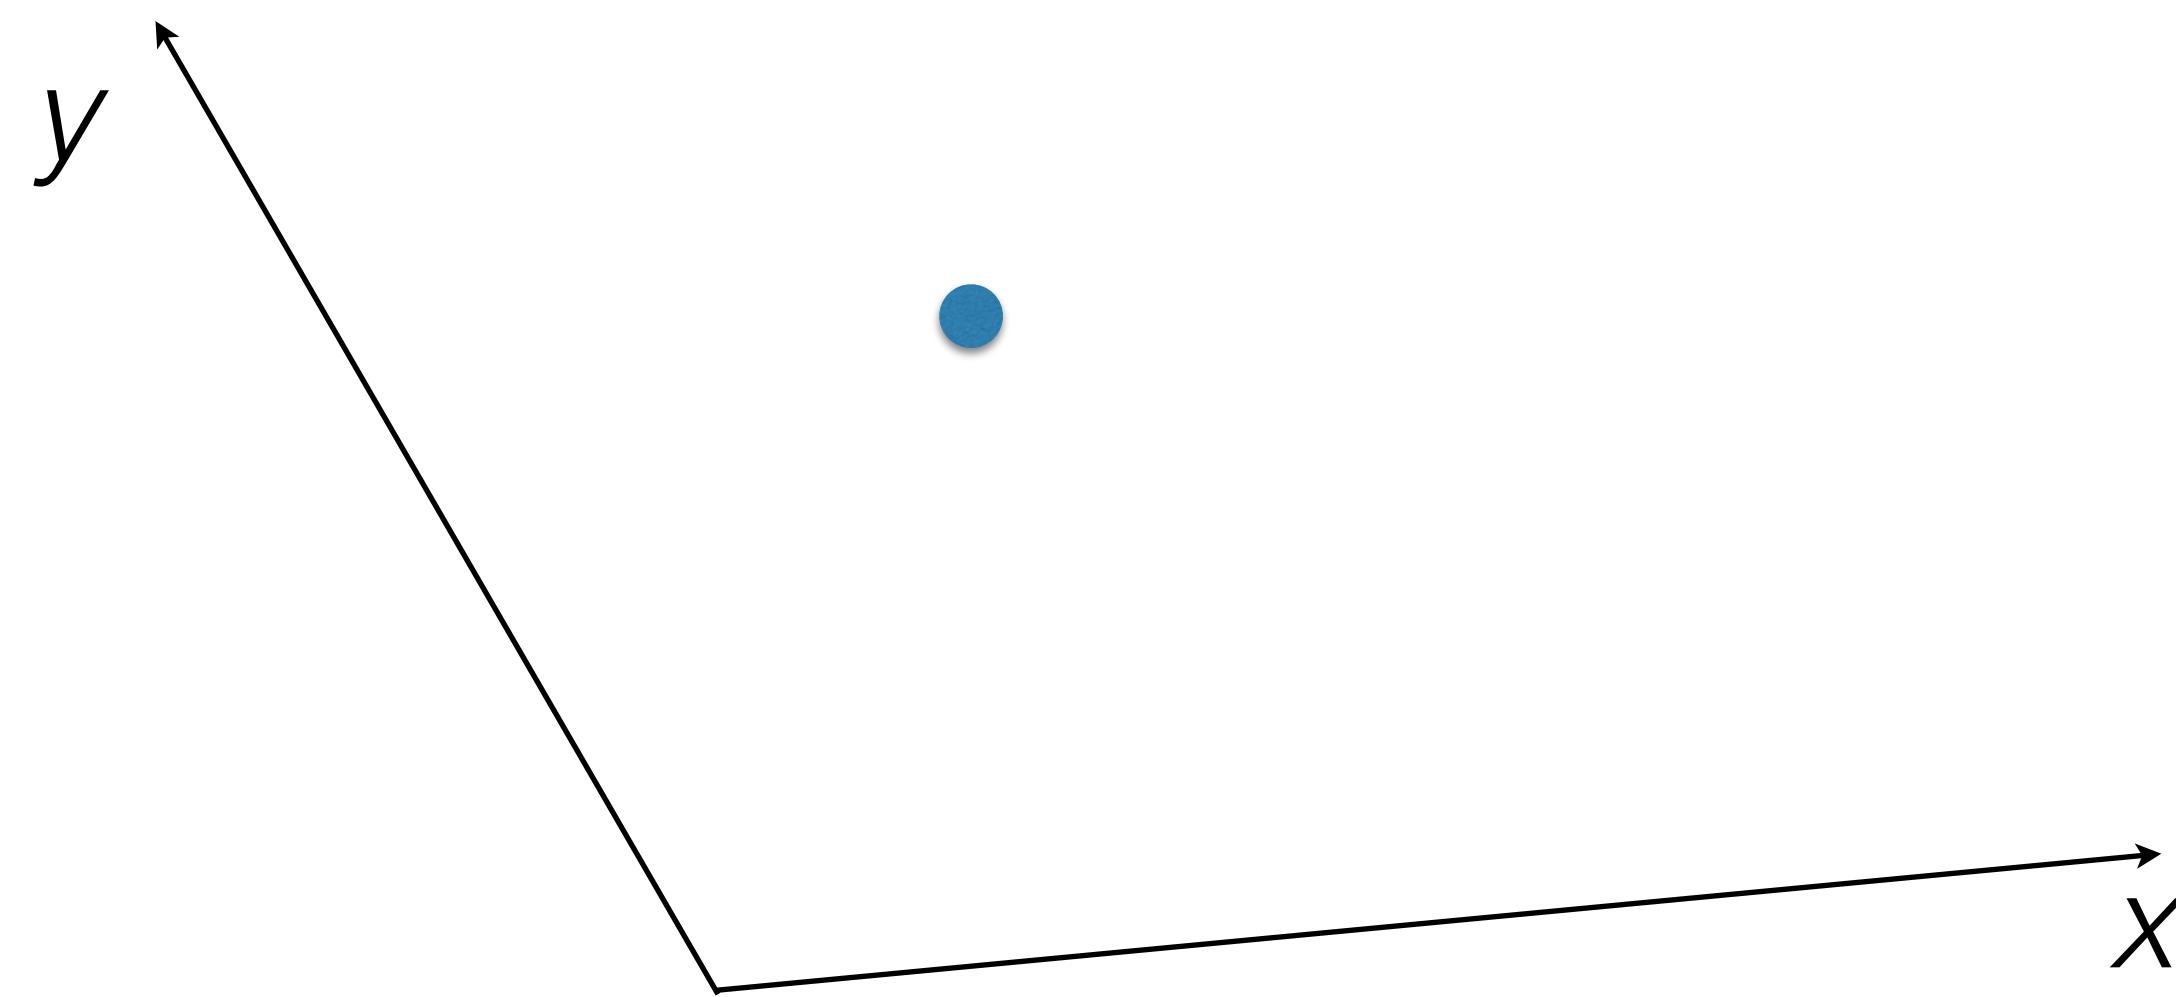
How about this coordinate system?

# Coordinate Systems



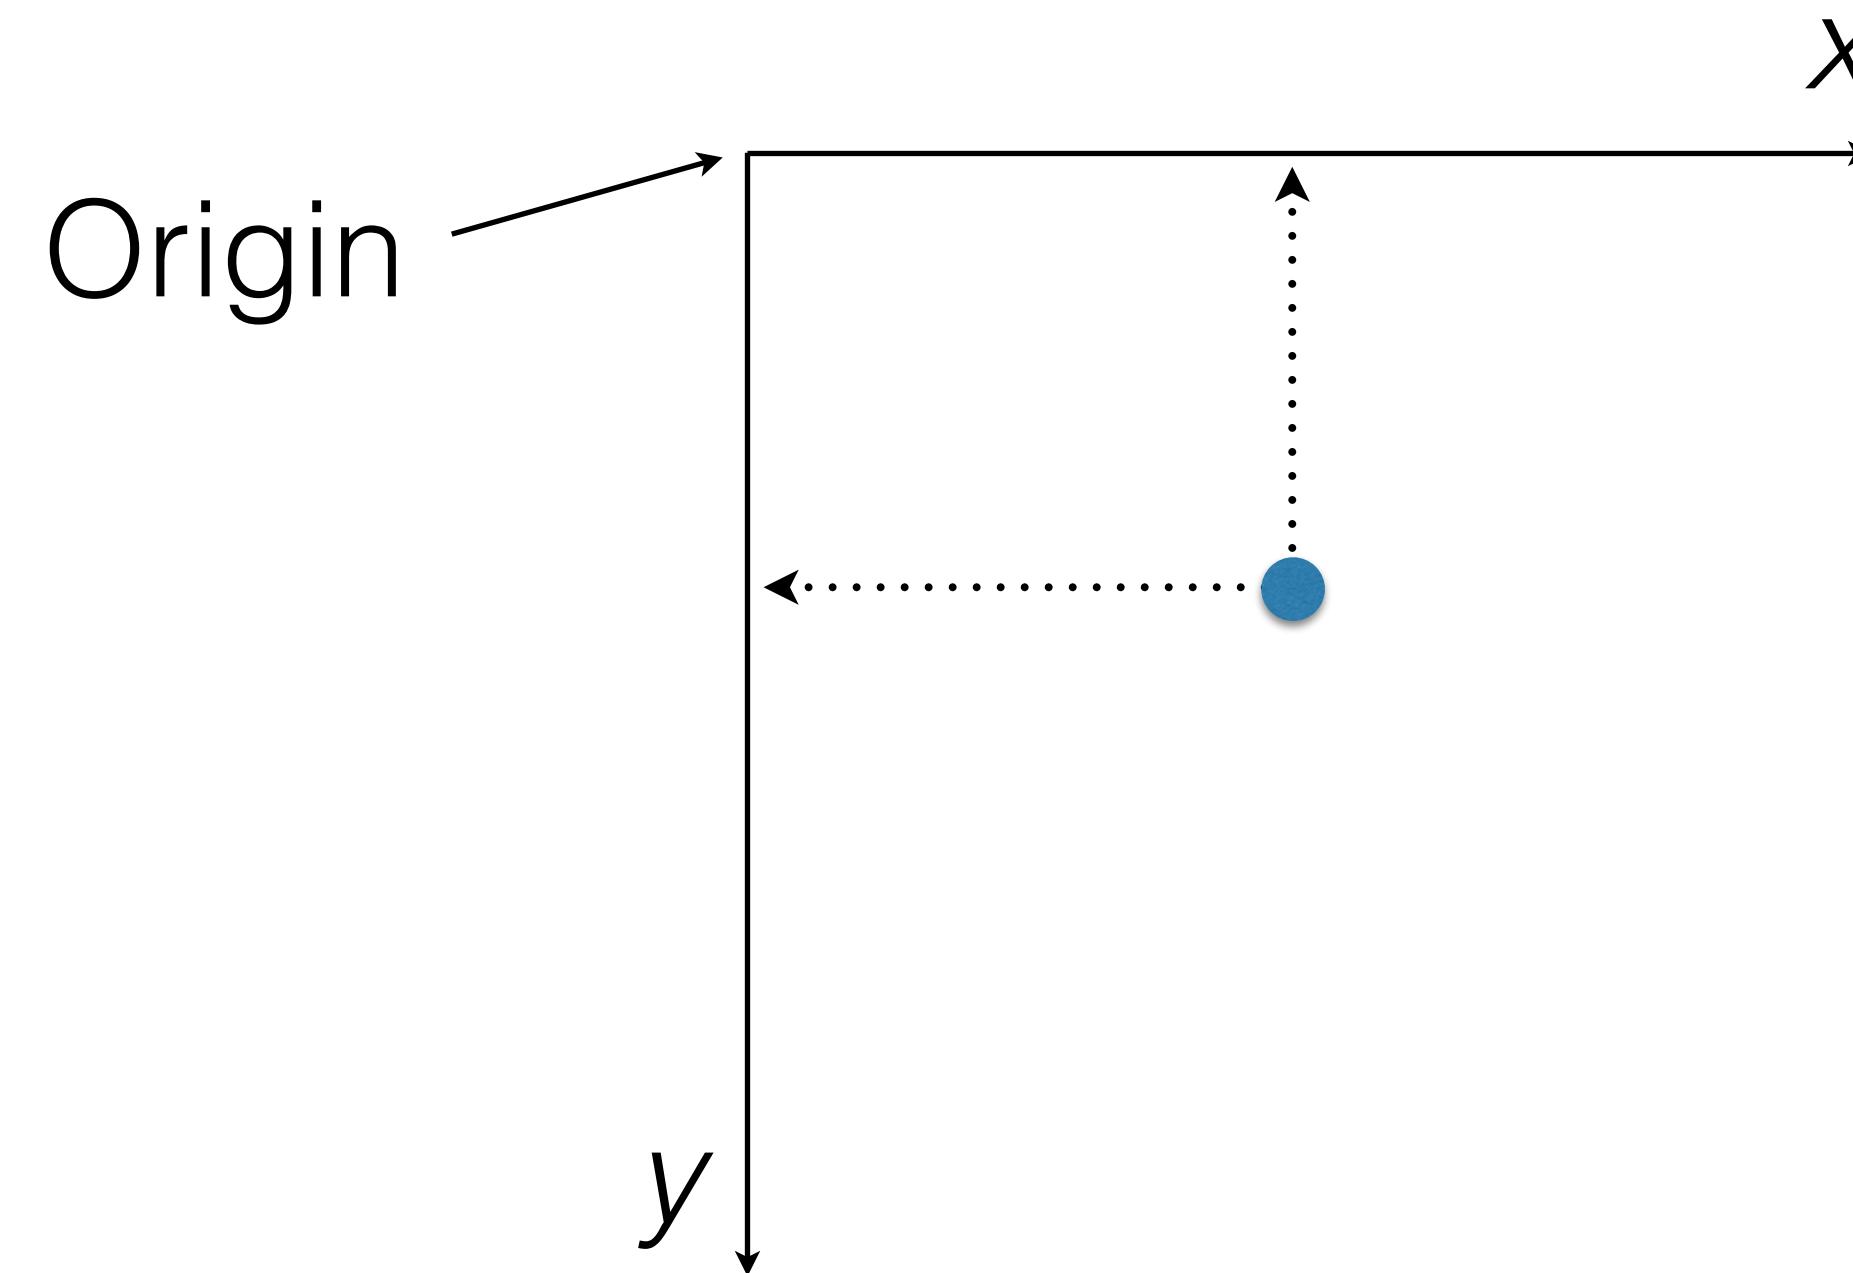
Or this one?

# Coordinate Systems



What about this one?

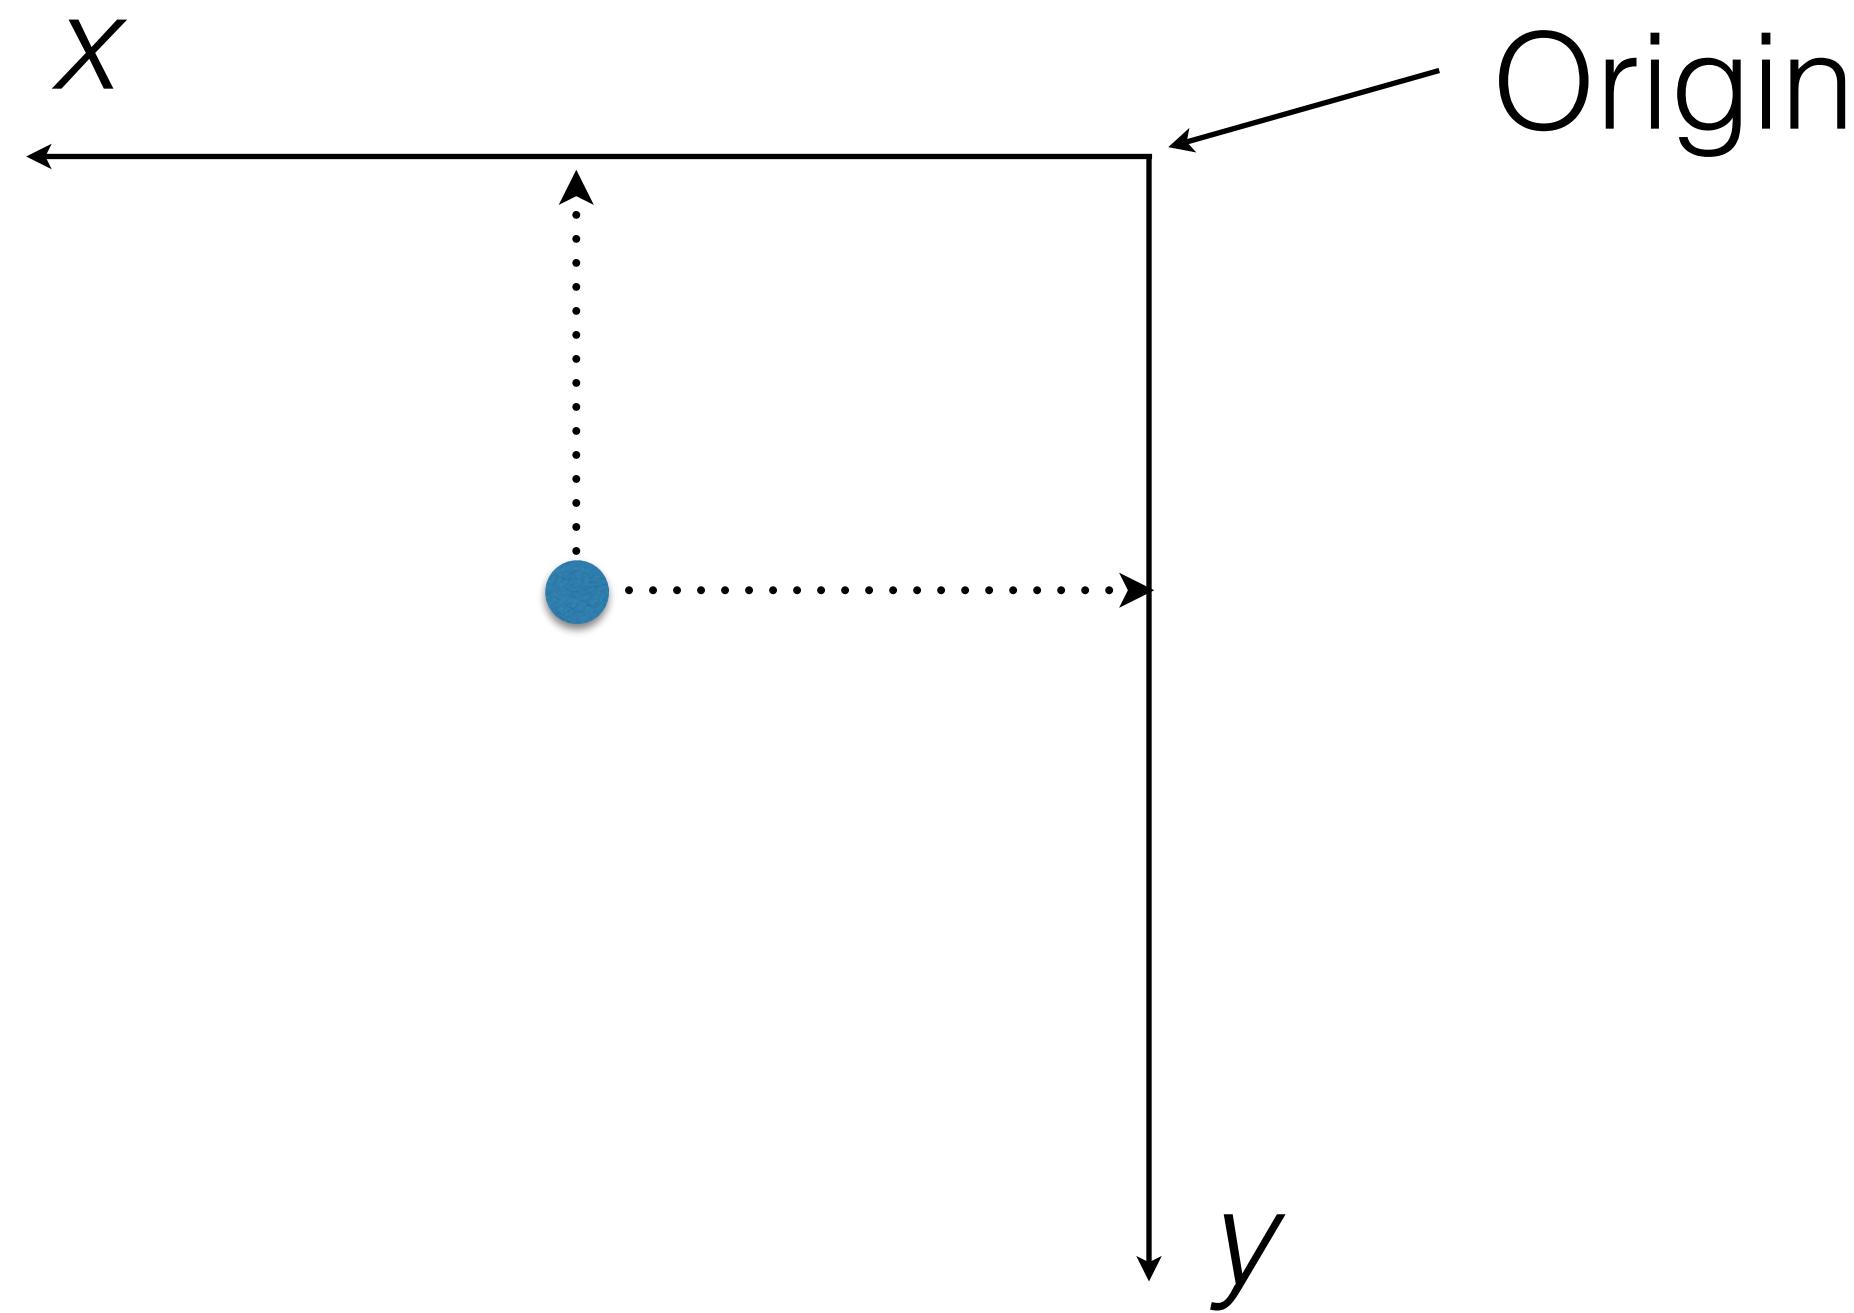
# Coordinate Systems



(This is what most screens use)

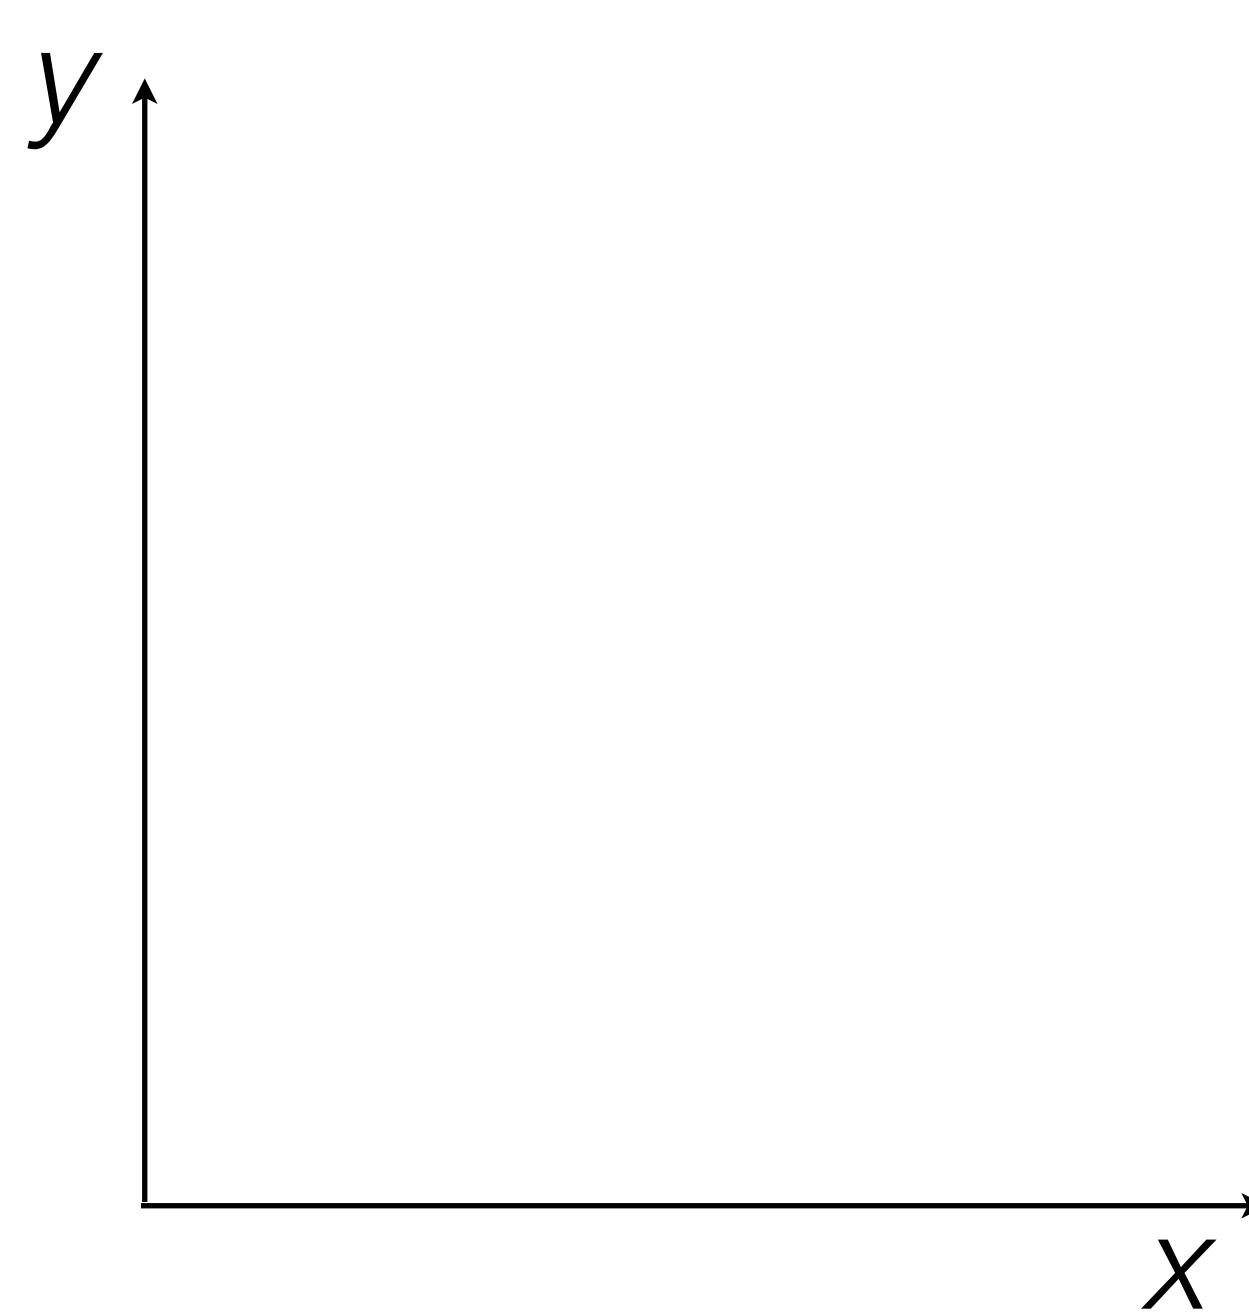
Why not this?

# Coordinate Systems

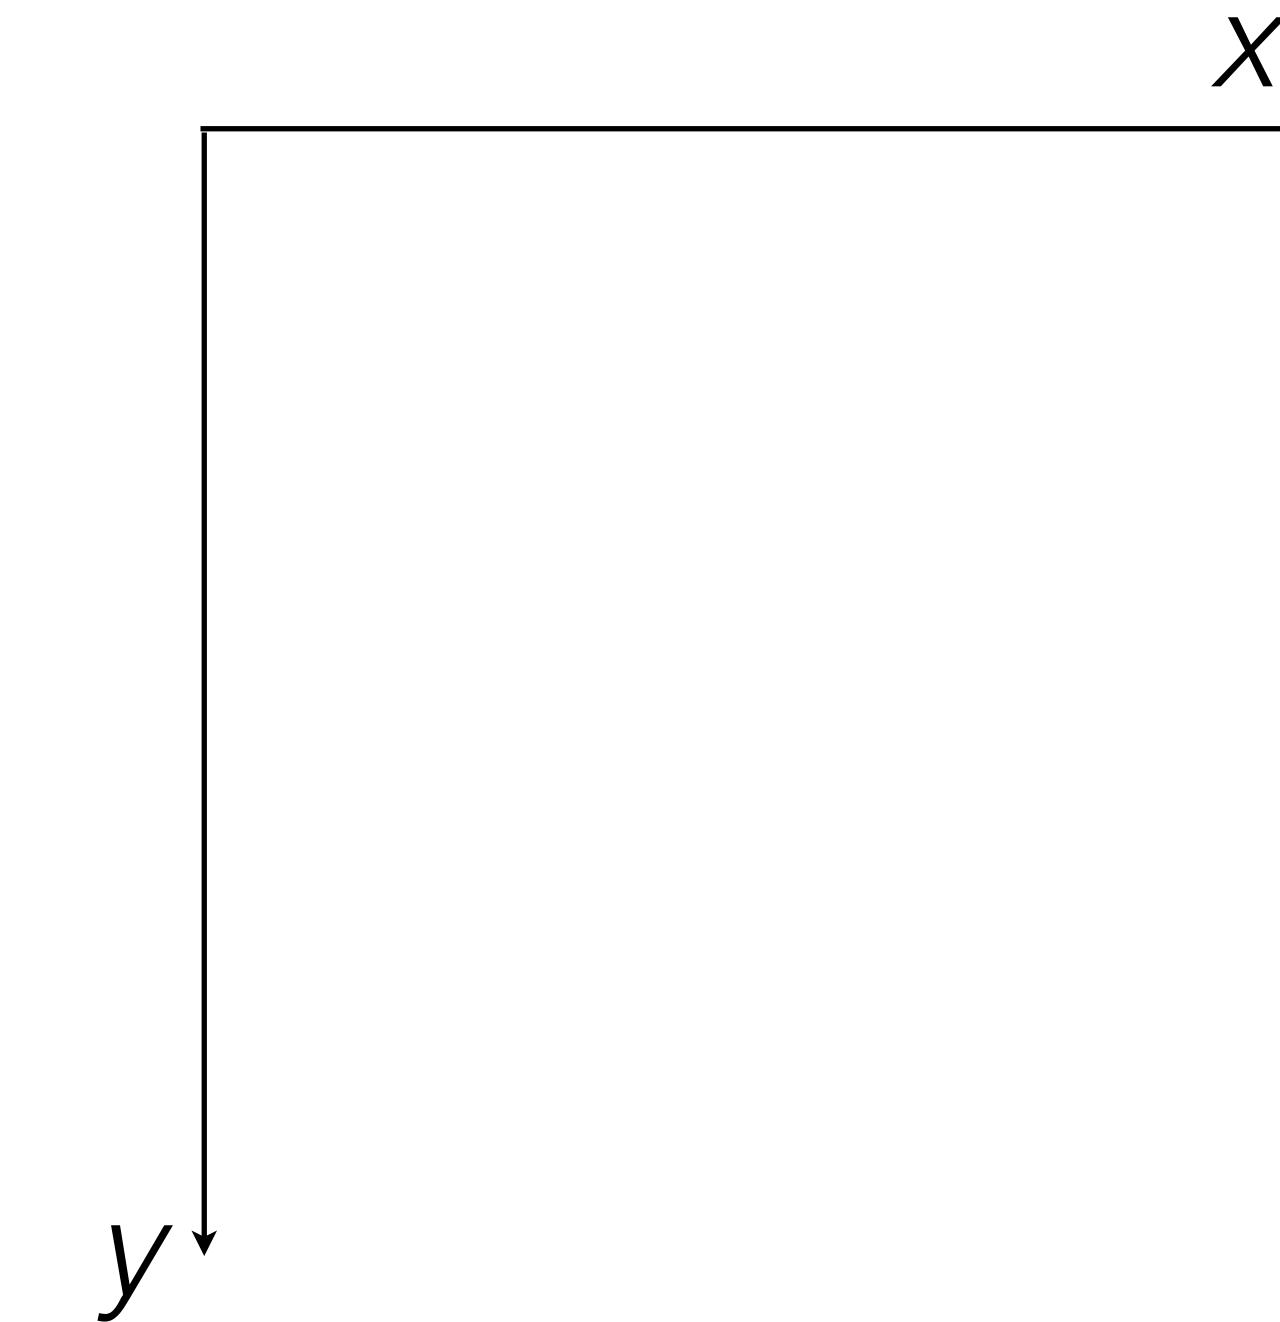


Or this?

# Coordinate Systems

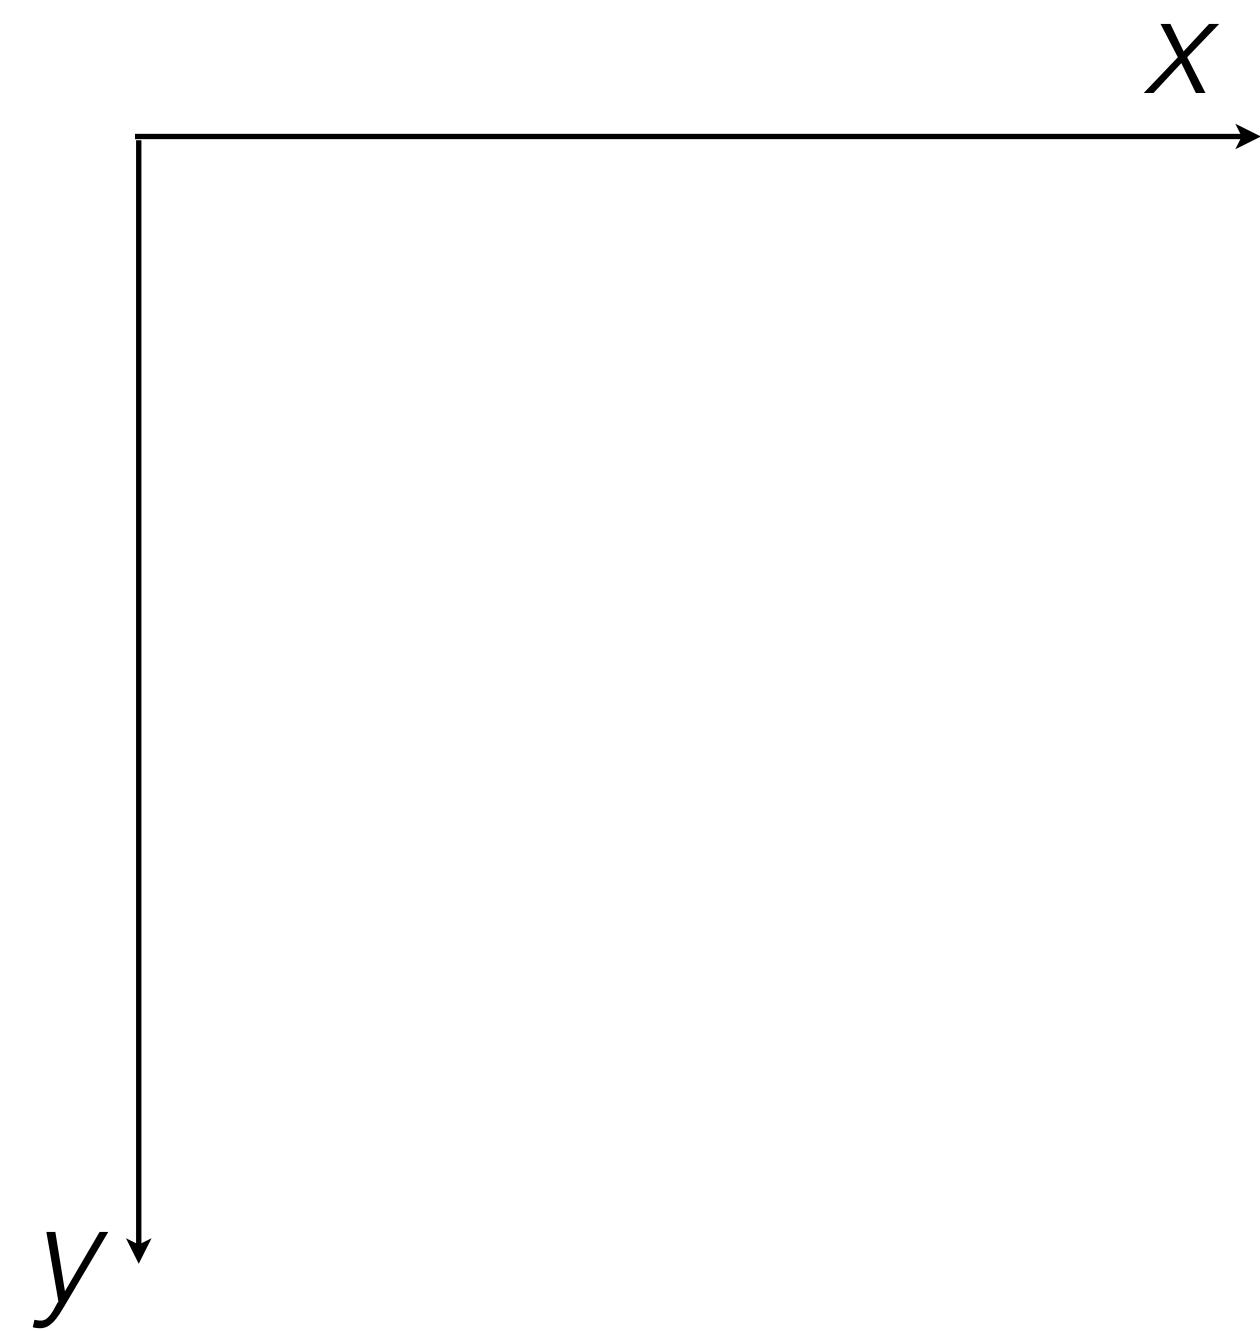


Math teachers  
("right handed")

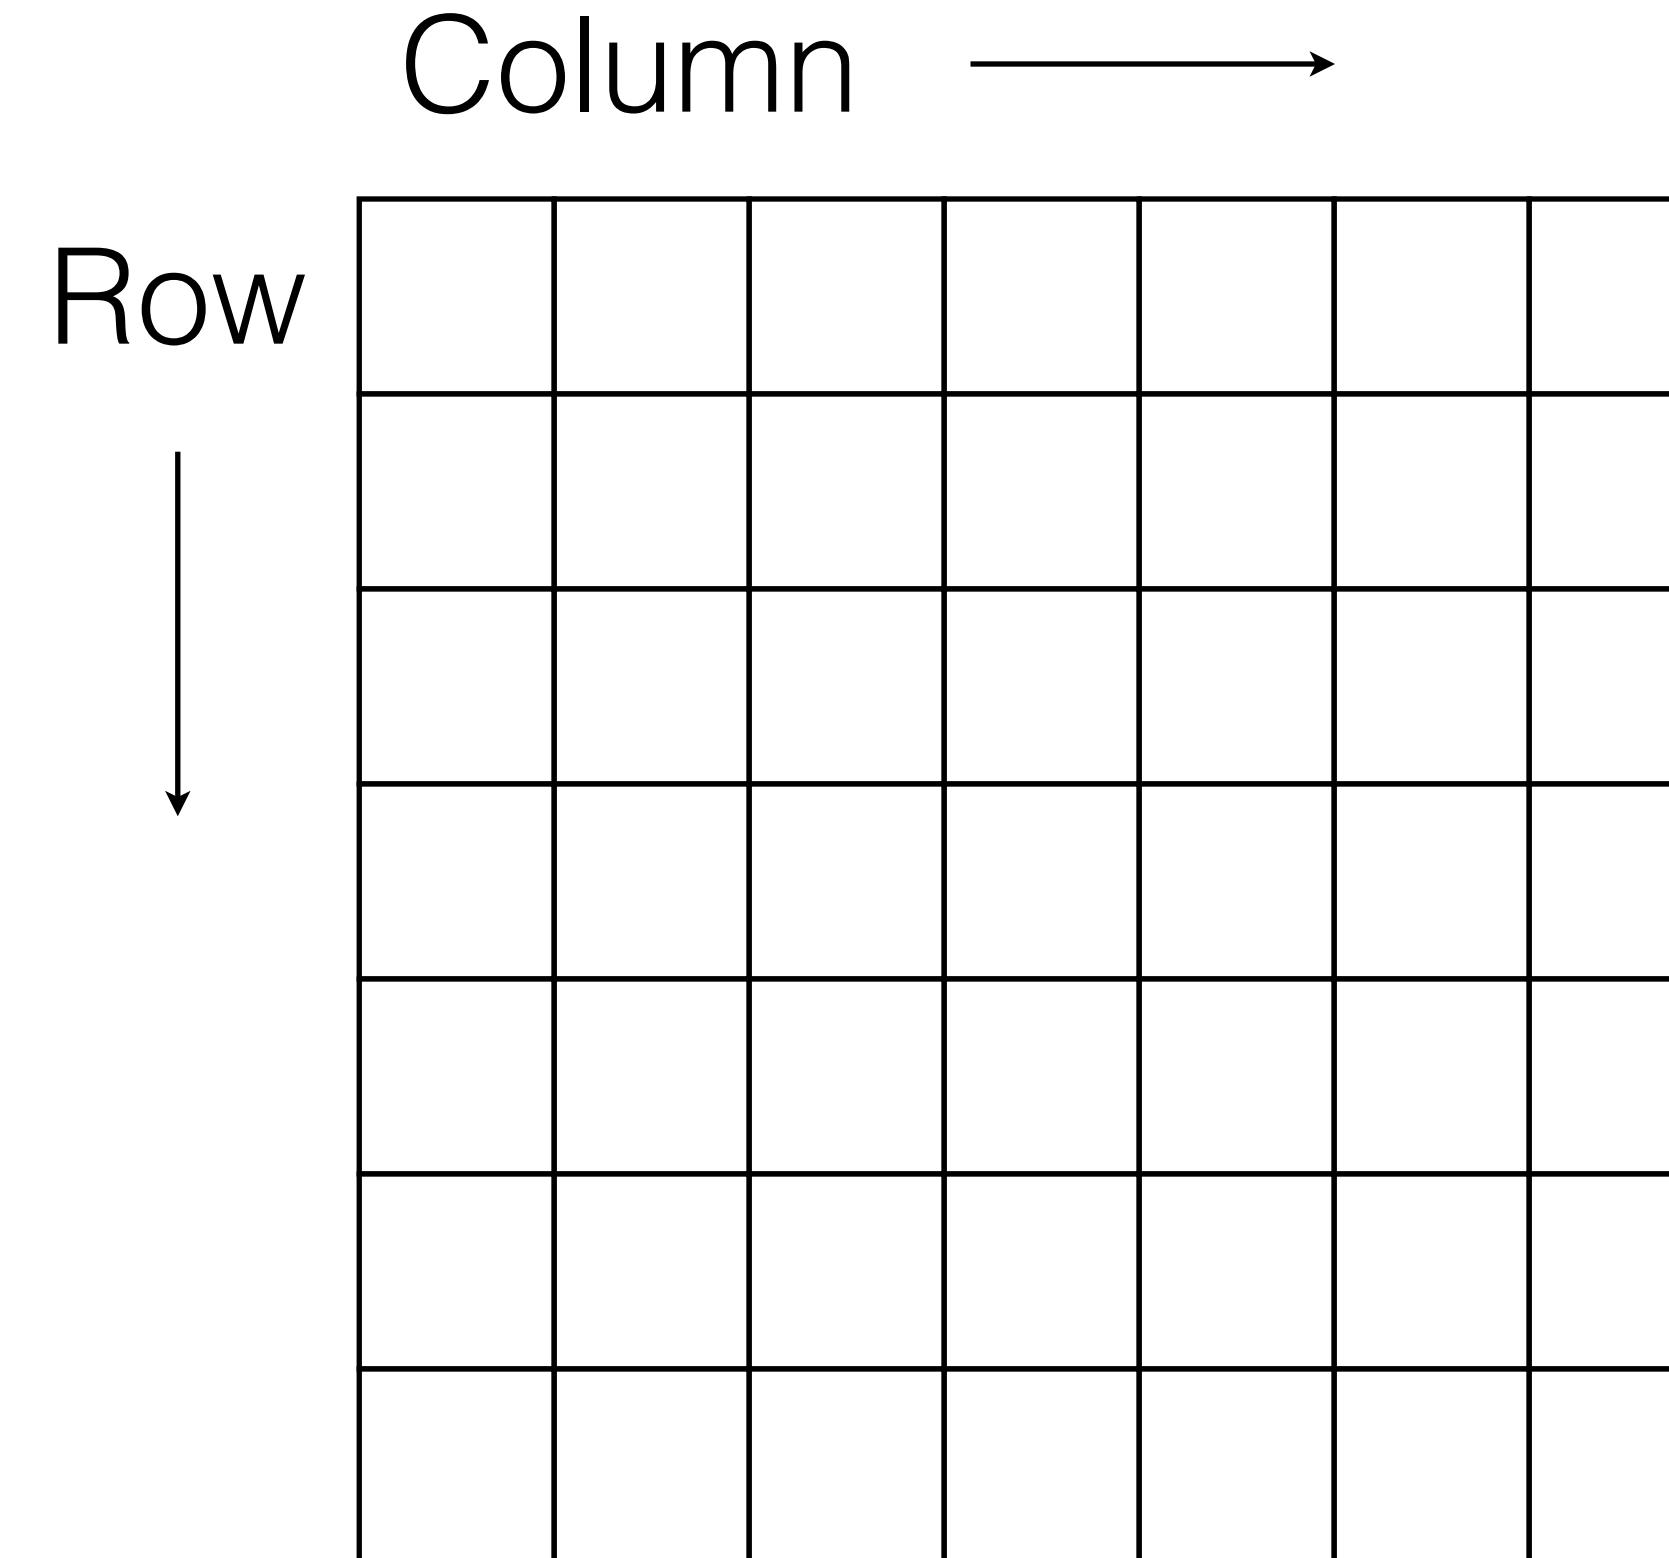


Computer screens  
("left handed")

# Math vs. Code



How we draw  
( $x,y$ )



How we store  
(row,col)

# What About 3D?

- Same ideas apply in 3D
- Right-handed vs. left-handed becomes more important  
(which way does z go?)

# Linear Transformations

# Translation

- *Translating* a coordinate system simply moves the origin
- Or can be thought of as moving the point
- Keeps the x and y directions the same
- Just add desired x, y offsets

$$(x', y') = (x + t_x, y + t_y)$$

OR

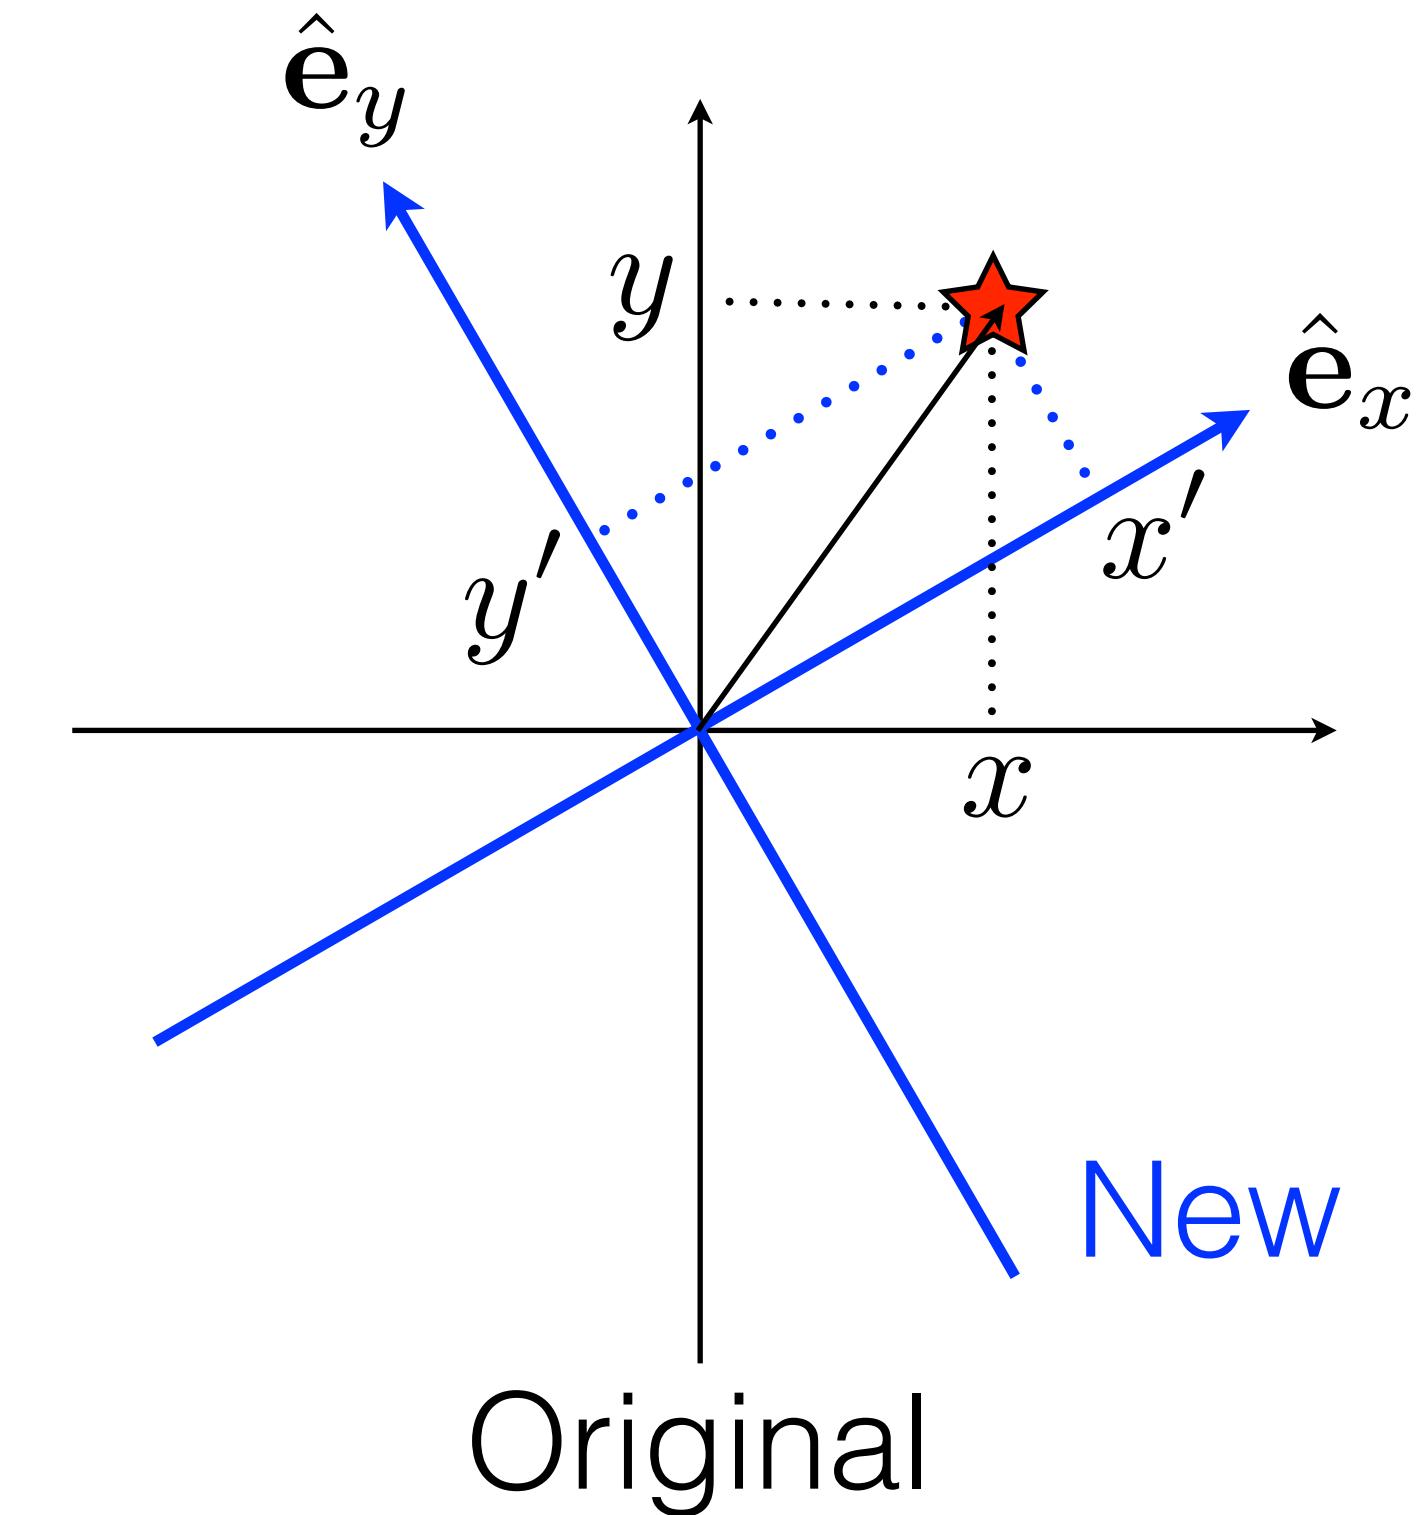
$$\mathbf{p}' = \mathbf{p} + \mathbf{t}$$

$$\mathbf{t} = \begin{bmatrix} t_x \\ t_y \end{bmatrix}$$

# Rotation

- Can again thing interchangeably about rotating a coordinate system and rotating points around the origin of that coordinate system
- To rotate a point's coordinates in a rotated system, project the point on to the new coordinate axes:

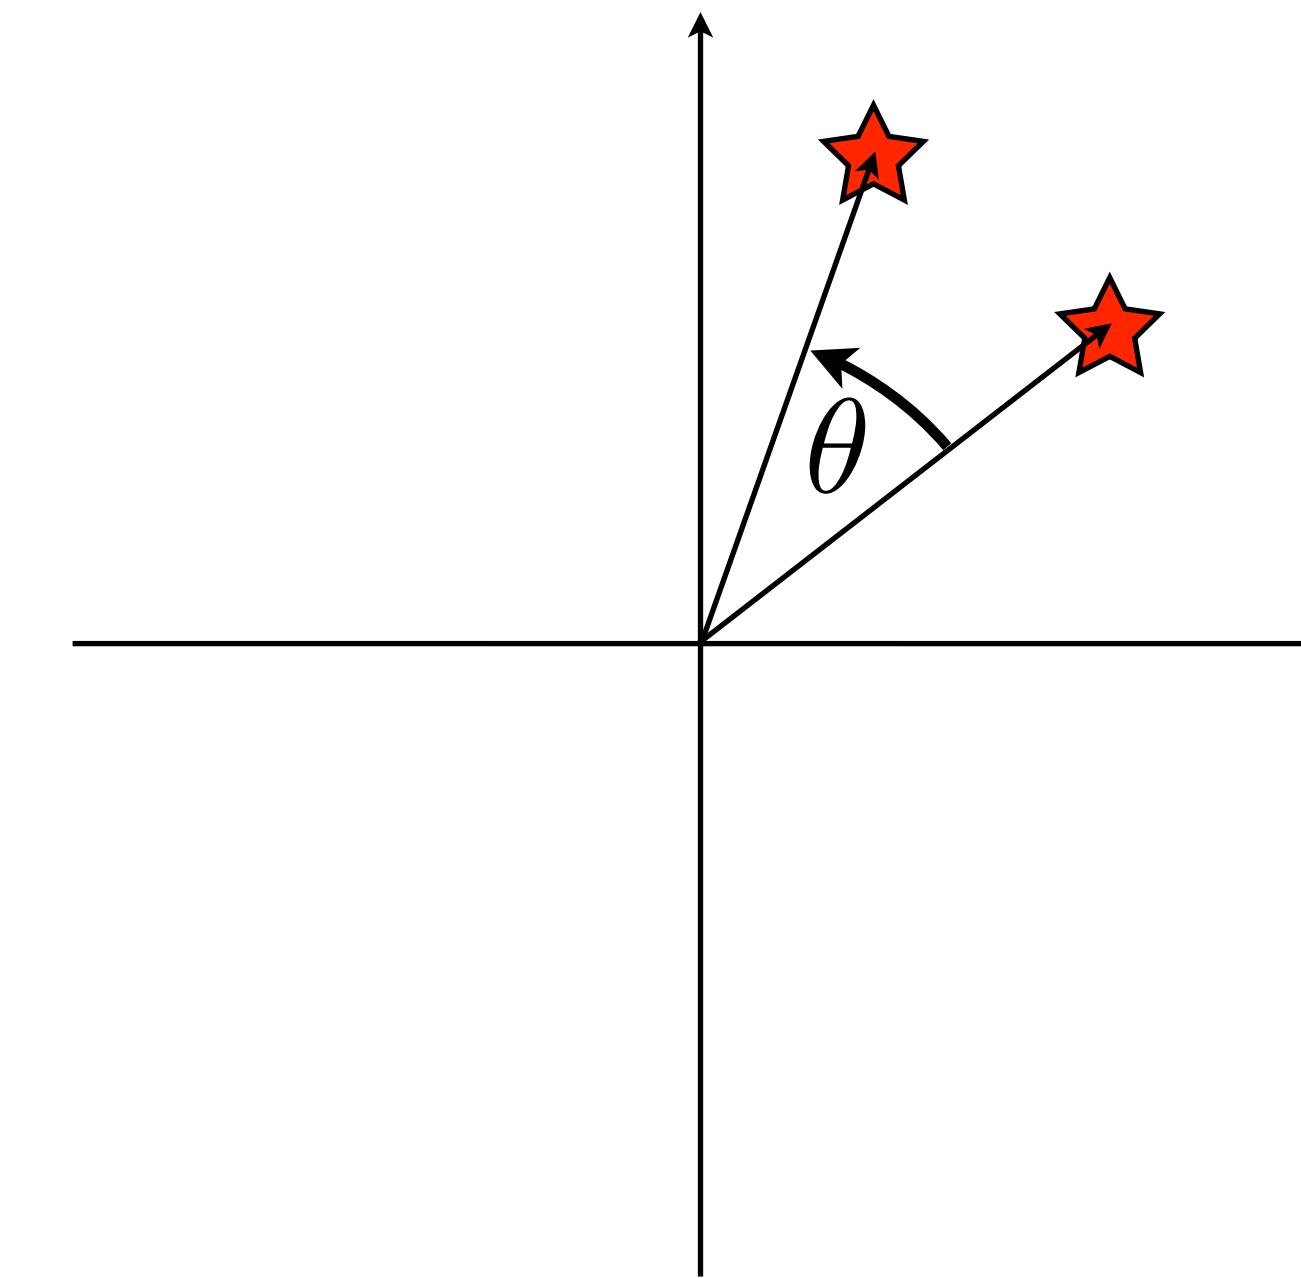
$$\mathbf{p}' = \begin{bmatrix} e_{x1} & e_{x2} \\ e_{y1} & e_{y2} \end{bmatrix} \mathbf{p}$$



# Rotation

$$\mathbf{R}(\theta) = \begin{bmatrix} \cos \theta & -\sin \theta \\ \sin \theta & \cos \theta \end{bmatrix}$$

$$\mathbf{R}^{-1}(\theta) = \begin{bmatrix} \cos \theta & \sin \theta \\ -\sin \theta & \cos \theta \end{bmatrix}$$

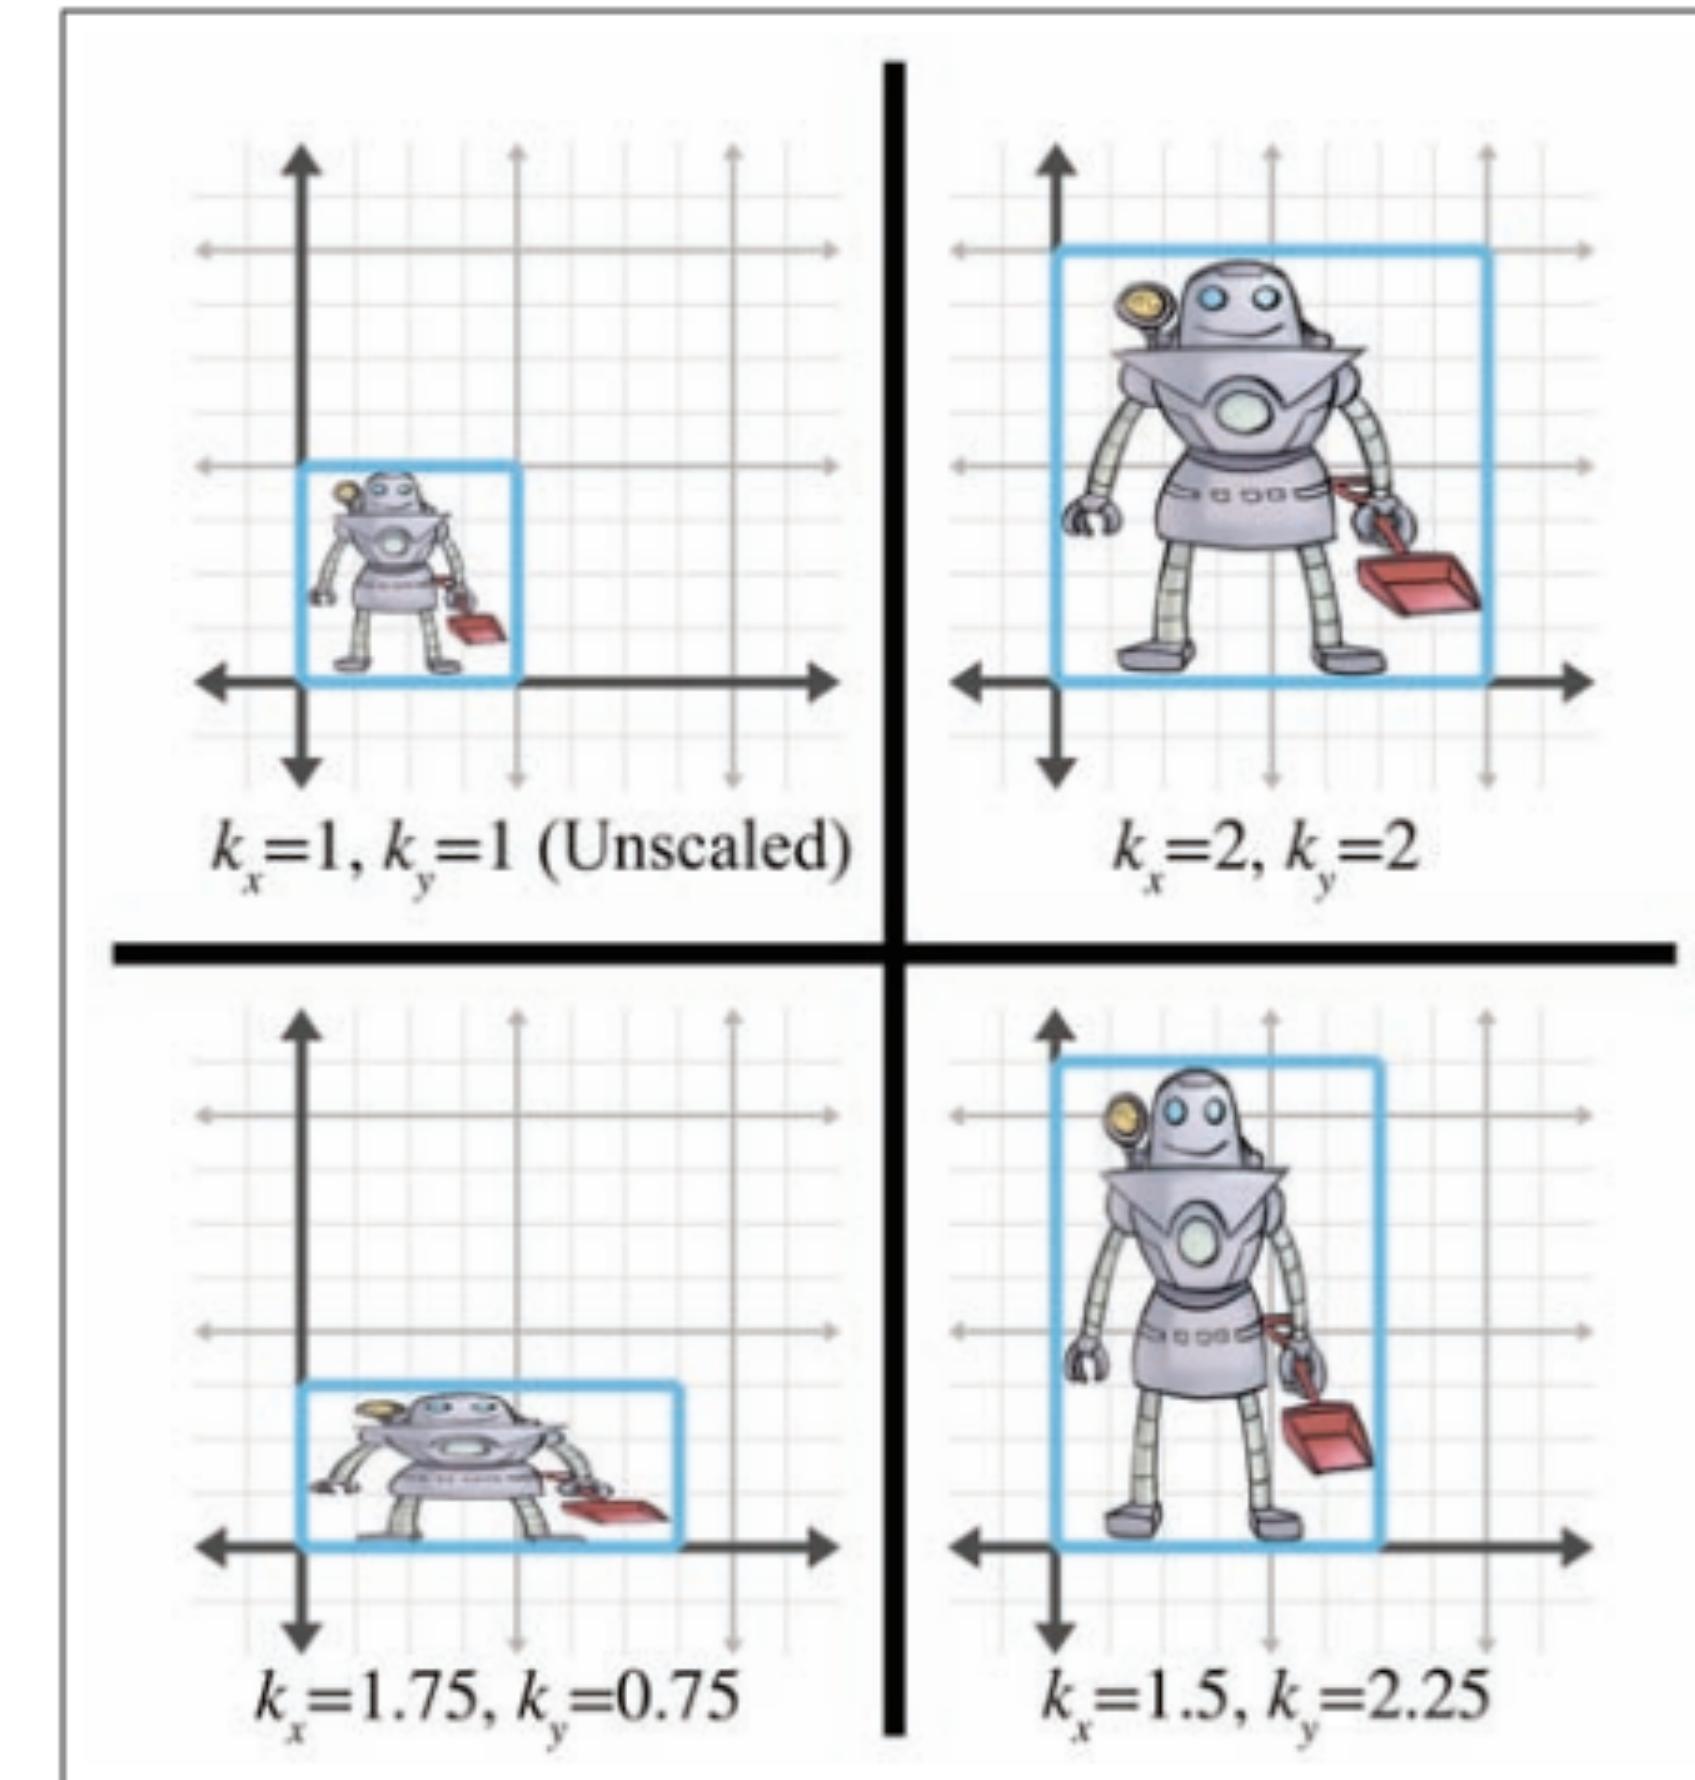


$$\begin{bmatrix} x' \\ y' \end{bmatrix} = \begin{bmatrix} \cos \theta & -\sin \theta \\ \sin \theta & \cos \theta \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix}$$

# Scaling

$$\begin{bmatrix} s & 0 \\ 0 & s \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix} = \begin{bmatrix} sx \\ sy \end{bmatrix}$$

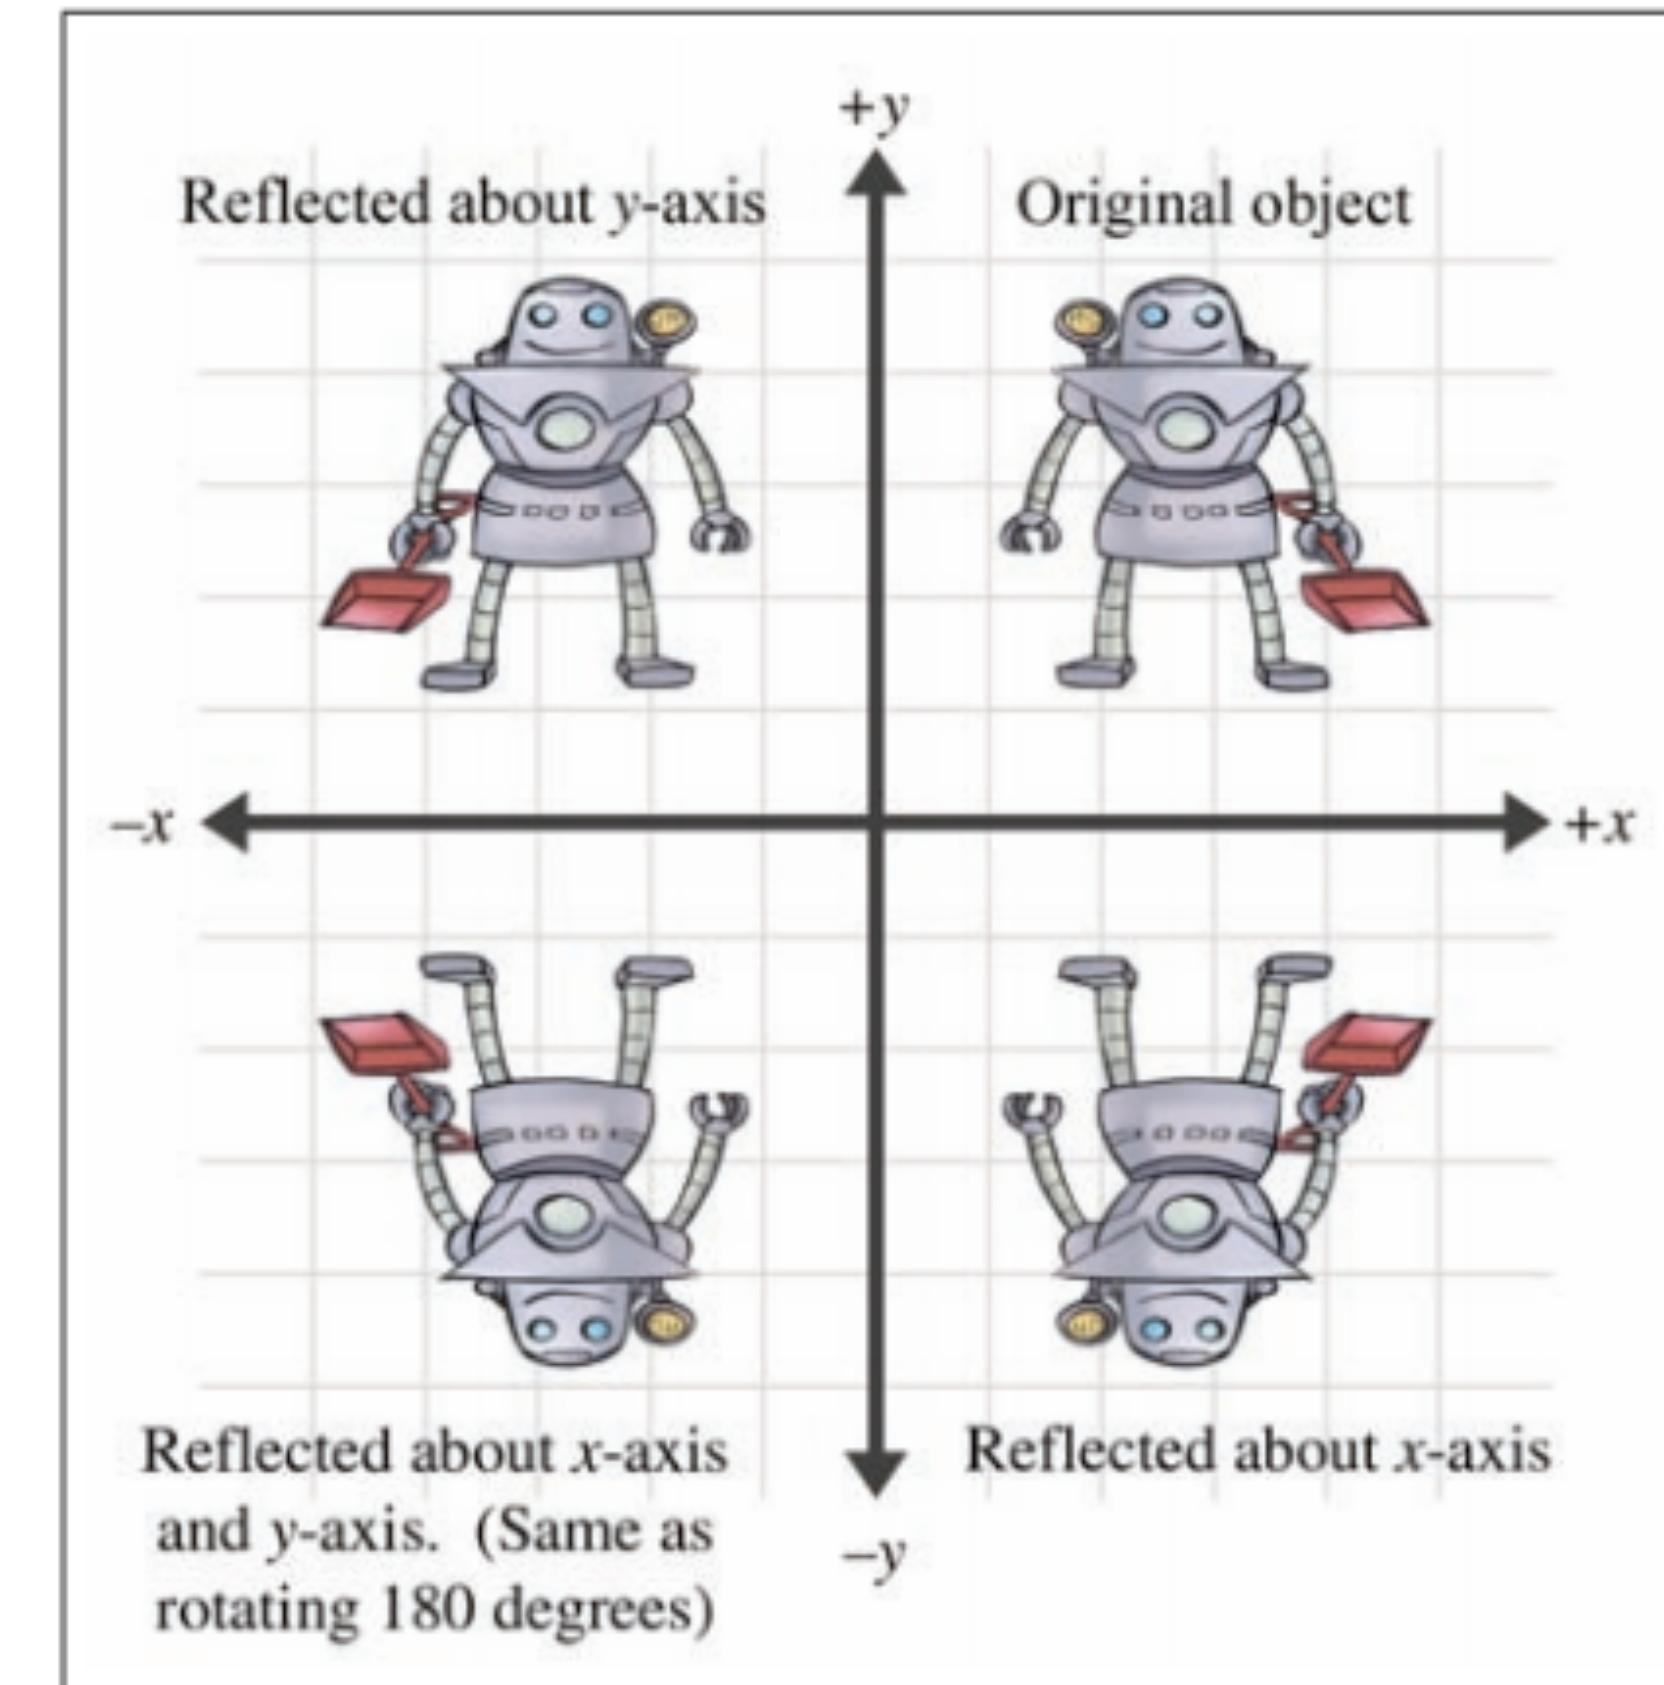
$$\begin{bmatrix} s_x & 0 \\ 0 & s_y \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix} = \begin{bmatrix} s_x x \\ s_y y \end{bmatrix}$$



# Reflection

$$\begin{bmatrix} -1 & 0 \\ 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix} = \begin{bmatrix} -x \\ y \end{bmatrix}$$

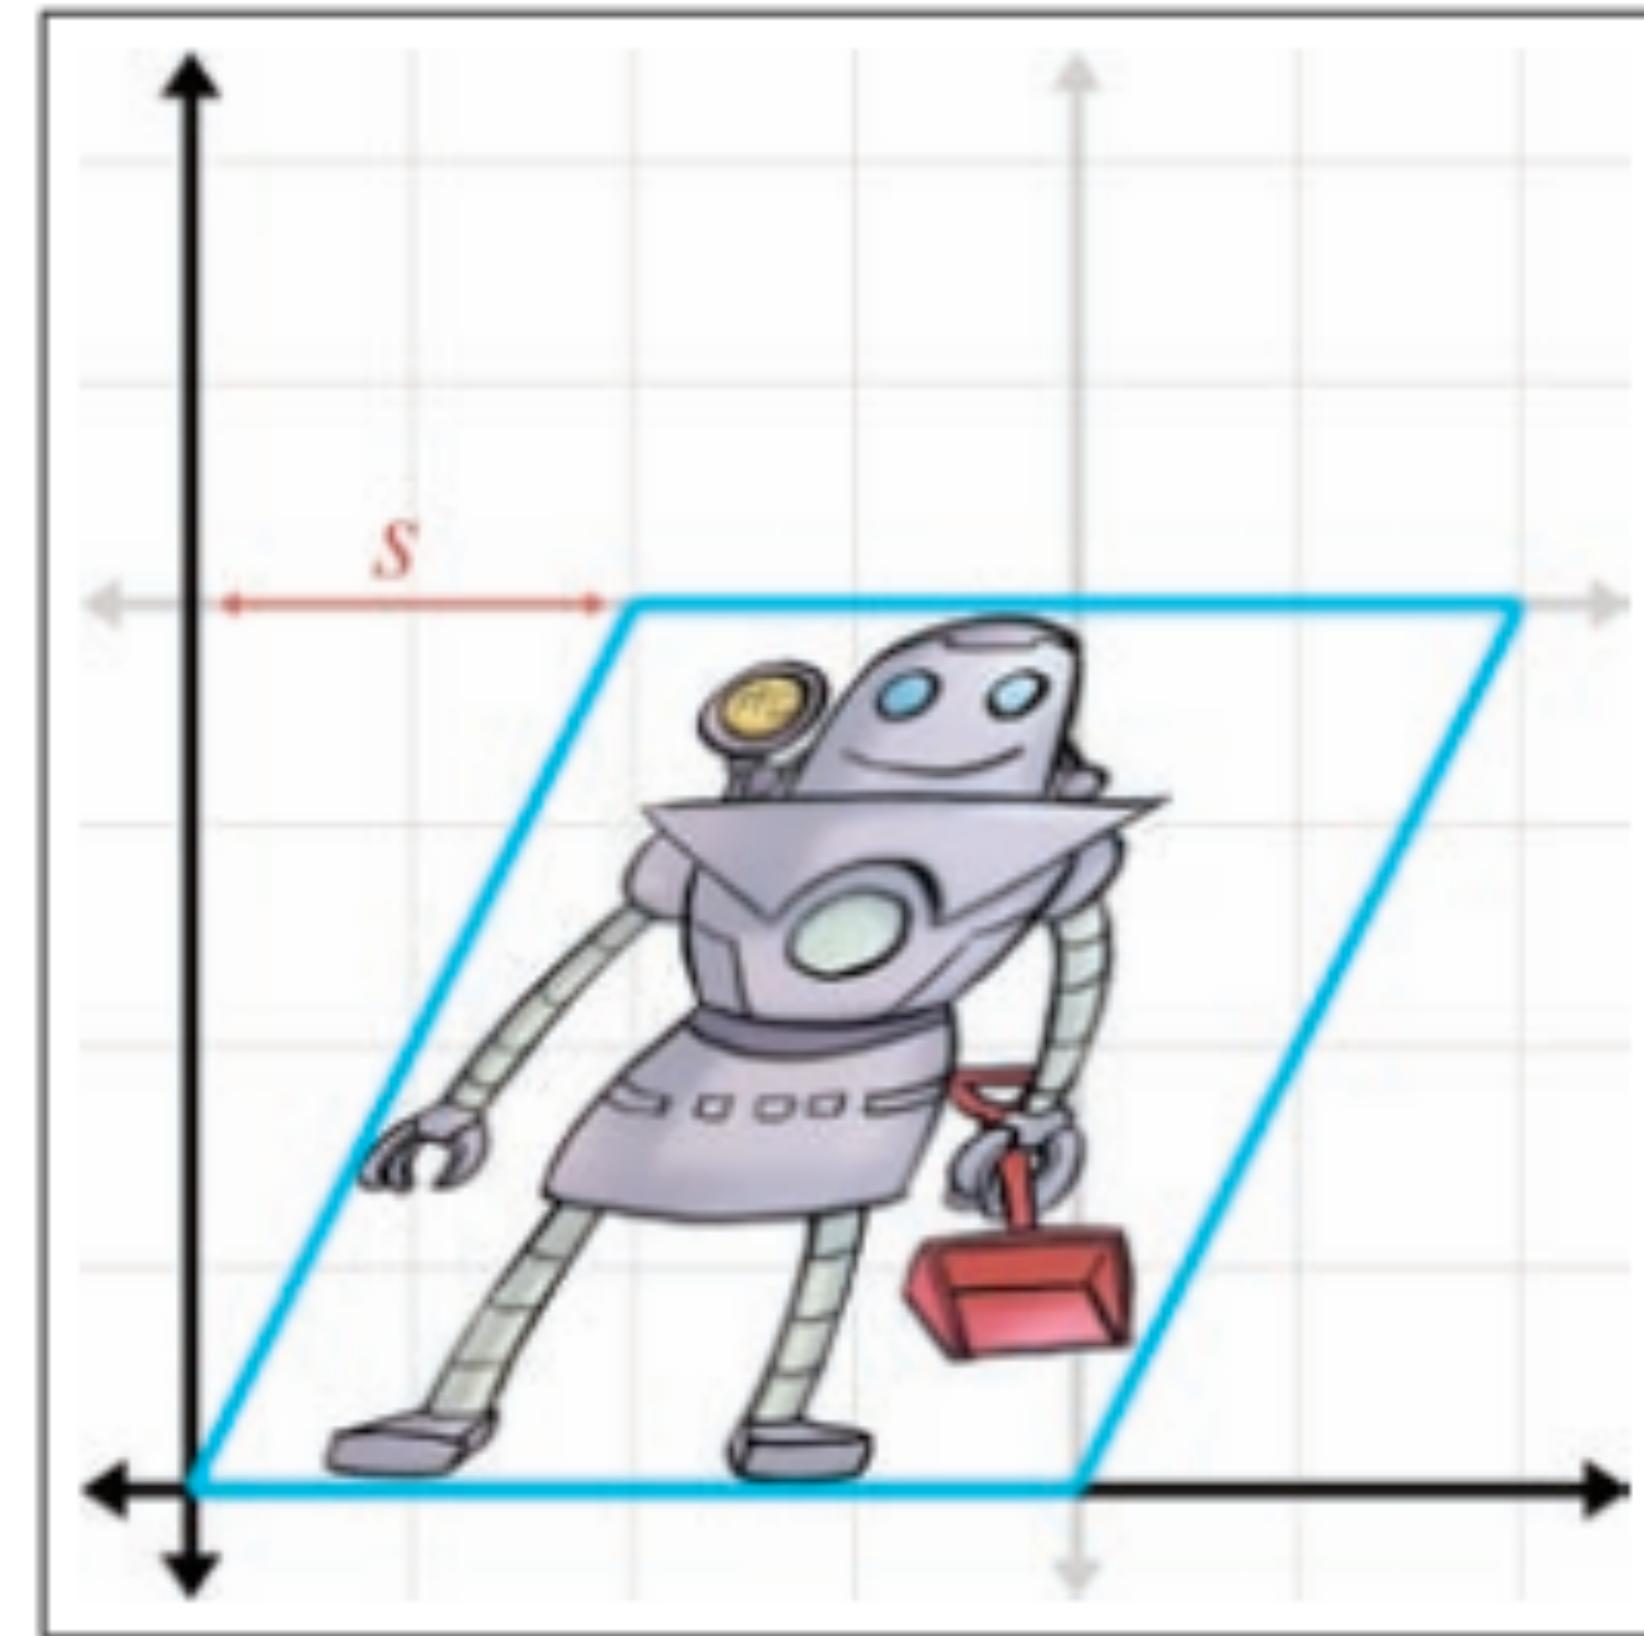
$$\begin{bmatrix} 1 & 0 \\ 0 & -1 \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix} = \begin{bmatrix} x \\ -y \end{bmatrix}$$



# Shearing

$$\begin{bmatrix} 1 & s \\ 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix} = \begin{bmatrix} x + sy \\ y \end{bmatrix}$$

$$\begin{bmatrix} 1 & 0 \\ s & 1 \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix} = \begin{bmatrix} x \\ sx + y \end{bmatrix}$$



Sometimes called a skew transform

# Composition

$$\begin{aligned} \mathbf{M}_2 &= \begin{bmatrix} s_x & 0 \\ 0 & s_y \end{bmatrix} & \mathbf{M}_1 &= \begin{bmatrix} \cos \theta & -\sin \theta \\ \sin \theta & \cos \theta \end{bmatrix} \\ &\searrow &&\downarrow \\ \mathbf{M}_{12} &= \mathbf{M}_2 \mathbf{M}_1 &&\longleftarrow \end{aligned}$$

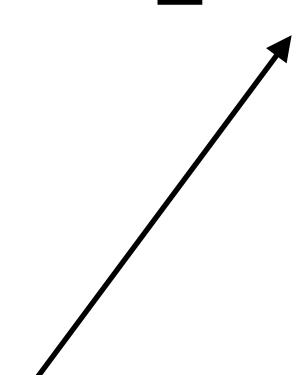
$$\mathbf{M}_{12}\mathbf{p} = (\mathbf{M}_2\mathbf{M}_1)\mathbf{p} = \mathbf{M}_2(\mathbf{M}_1\mathbf{p})$$

# But...

- Composition of transformations (matrix multiplication) is really useful
- But translation must be done as a separate operation
- *Can you include translation as a linear transformation?*

# Homogeneous Coordinates

- Can include translation as a linear transformation using homogeneous coordinates
- Also useful for perspective (coming up in a bit)

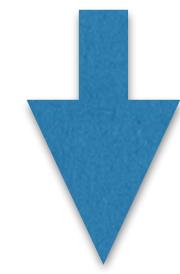
$$\mathbf{p} = \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$


Add additional “homogeneous” element

# Homogeneous Matrices

- To make a matrix work with homogeneous coordinates, add an extra row and column
- To do nothing else, make these  $[0 \ 0 \ 1]$

$$\begin{bmatrix} m_{11} & m_{12} \\ m_{21} & m_{22} \end{bmatrix}$$



$$\begin{bmatrix} m_{11} & m_{12} & 0 \\ m_{21} & m_{22} & 0 \\ 0 & 0 & 1 \end{bmatrix}$$

# Translation

$$\begin{bmatrix} 1 & 0 & t_x \\ 0 & 1 & t_y \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix} = \begin{bmatrix} x + t_x \\ y + t_y \\ 1 \end{bmatrix}$$

# Points vs. Vectors

$$\mathbf{p} = \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

allows translation

$$\mathbf{v} = \begin{bmatrix} x \\ y \\ 0 \end{bmatrix}$$

ignores translation

Remember: points have a position and *can* be translated,  
but vectors are only a displacement and *cannot* be translated

# Classes of Transformations

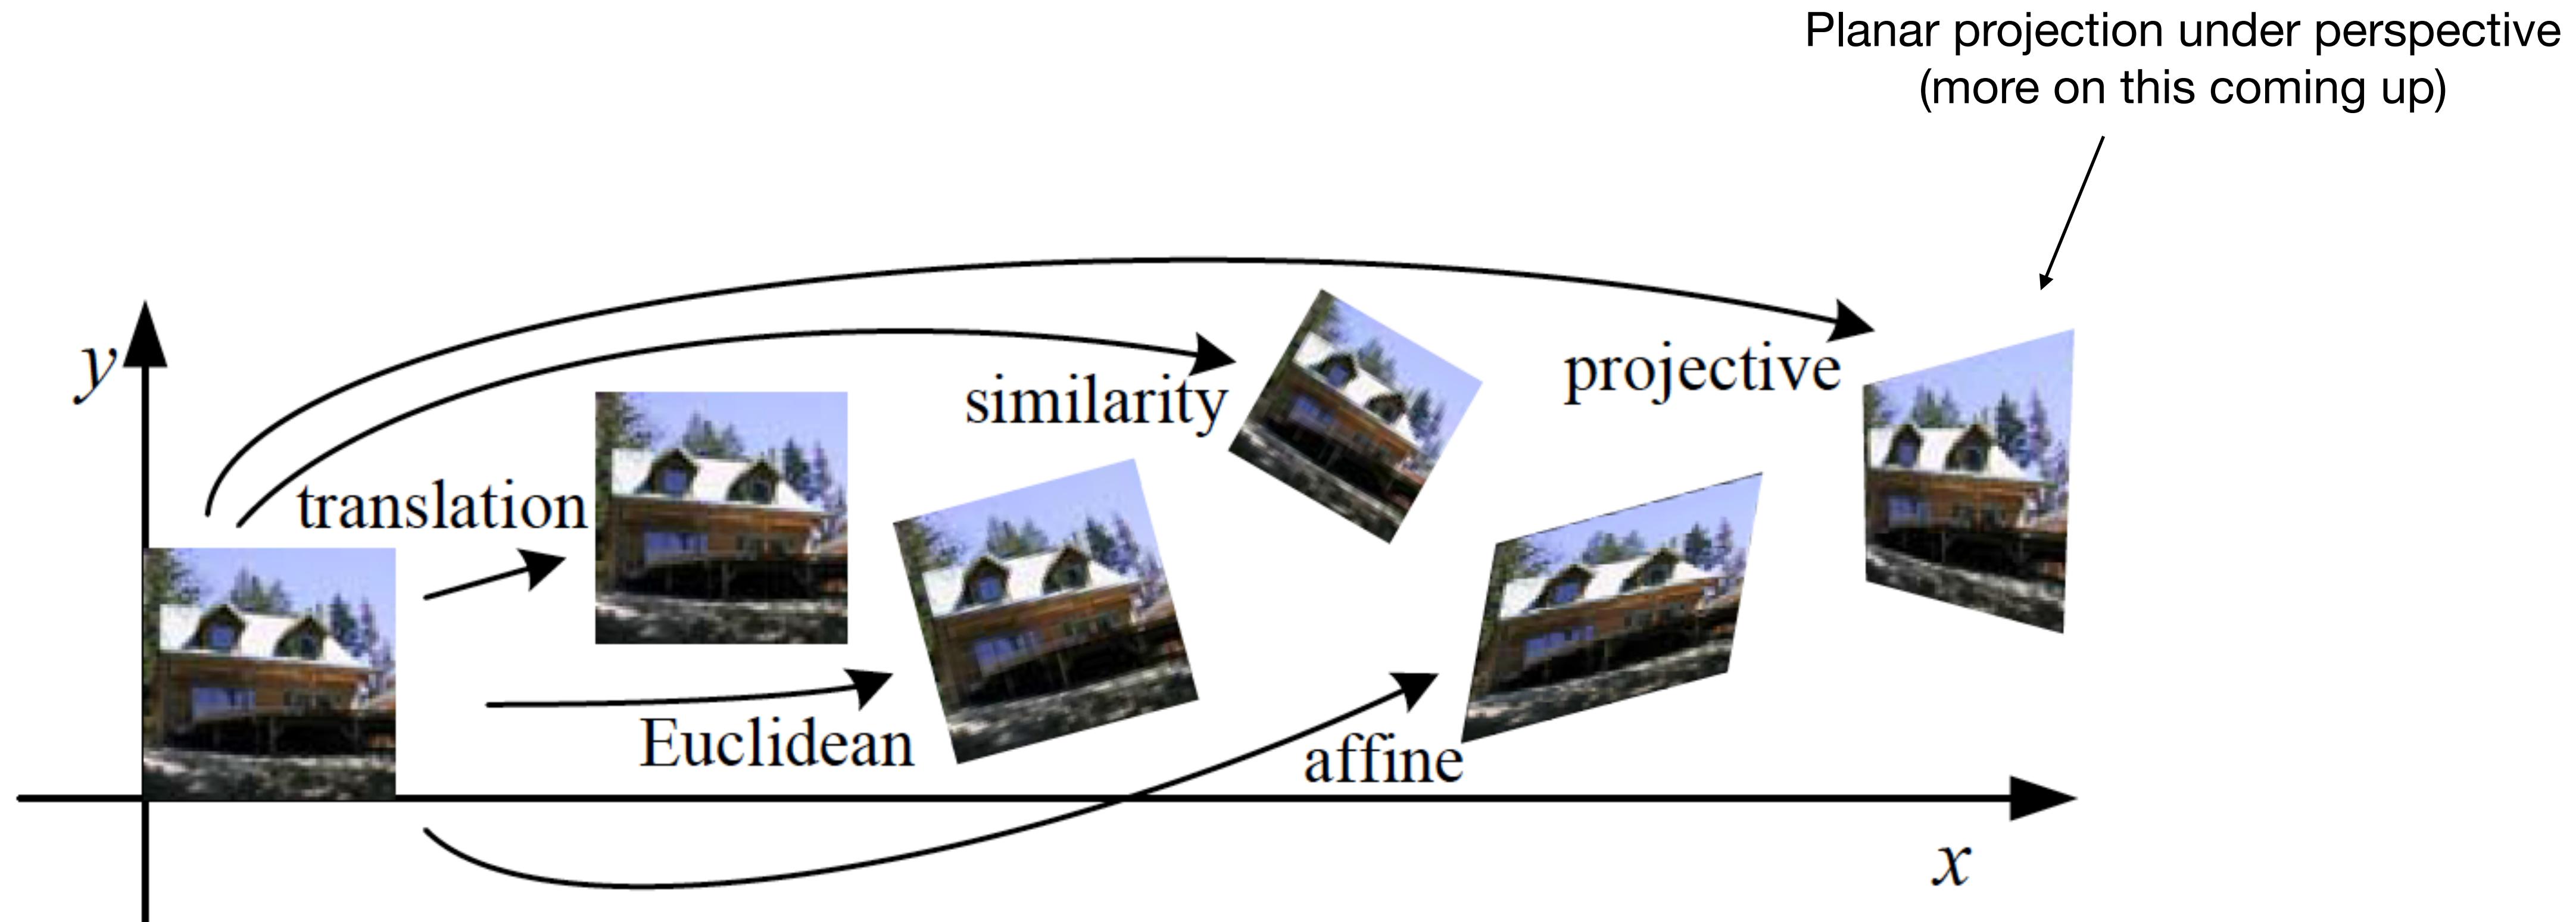
- Euclidean / Rigid body (preserves lengths):  
Rotation and translation
- Similarity (preserves angles):  
Rotation, translation, scale
- Affine (preserves straight lines):  
Arbitrary linear combinations

$$\begin{bmatrix} m_{11} & m_{12} & m_{13} \\ m_{21} & m_{22} & m_{23} \\ 0 & 0 & 1 \end{bmatrix}$$



General form of an  
affine transformation

# Classes of Transformations



# 3D Linear Transformations

- Scaling has the same form as in 2D
- Translation has the same form as in 2D
- Rotation has the same form as in 2D if you begin with unit vectors for the new coordinate axes
  - If using orthonormal coordinate systems, this is always an orthonormal matrix
  - Pro tip: any orthonormal matrix is doing a rotation or some form

$$\mathbf{S} = \begin{bmatrix} s_x & 0 & 0 & 0 \\ 0 & s_y & 0 & 0 \\ 0 & 0 & s_z & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix}$$

$$\mathbf{T} = \begin{bmatrix} 1 & 0 & 0 & t_x \\ 0 & 1 & 0 & t_y \\ 0 & 0 & 1 & t_z \\ 0 & 0 & 0 & 1 \end{bmatrix}$$

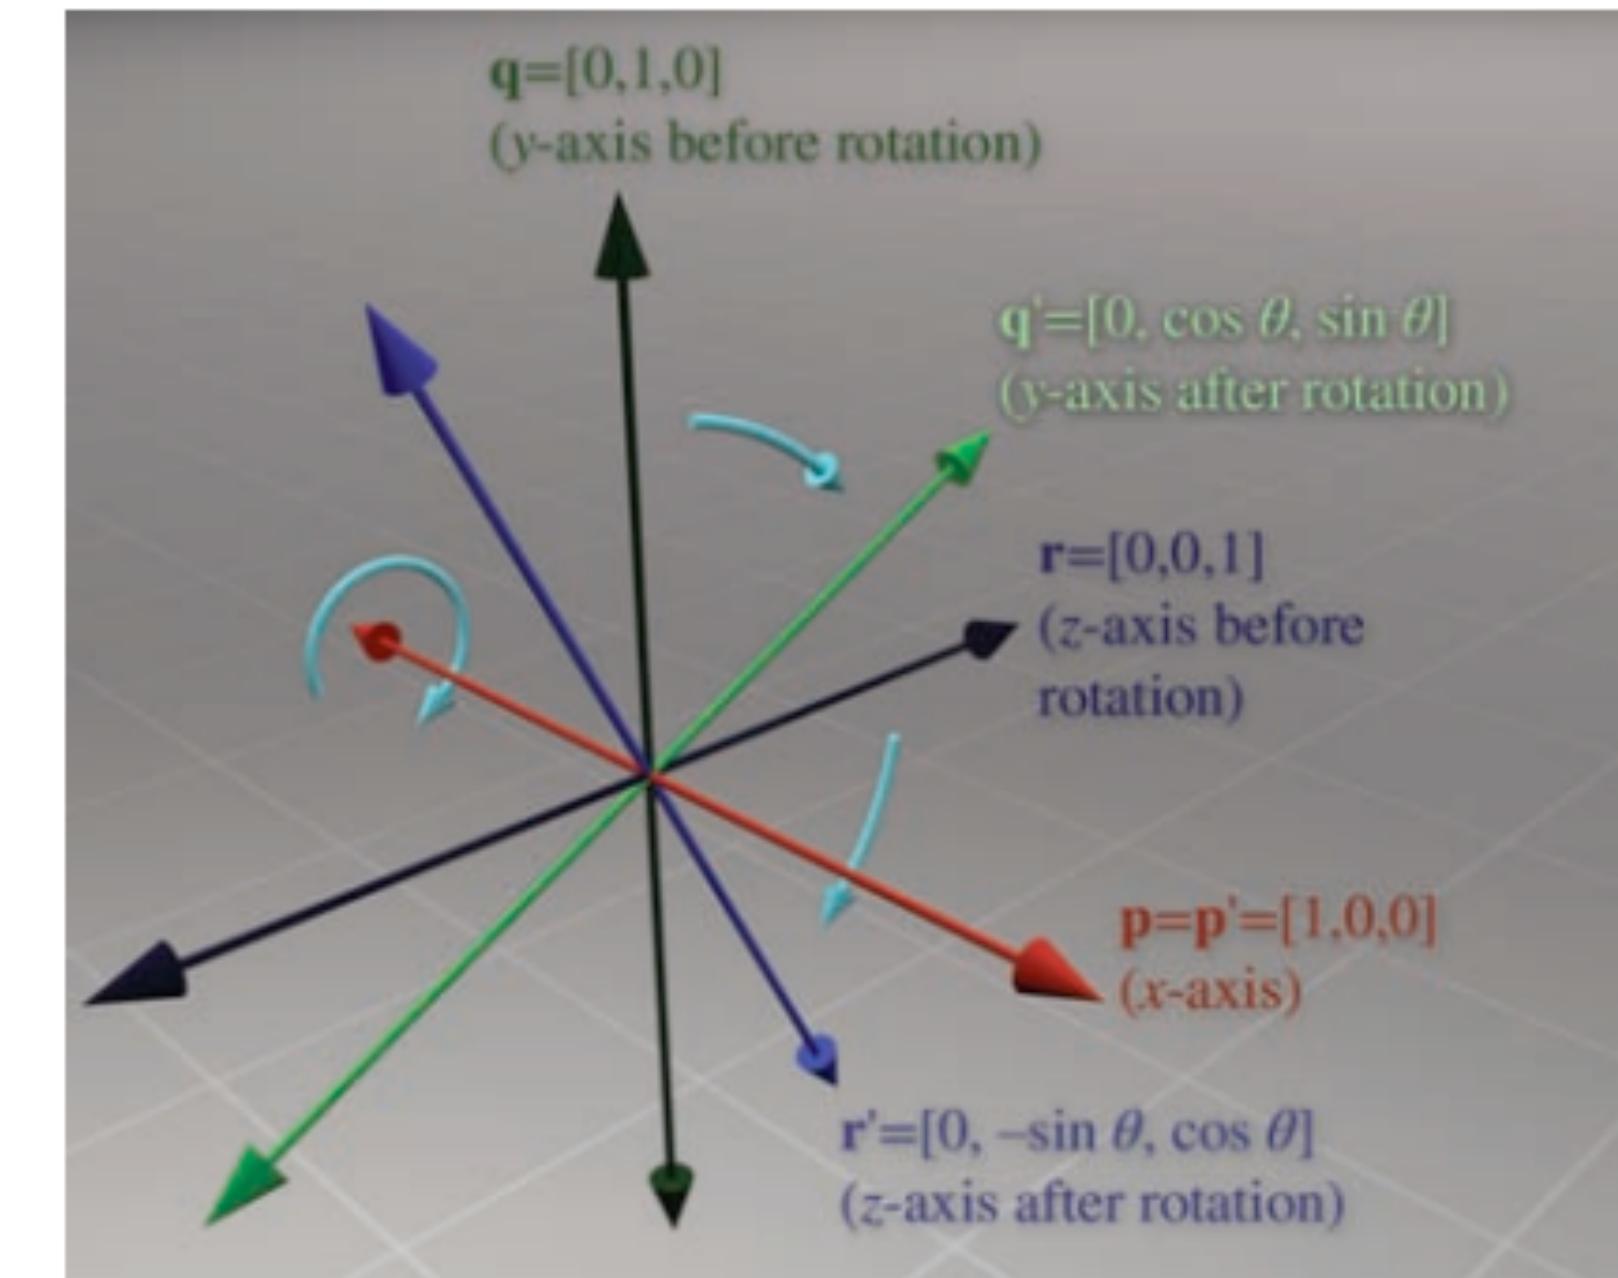
$$\mathbf{R} = \begin{bmatrix} e_{x1} & e_{x2} & e_{x3} & 0 \\ e_{y1} & e_{y2} & e_{y3} & 0 \\ e_{z1} & e_{z2} & e_{z3} & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix}$$

# Rotation in 3D

$$\mathbf{R}_x(\theta) = \begin{bmatrix} 1 & 0 & 0 \\ 0 & \cos \theta & -\sin \theta \\ 0 & \sin \theta & \cos \theta \end{bmatrix}$$

$$\mathbf{R}_y(\theta) = \begin{bmatrix} \cos \theta & 0 & \sin \theta \\ 0 & 1 & 0 \\ -\sin \theta & 0 & \cos \theta \end{bmatrix}$$

$$\mathbf{R}_z(\theta) = \begin{bmatrix} \cos \theta & -\sin \theta & 0 \\ \sin \theta & \cos \theta & 0 \\ 0 & 0 & 1 \end{bmatrix}$$



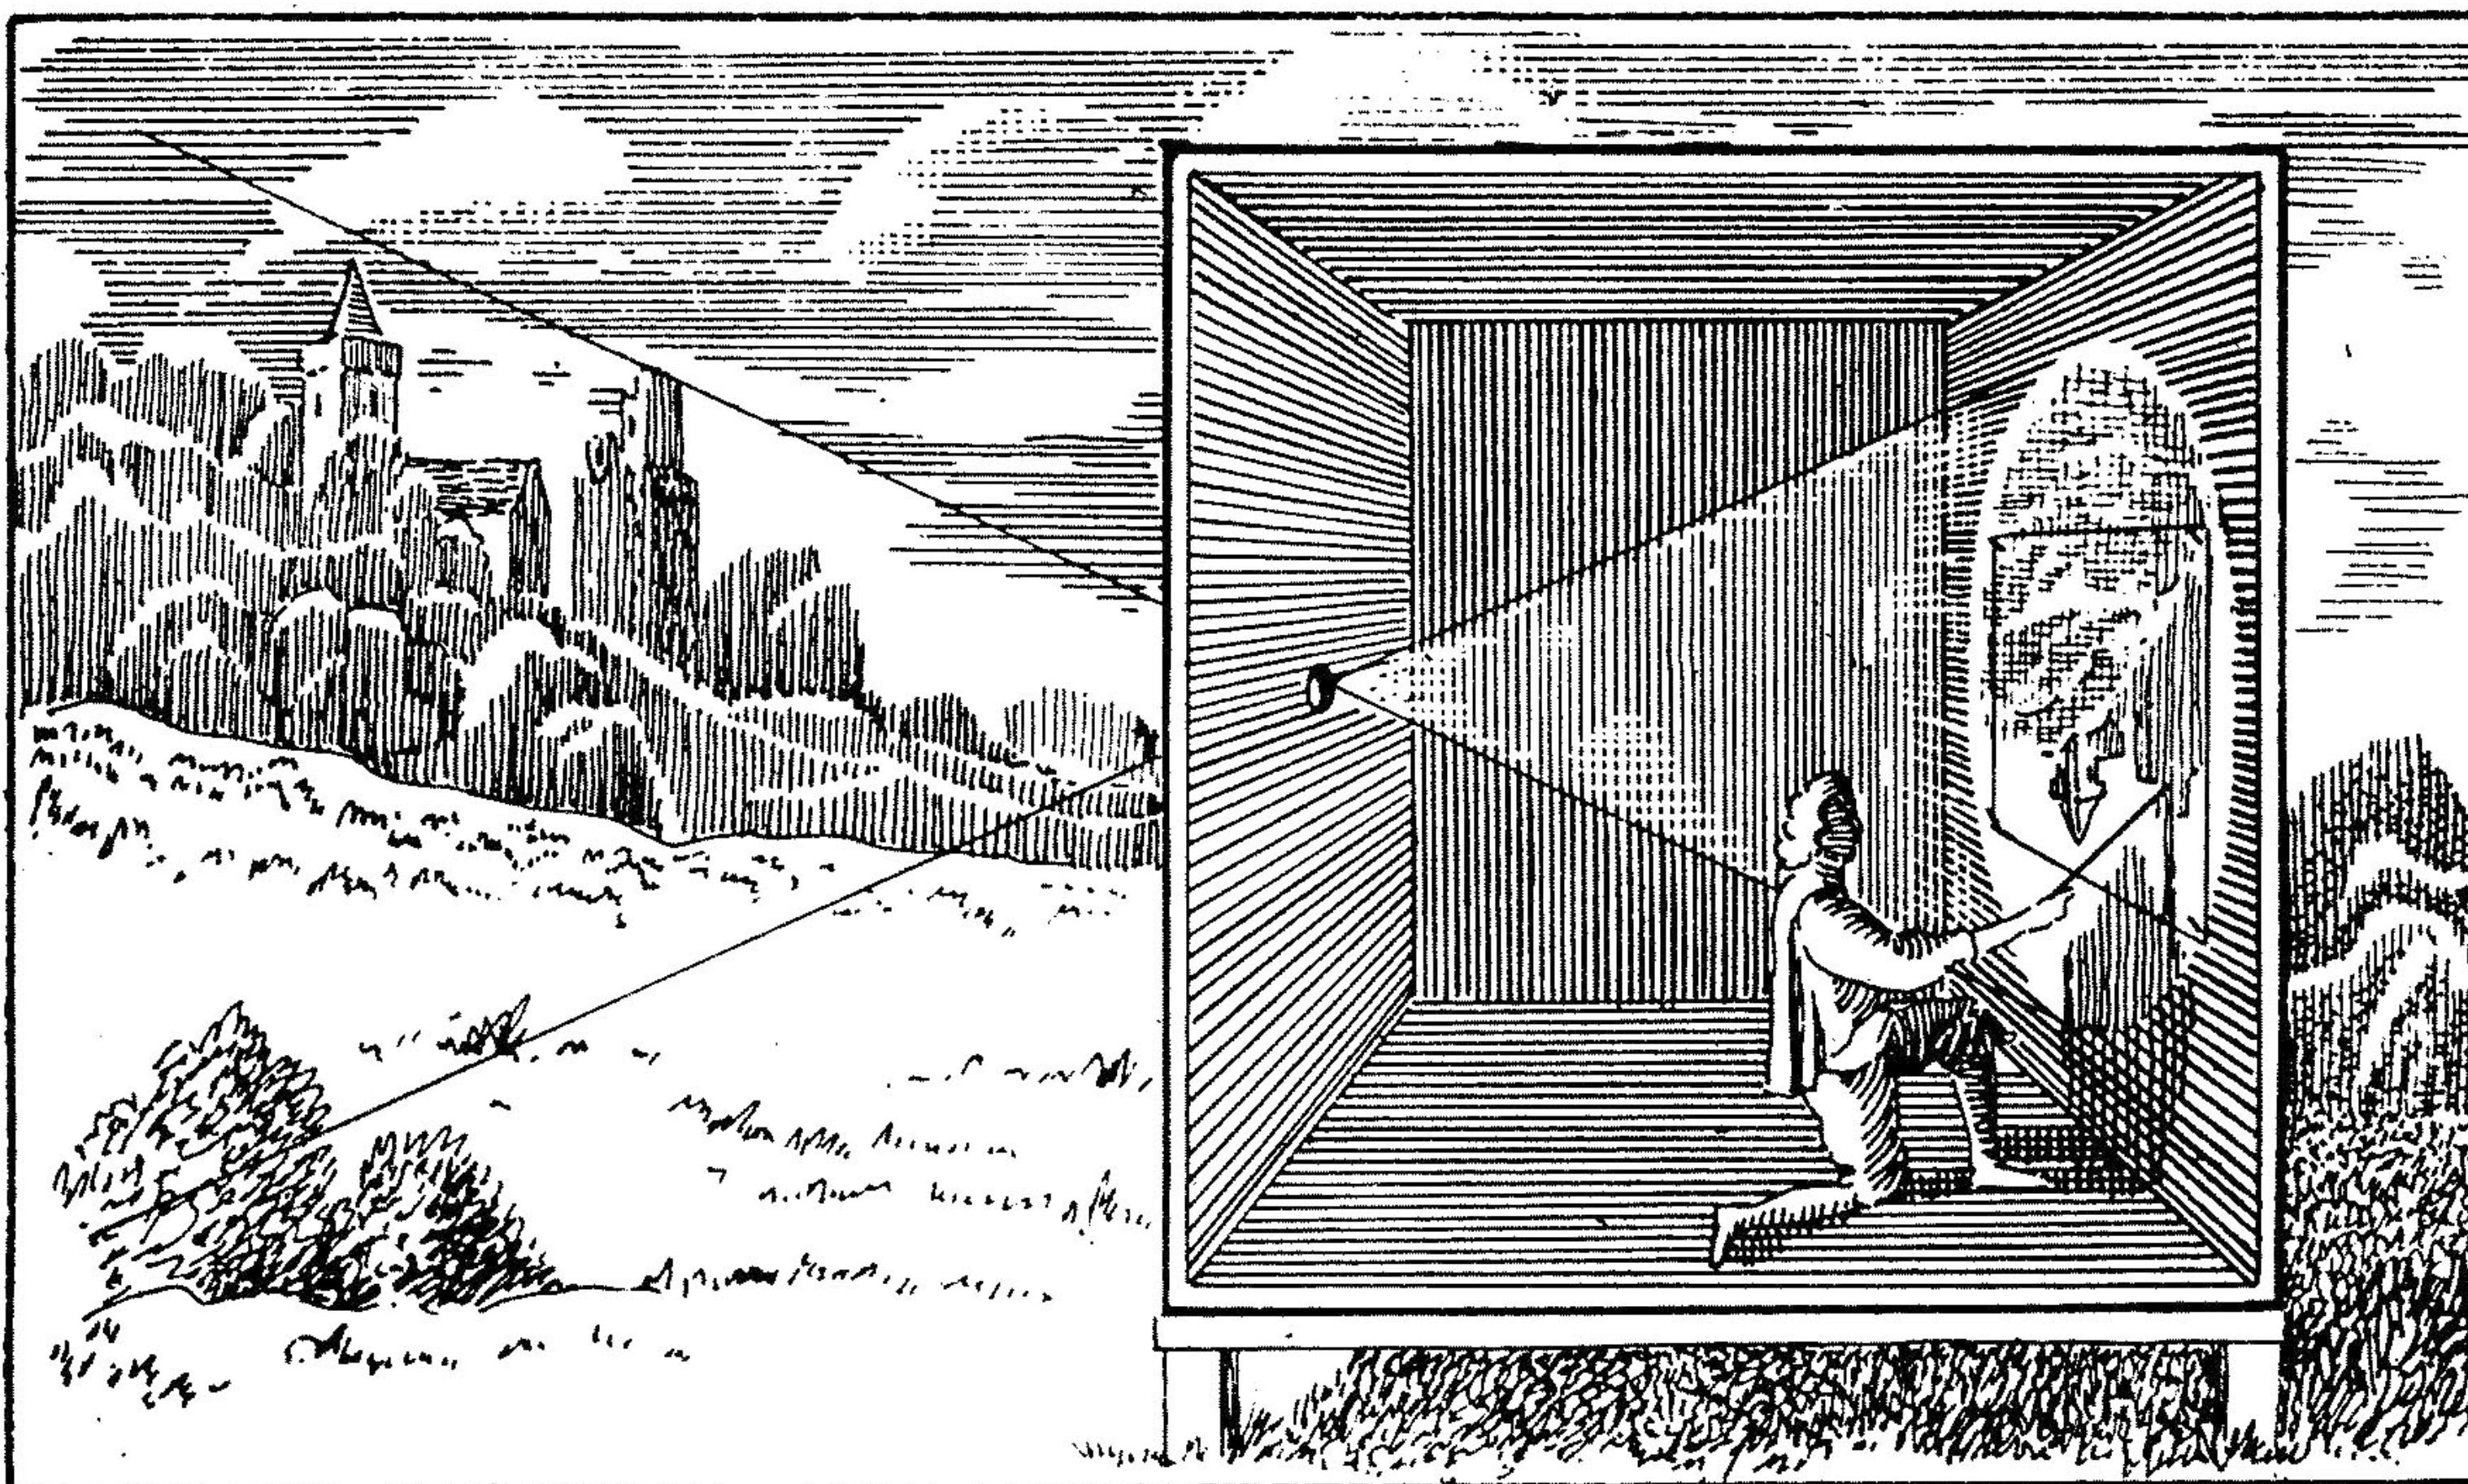
# Perspective Projection

# Perspective Projection

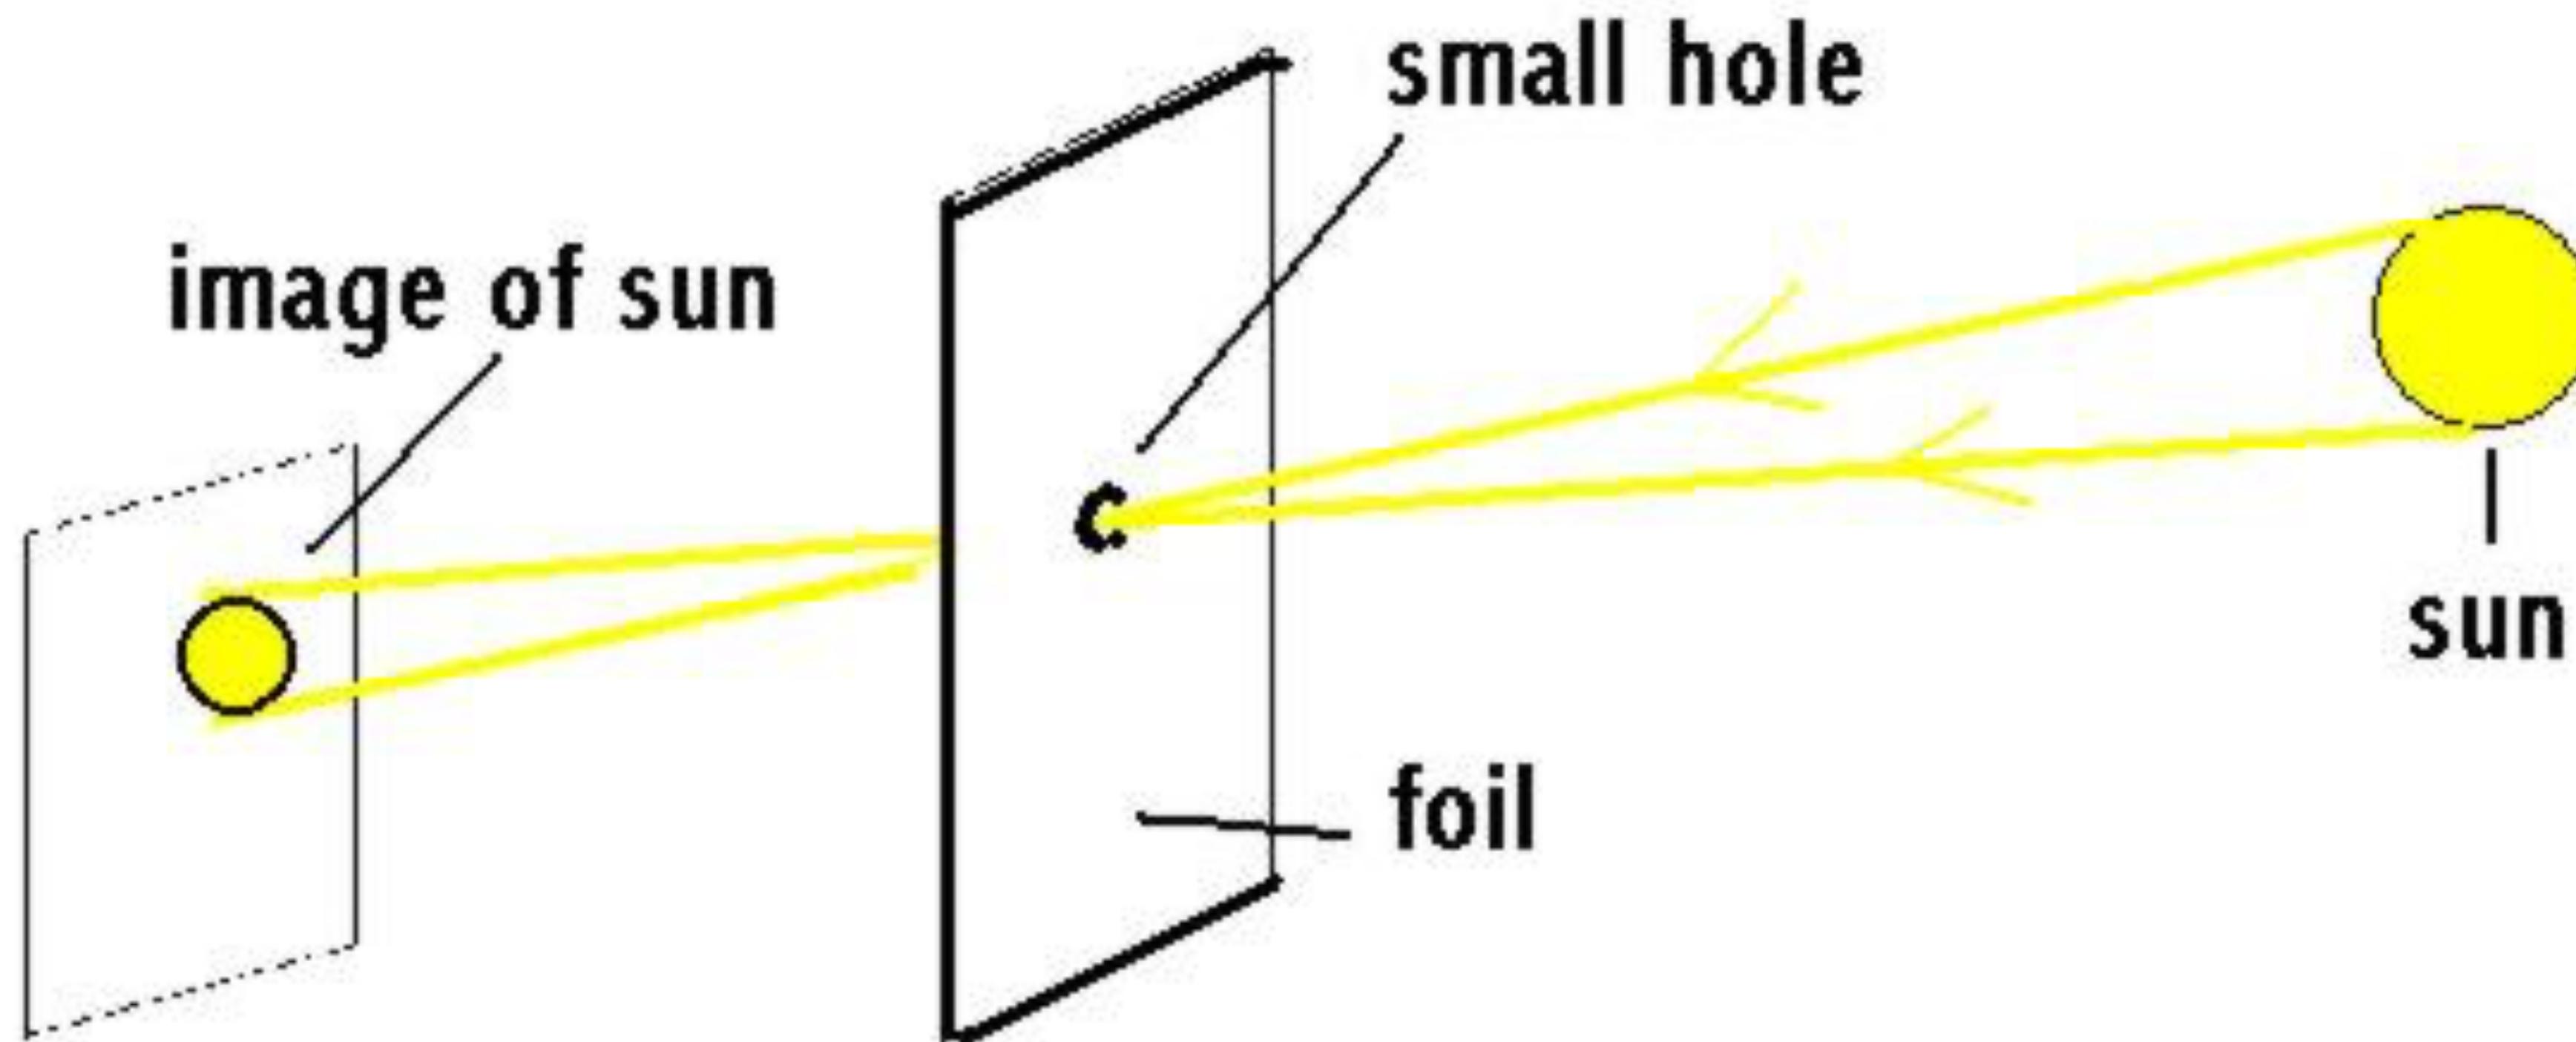
- What we're used to seeing
- Things farther away appear smaller  
(size inversely proportional to distance)
- Straight lines stay straight
- Parallel lines intersect in *vanishing points*



# Camera Obscura



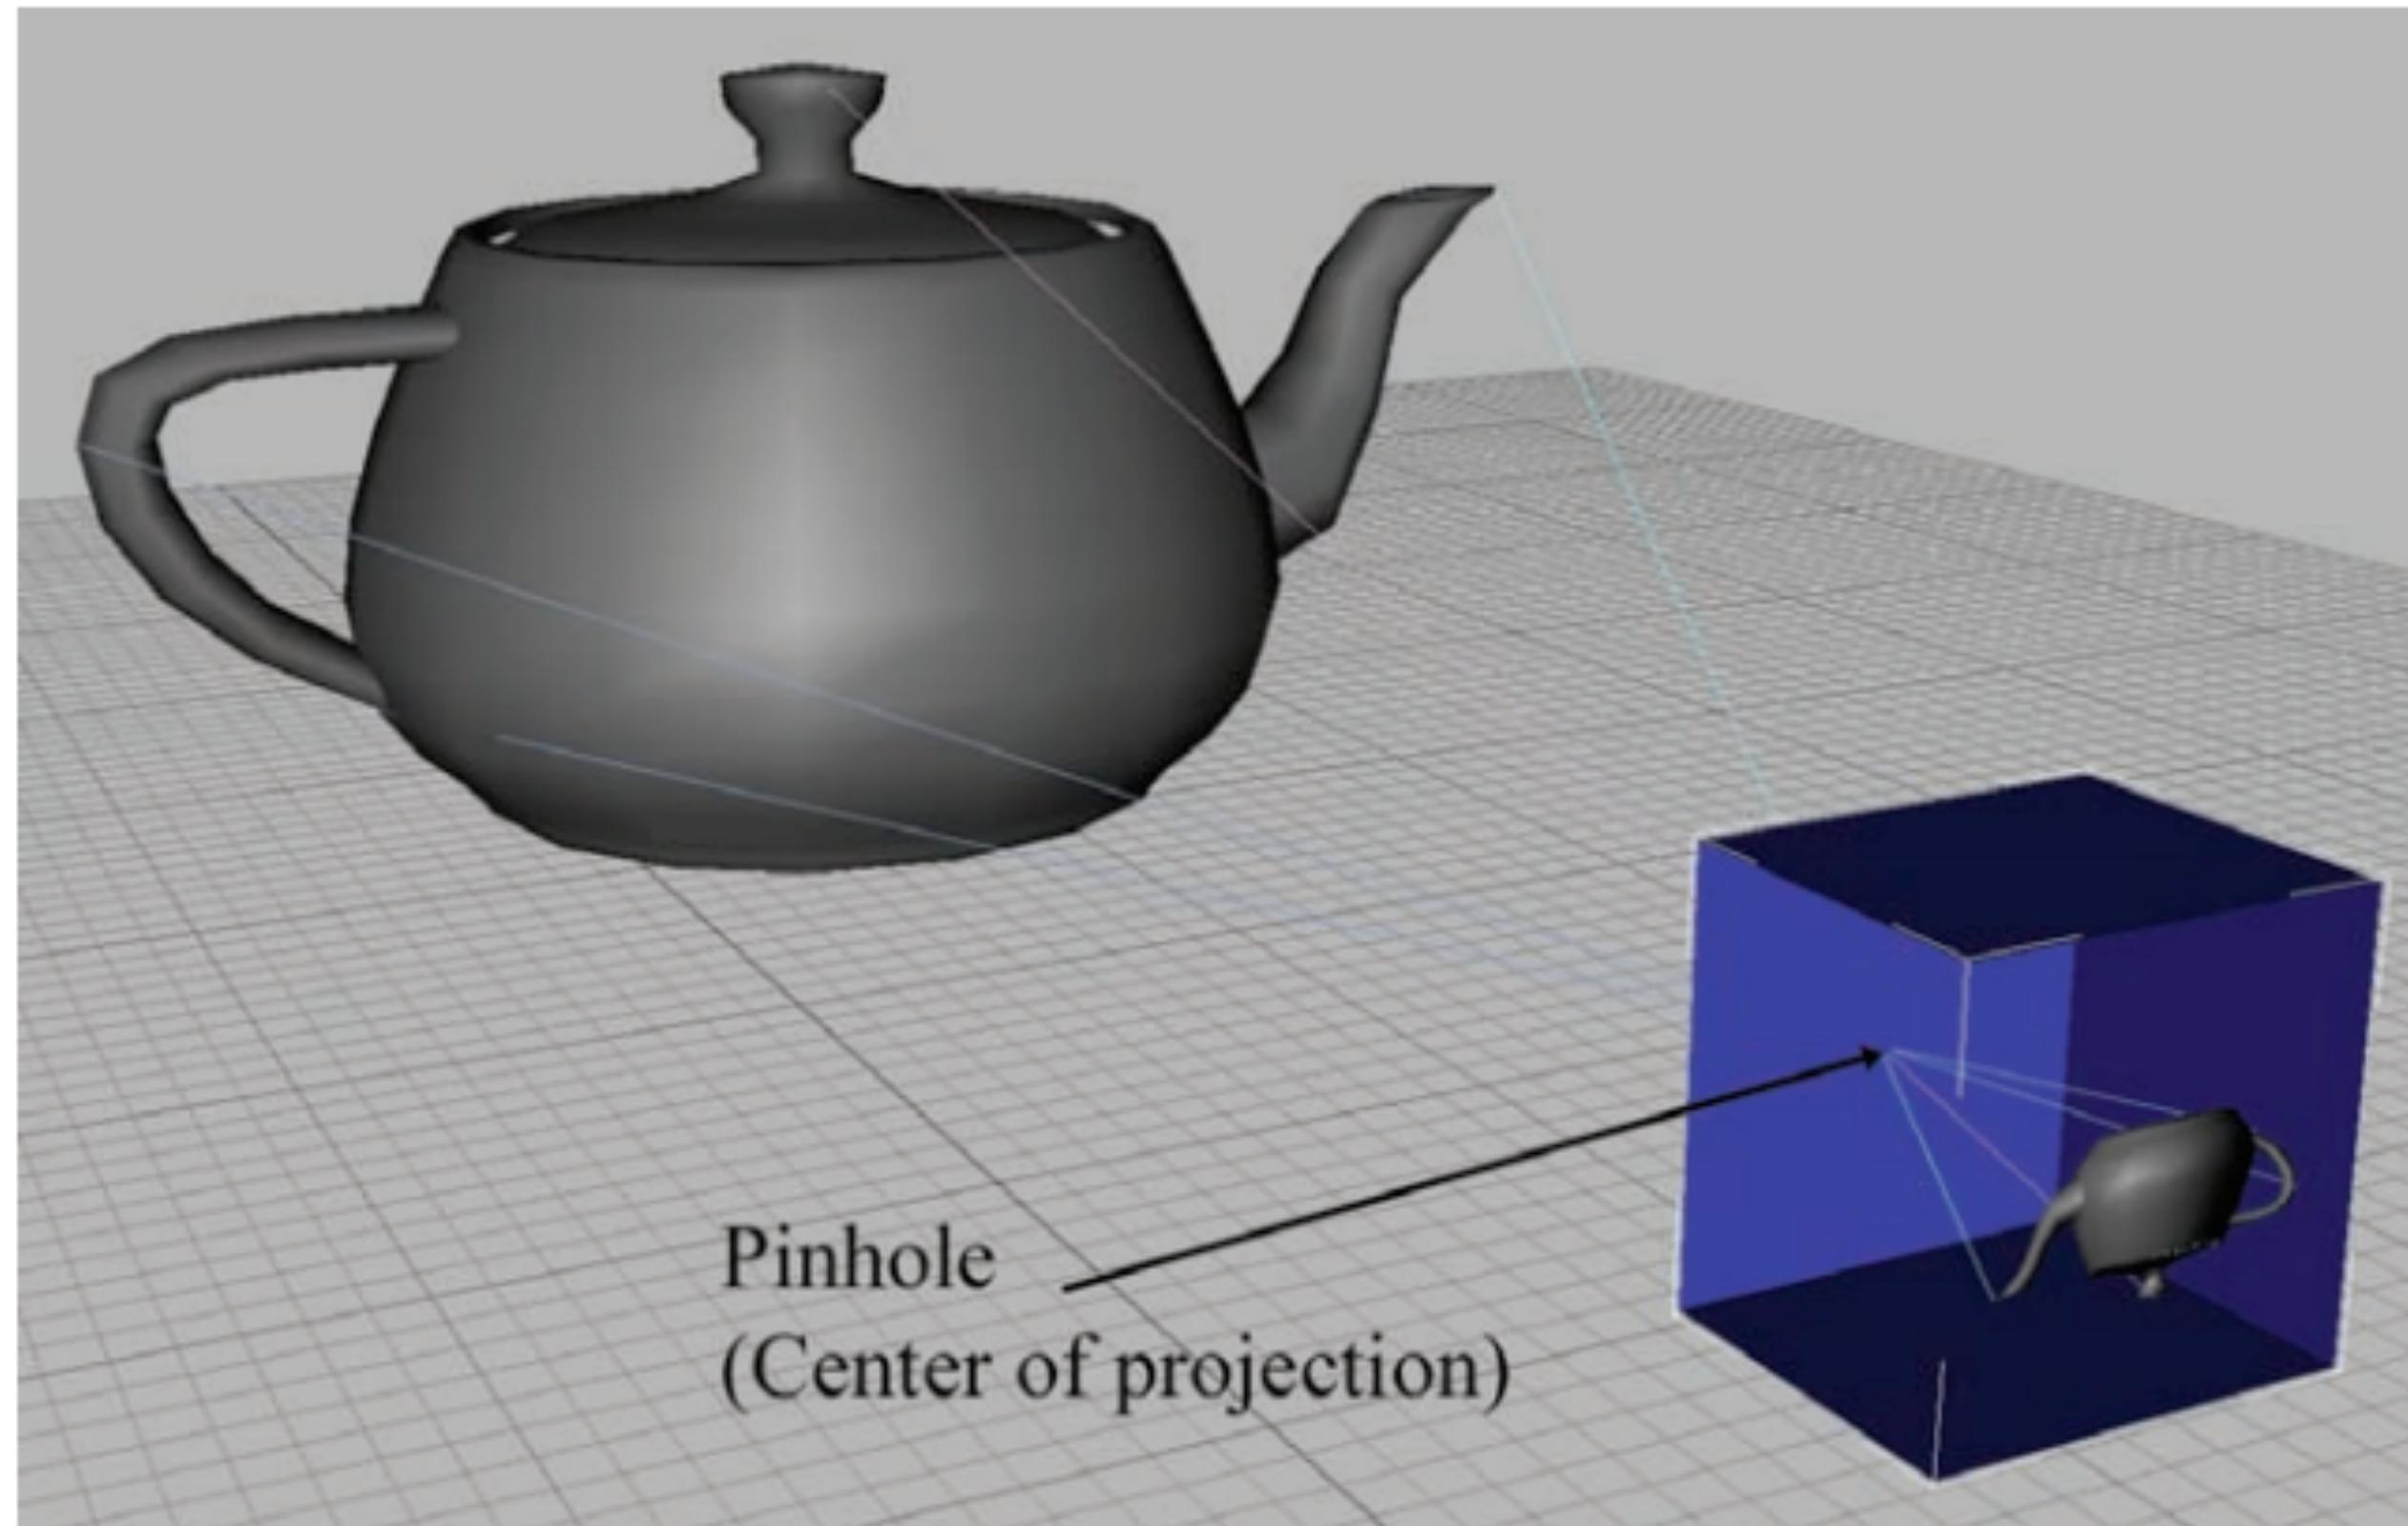
# Eclipse Viewing



# Eclipse Viewing

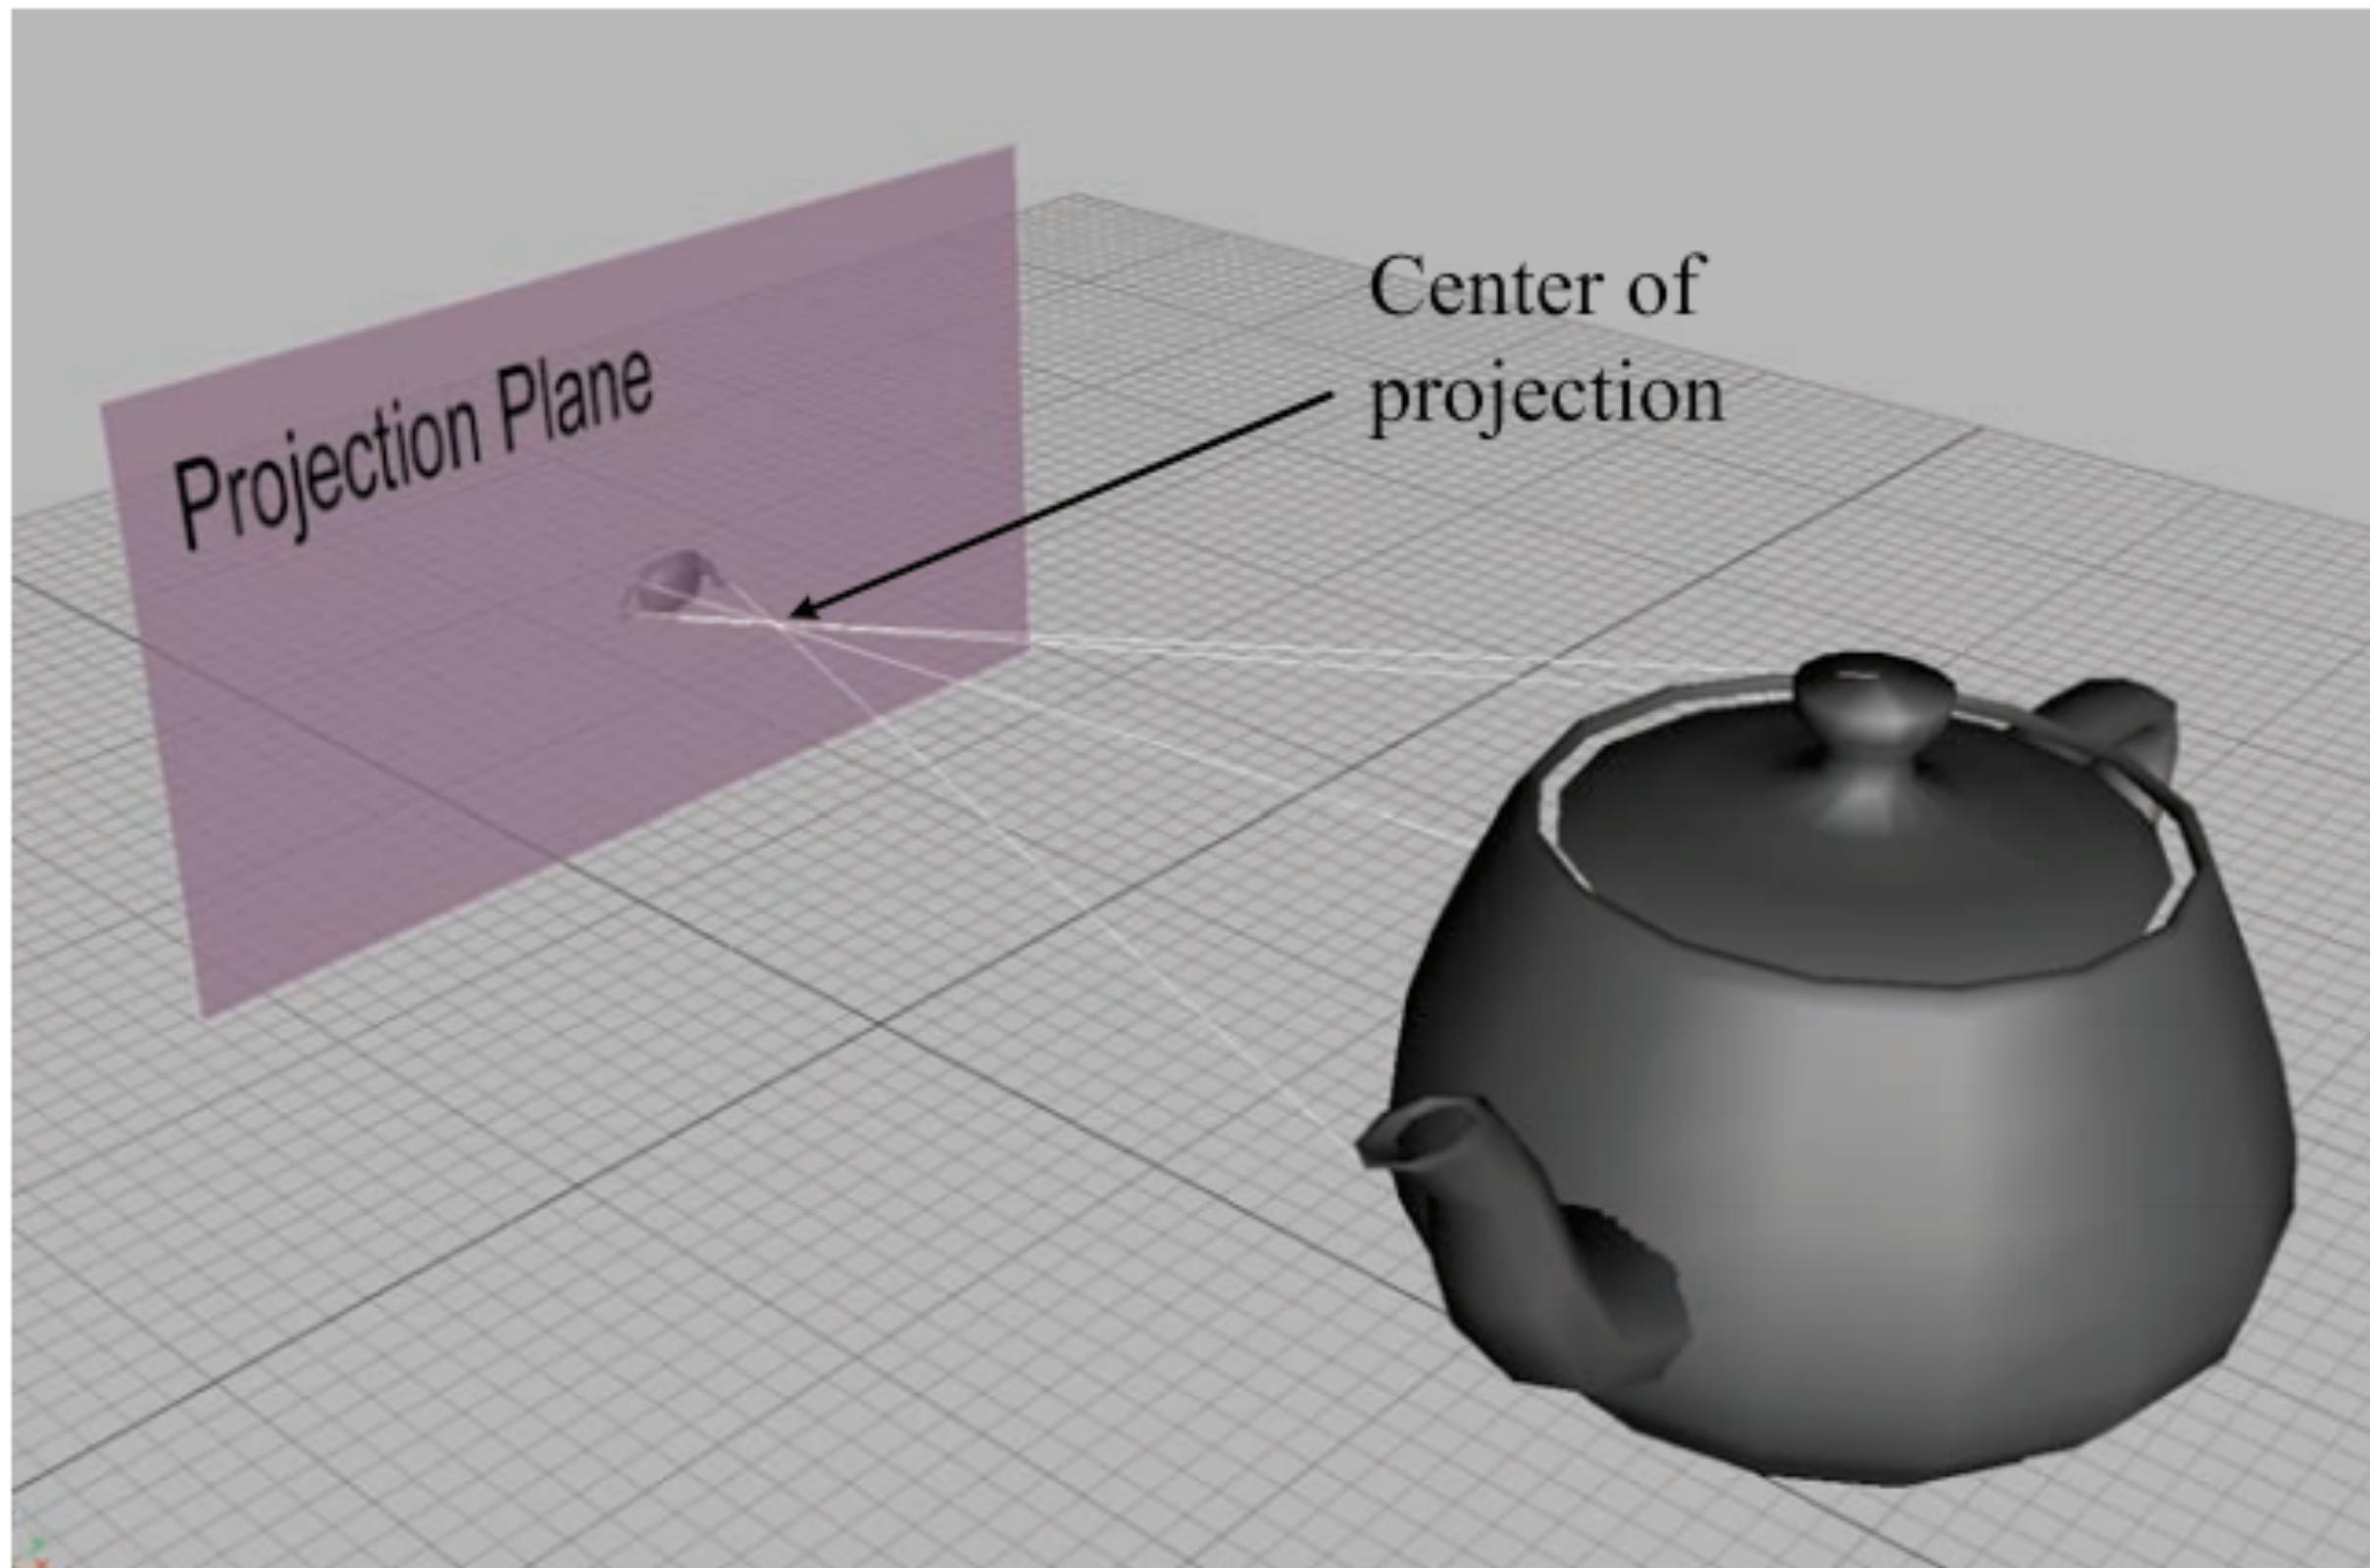


# Perspective Projection



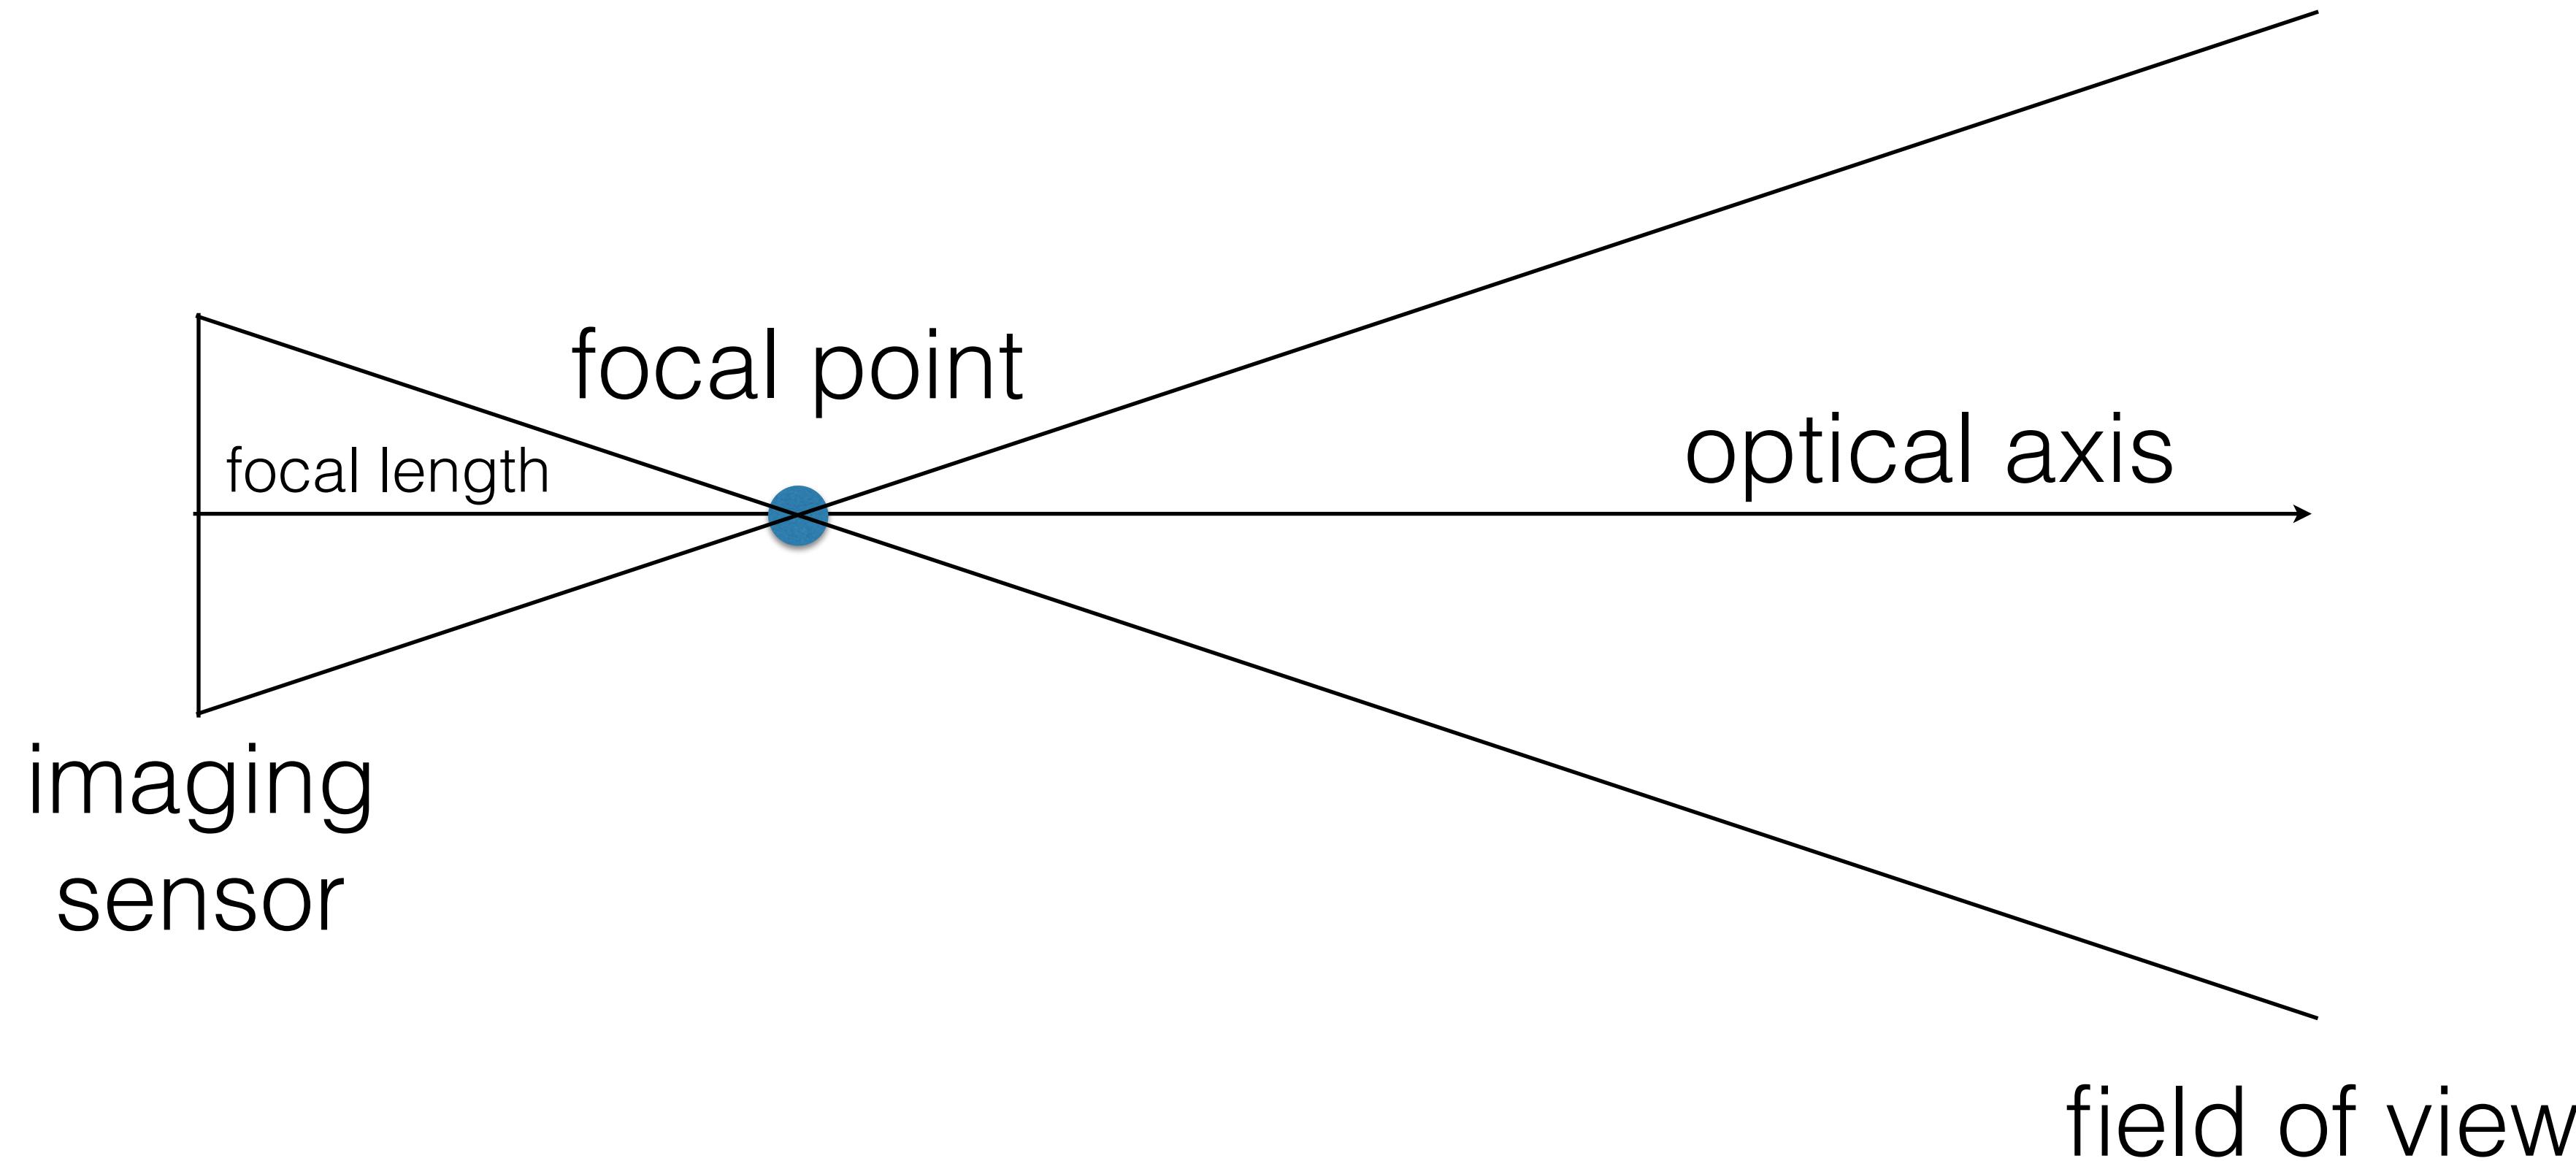
Many graphics/vision systems assume a simple pinhole camera model

# Perspective Projection

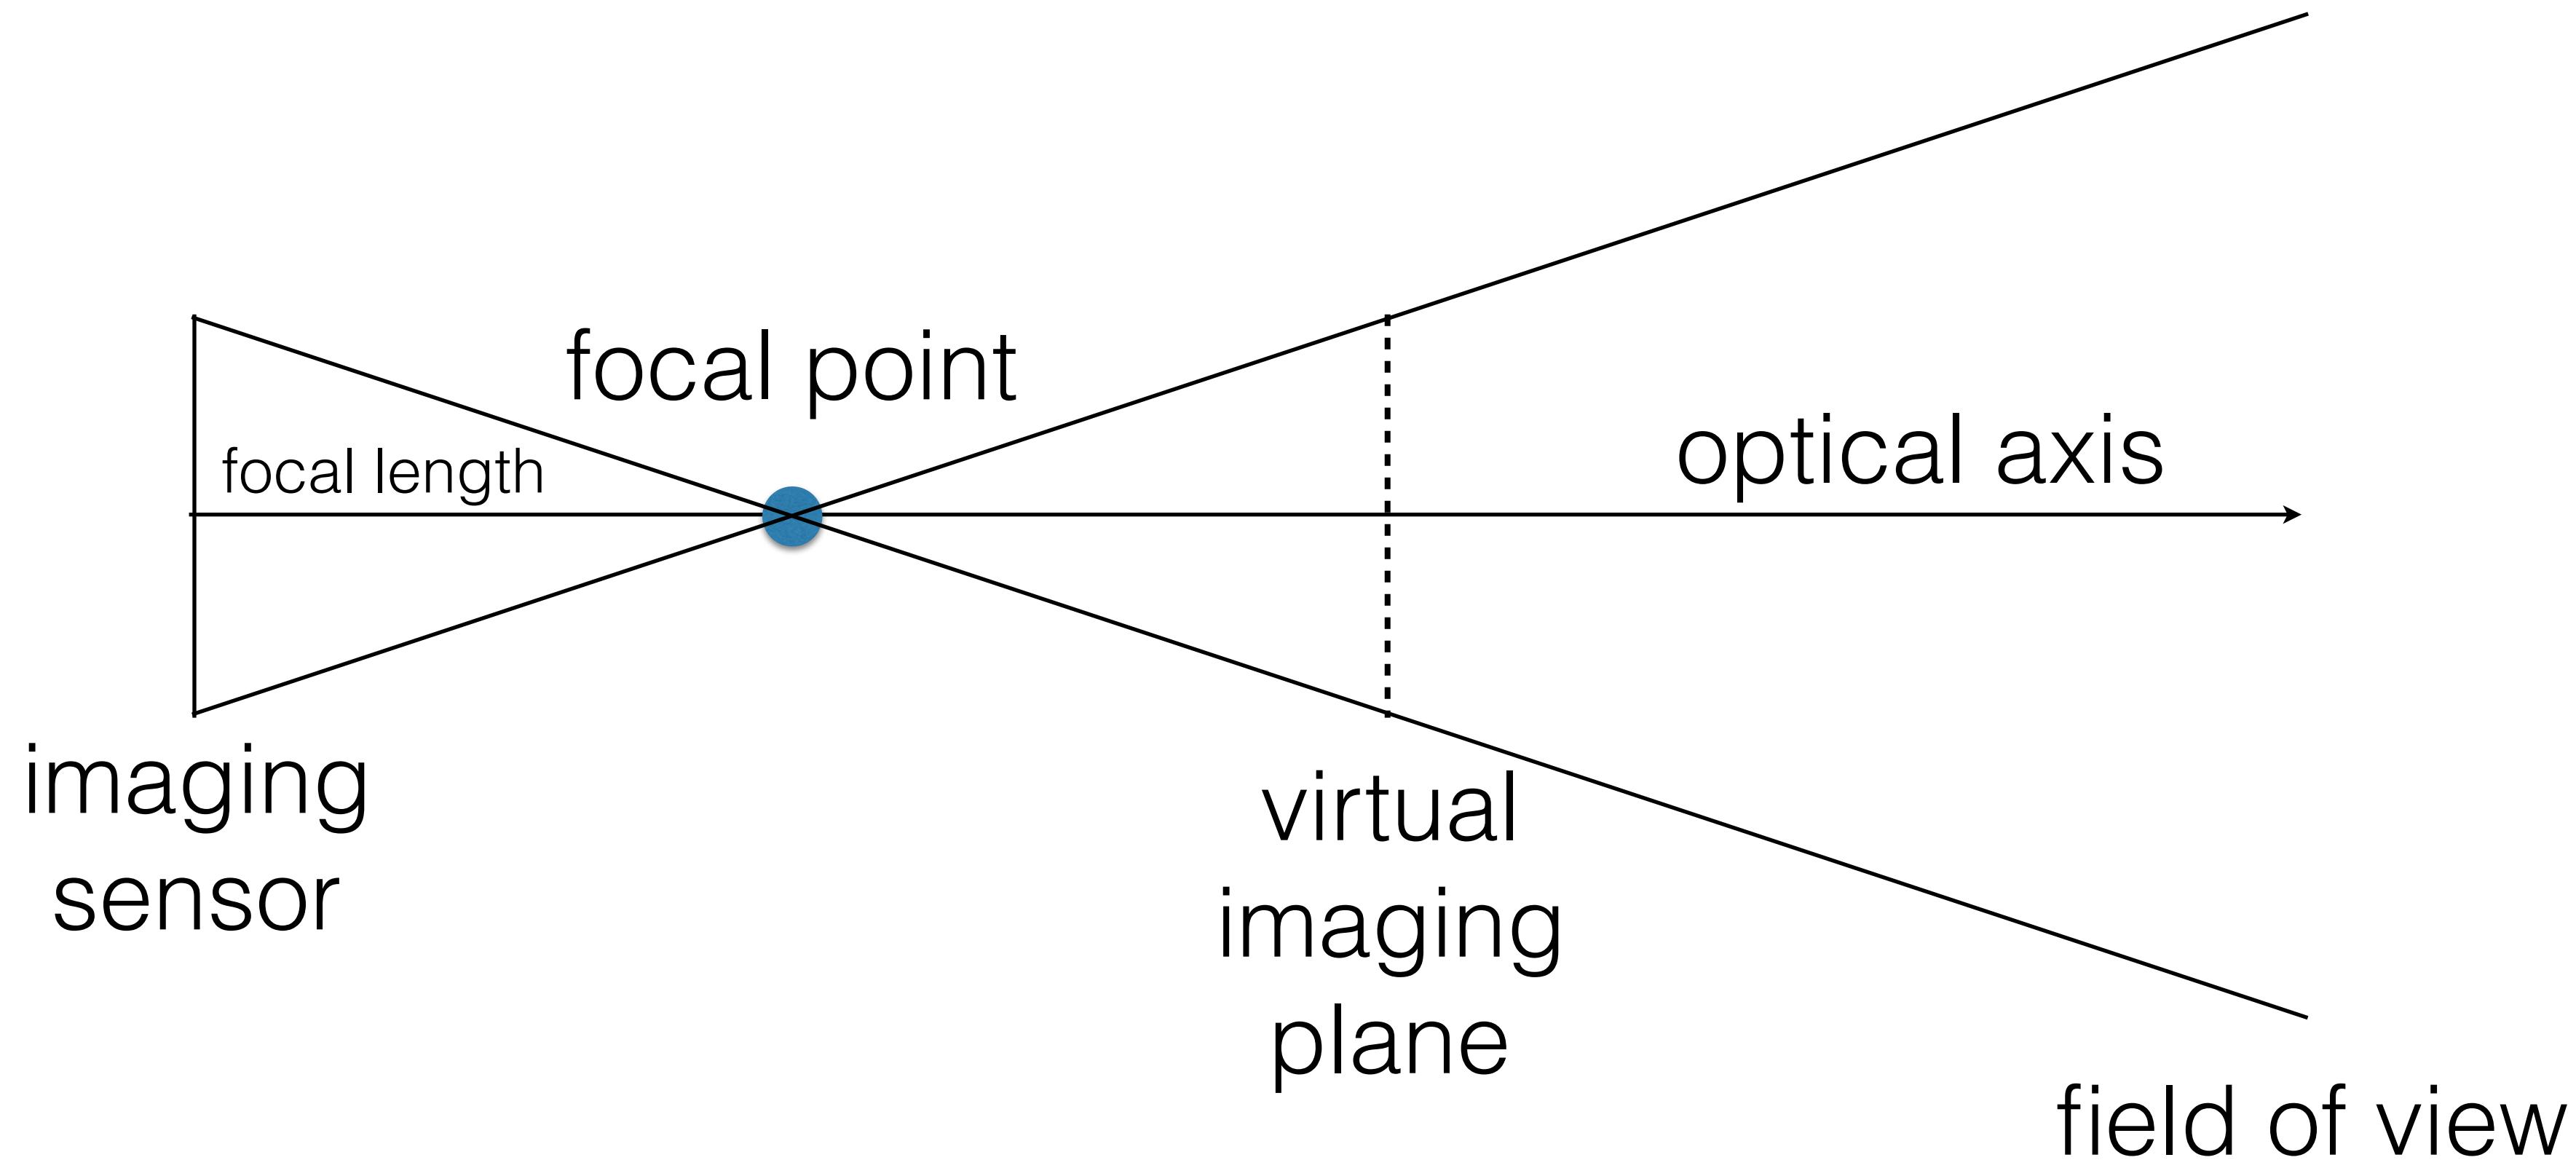


Pinhole camera model

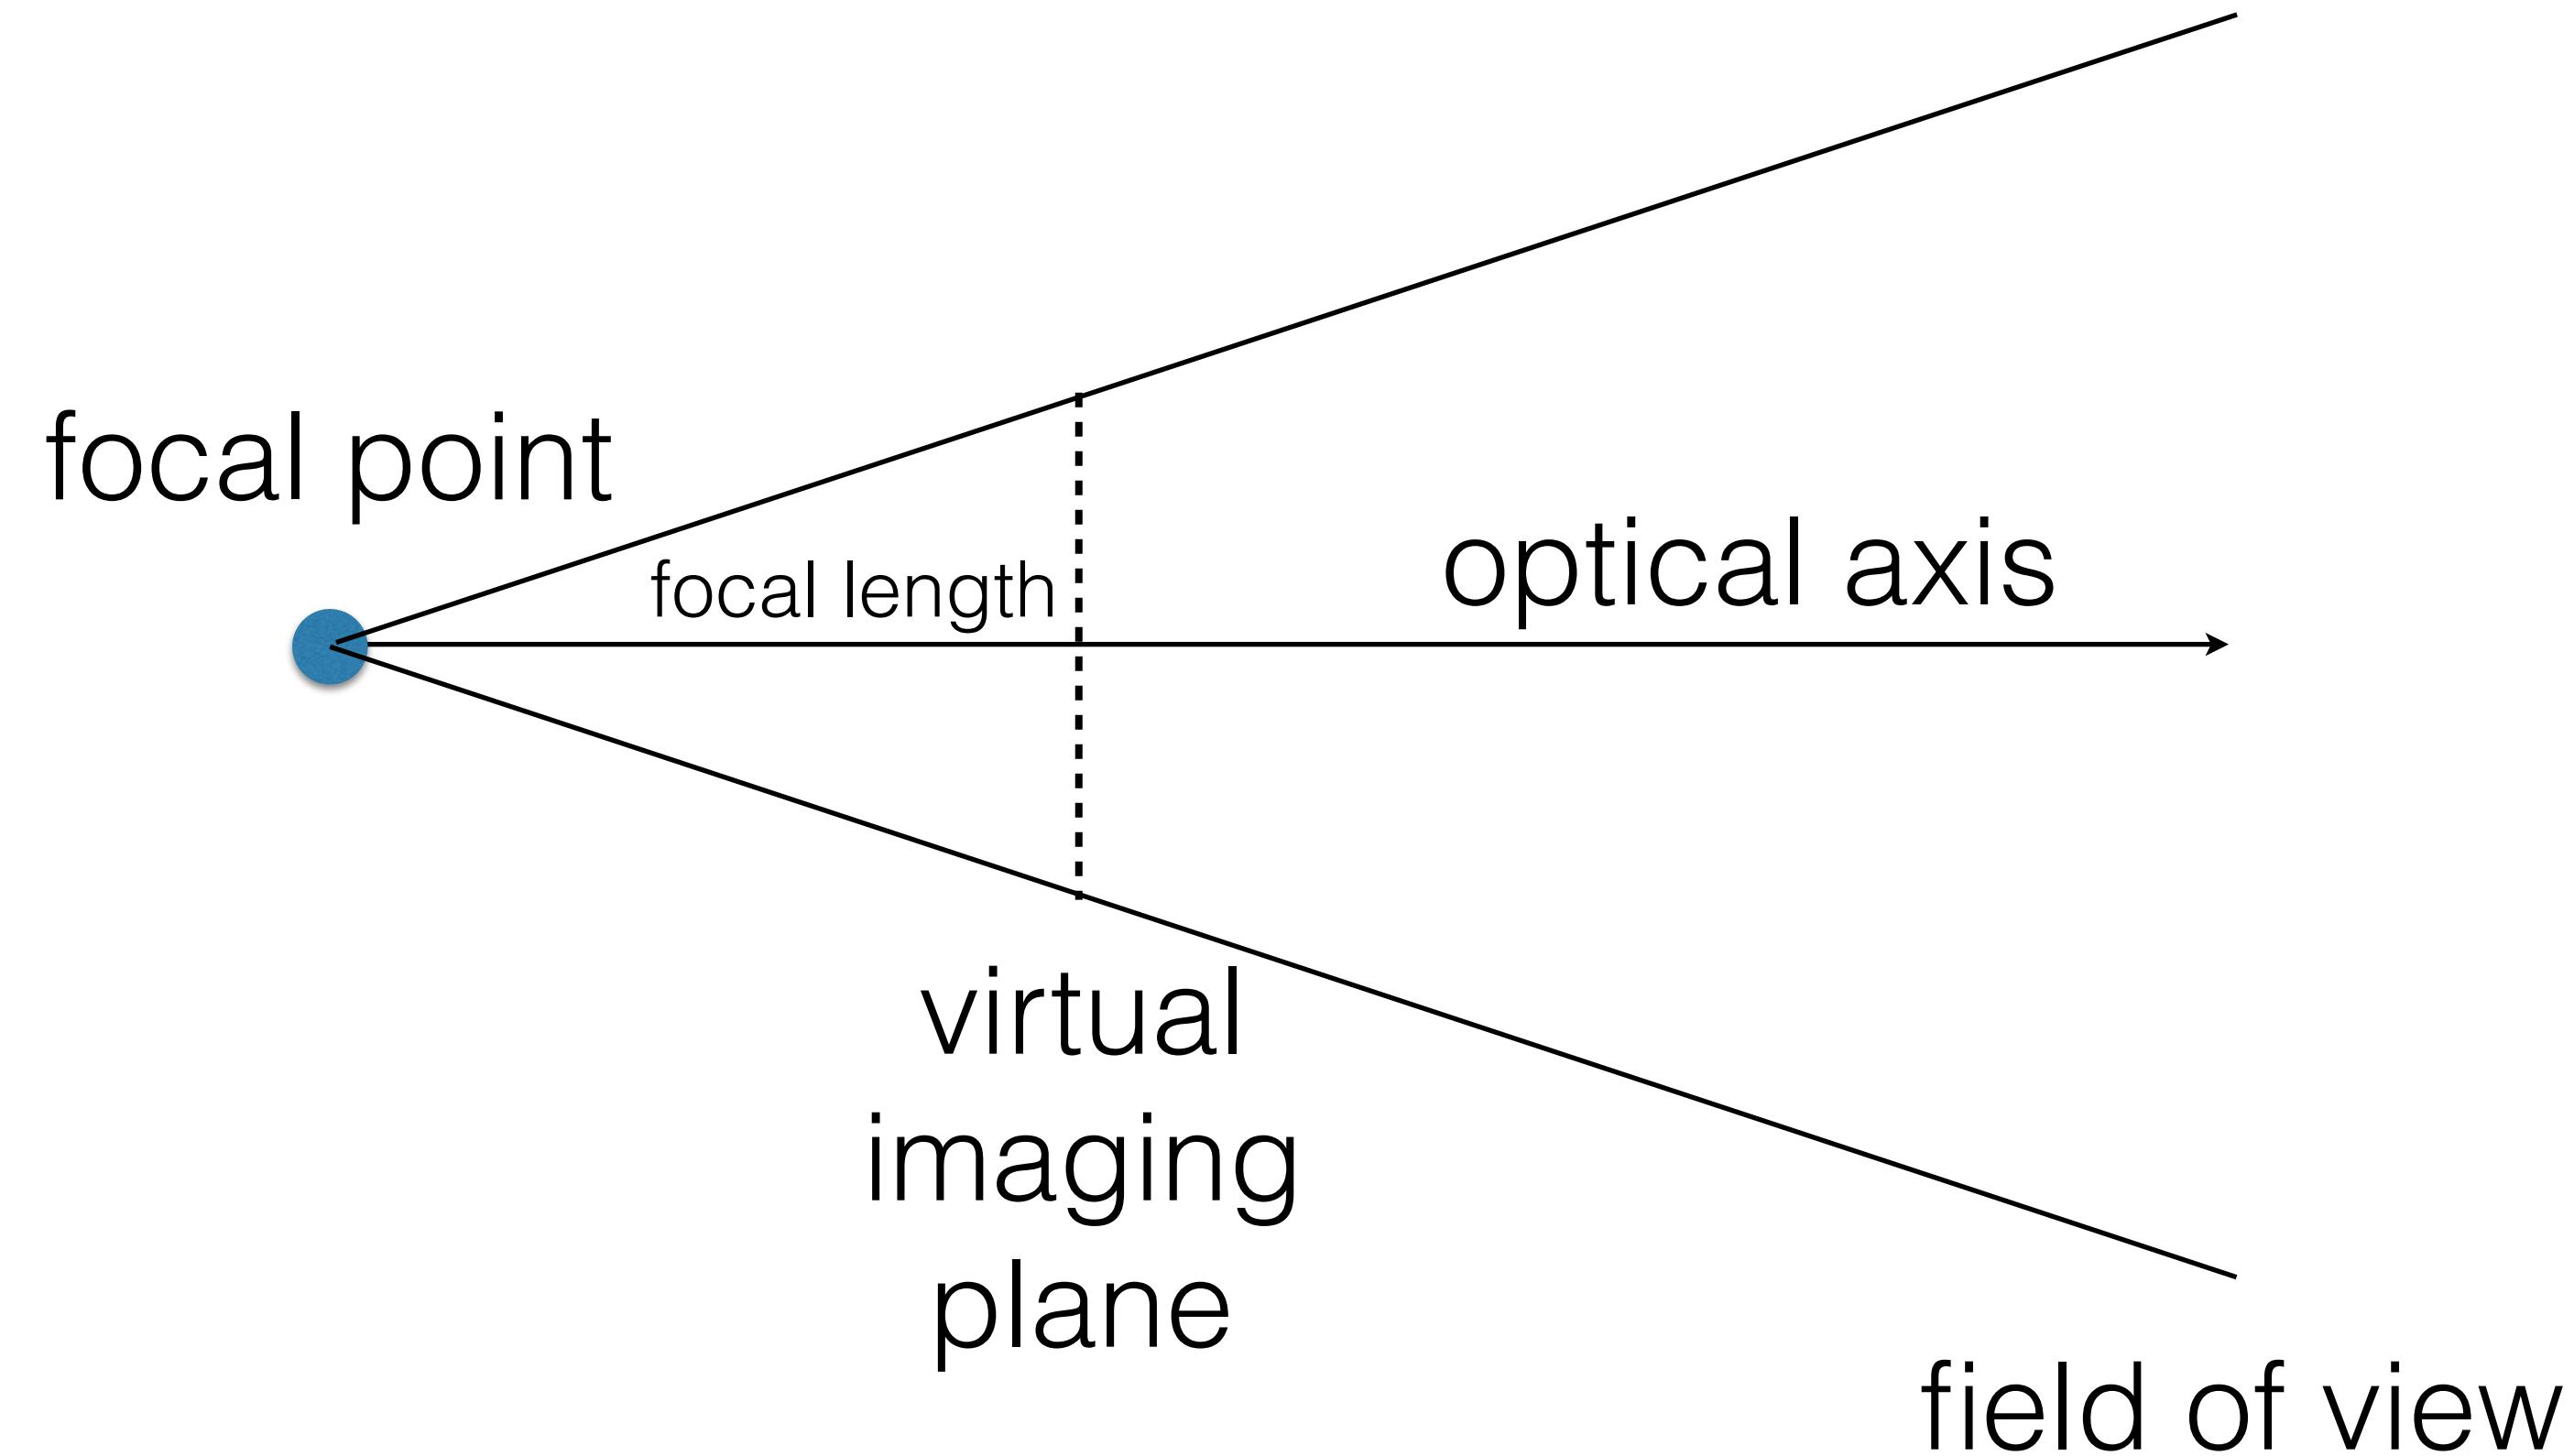
# Pinhole Cameras



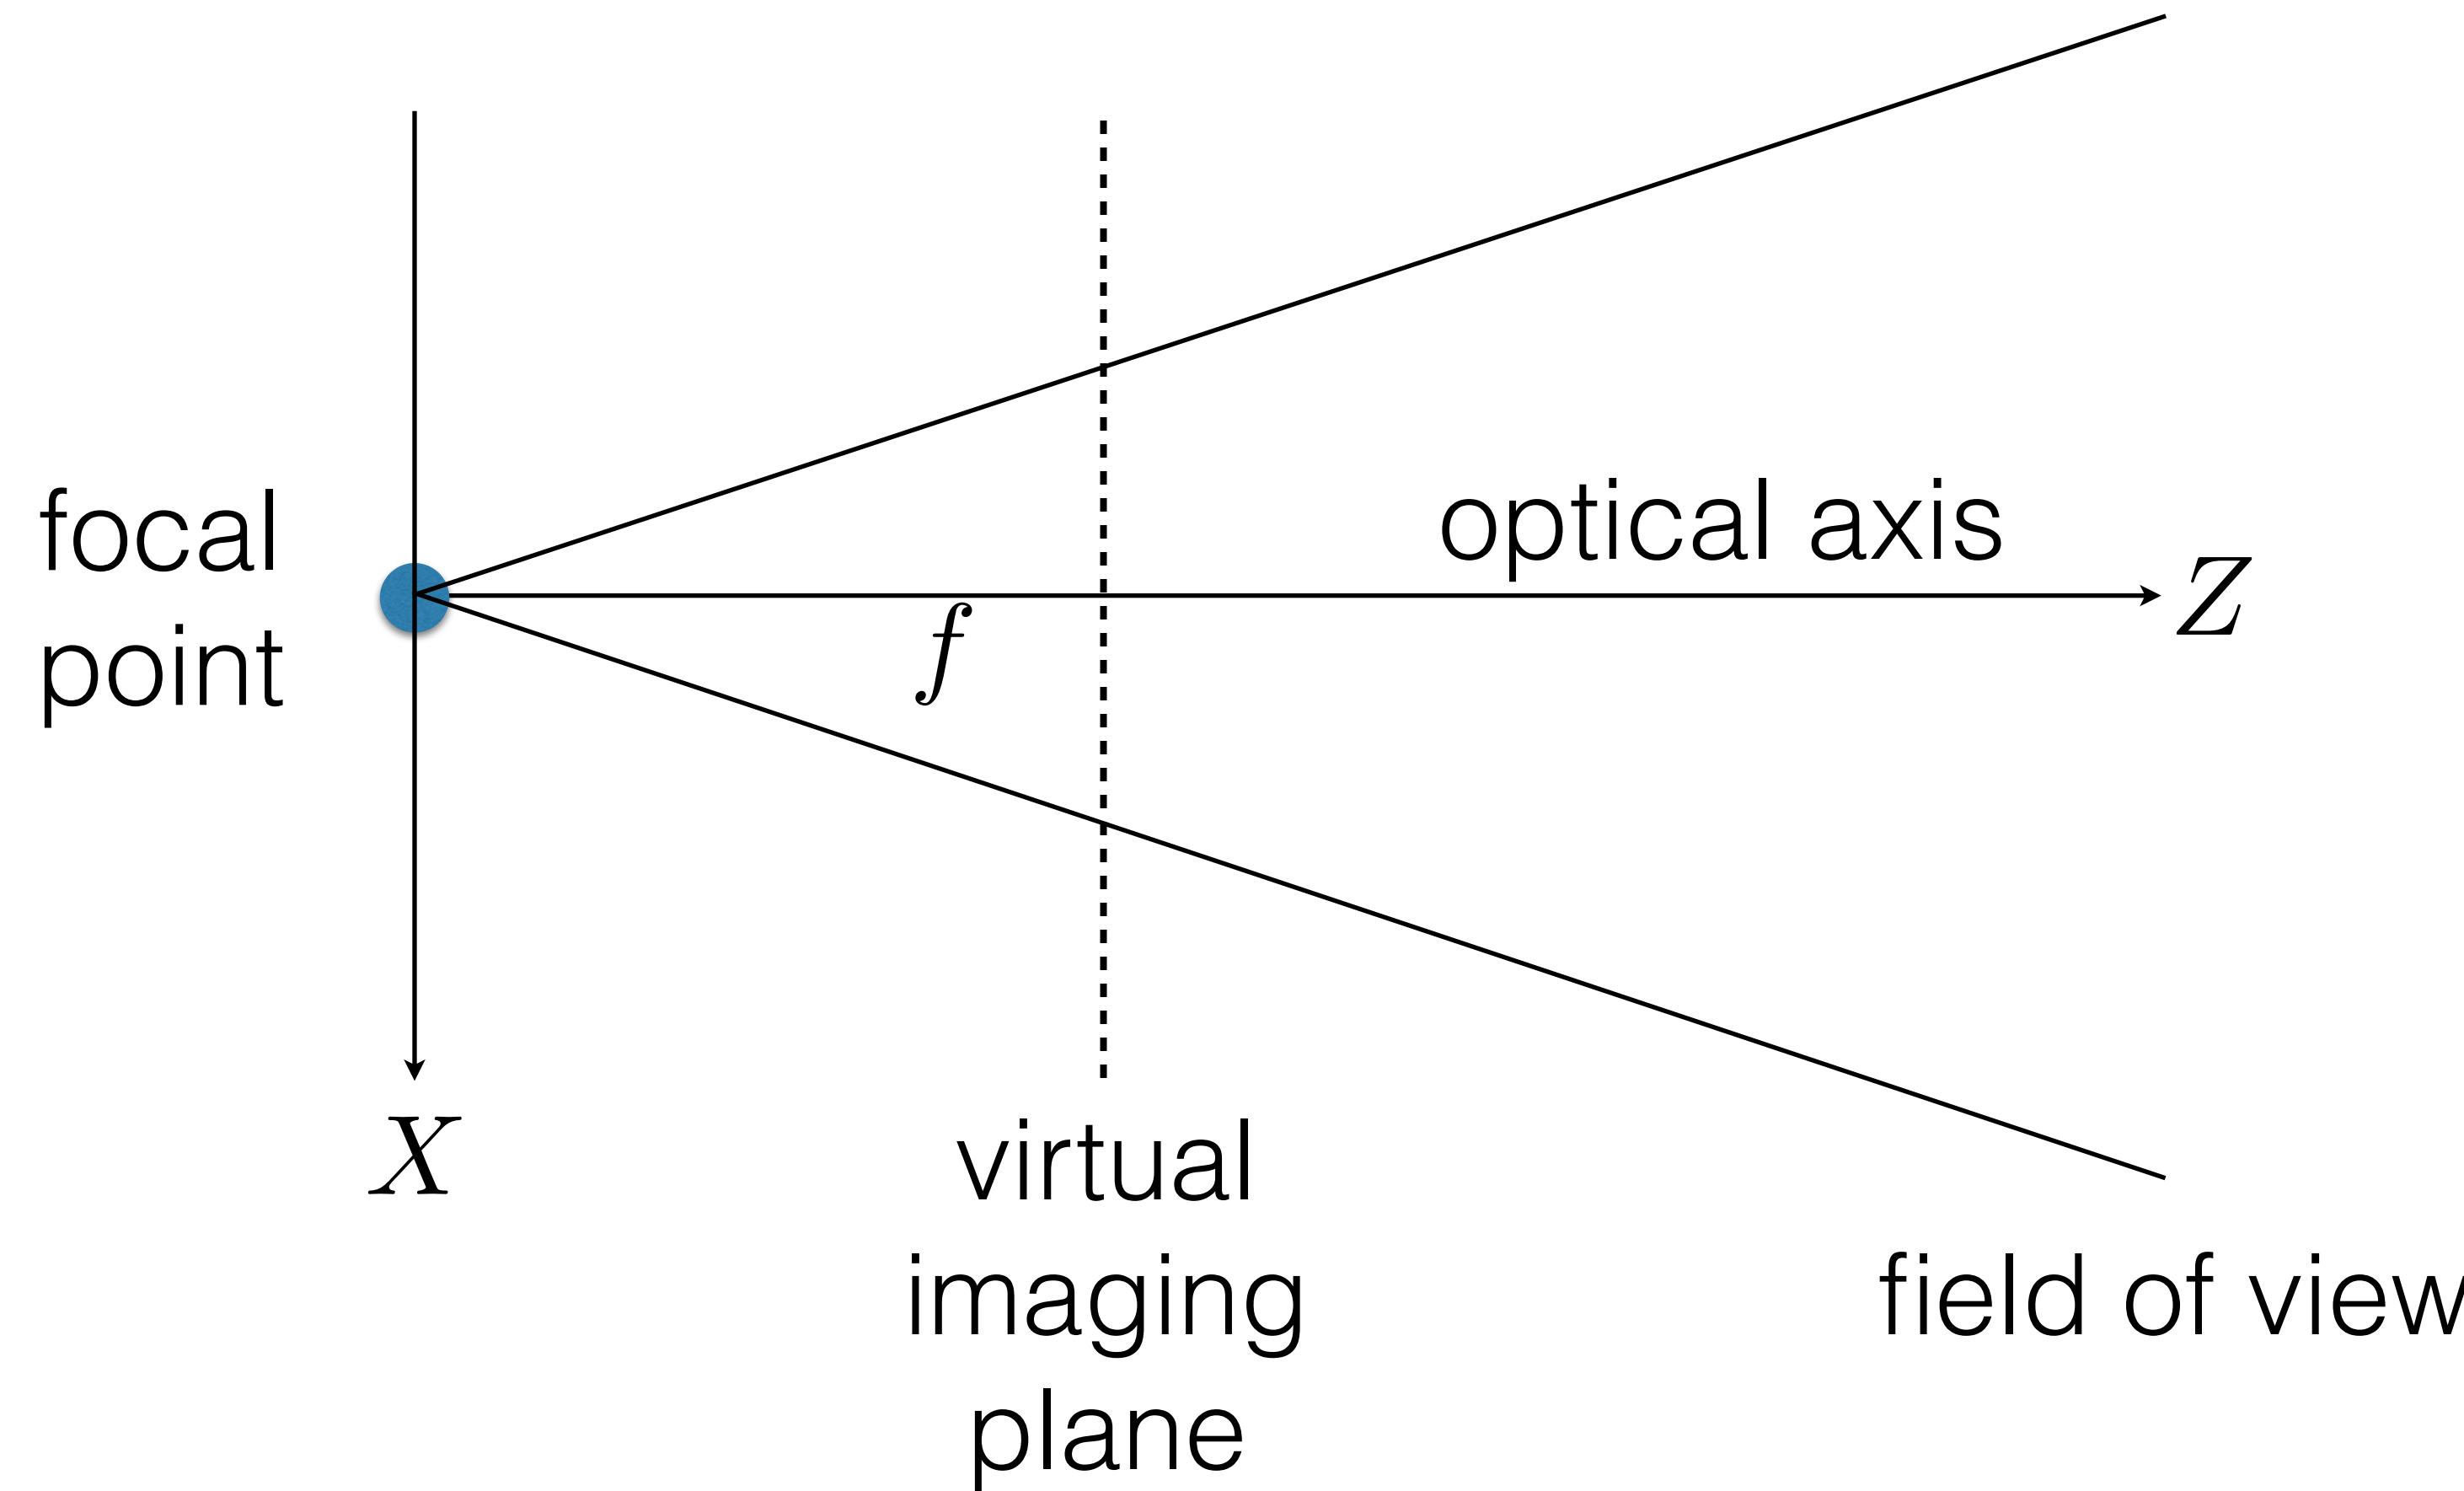
# Geometric Model



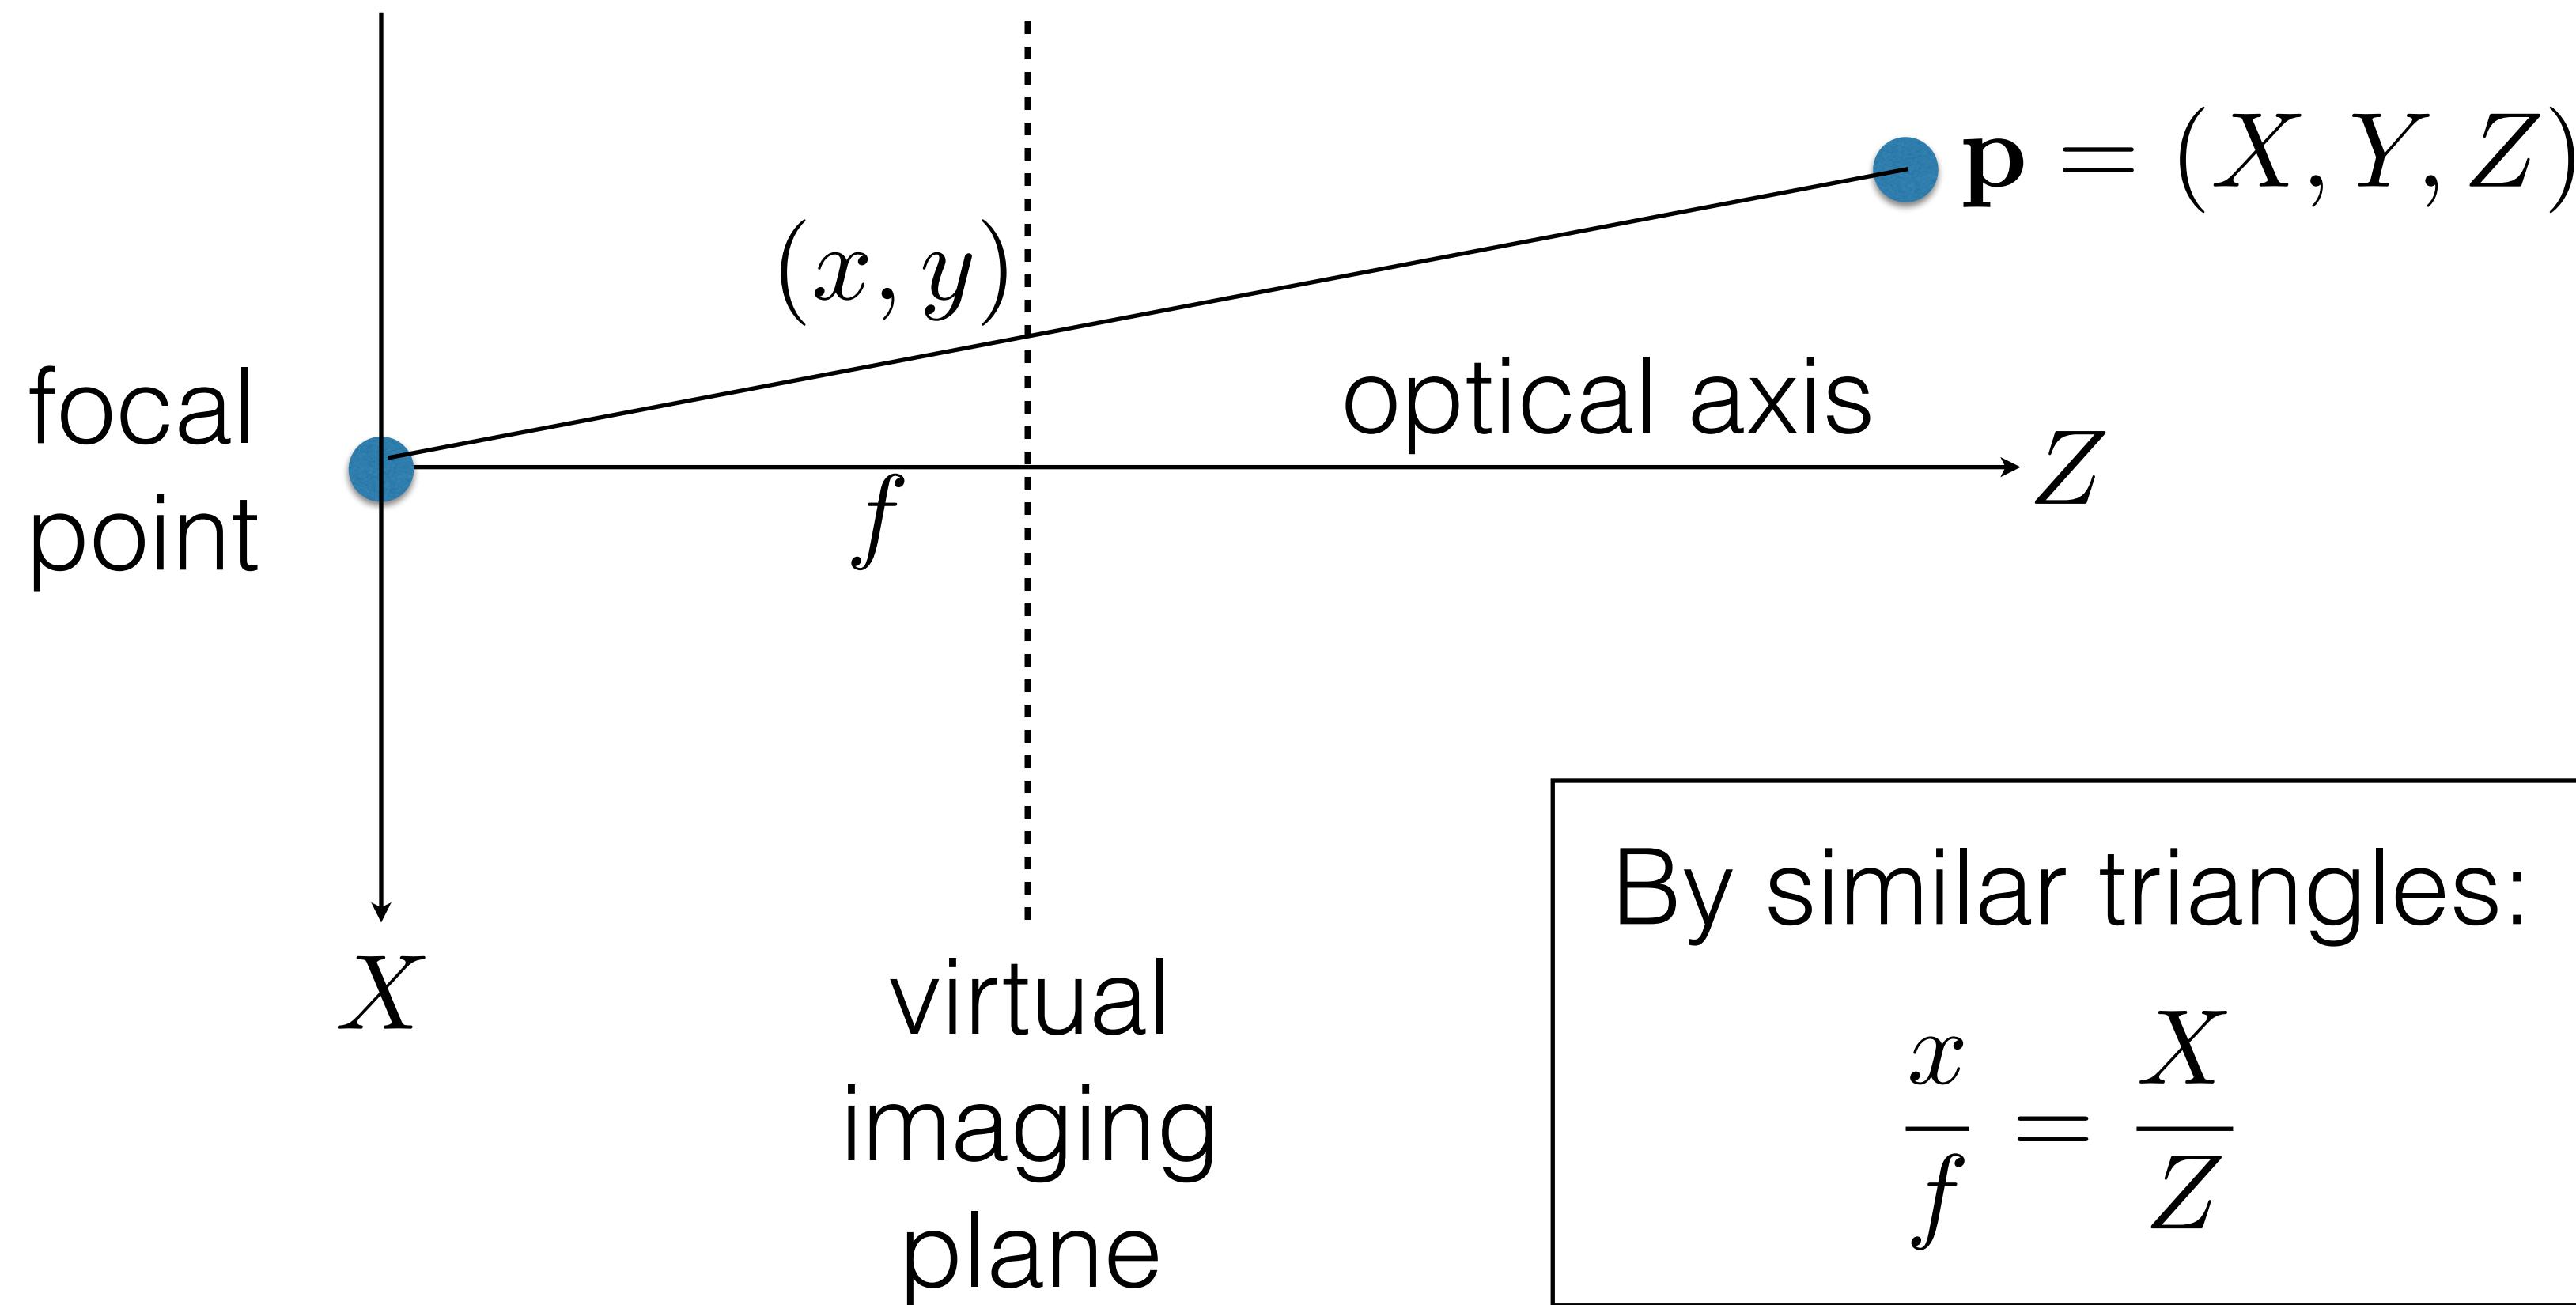
# Geometric Model



# Camera Coordinates



# Projection



# Projection

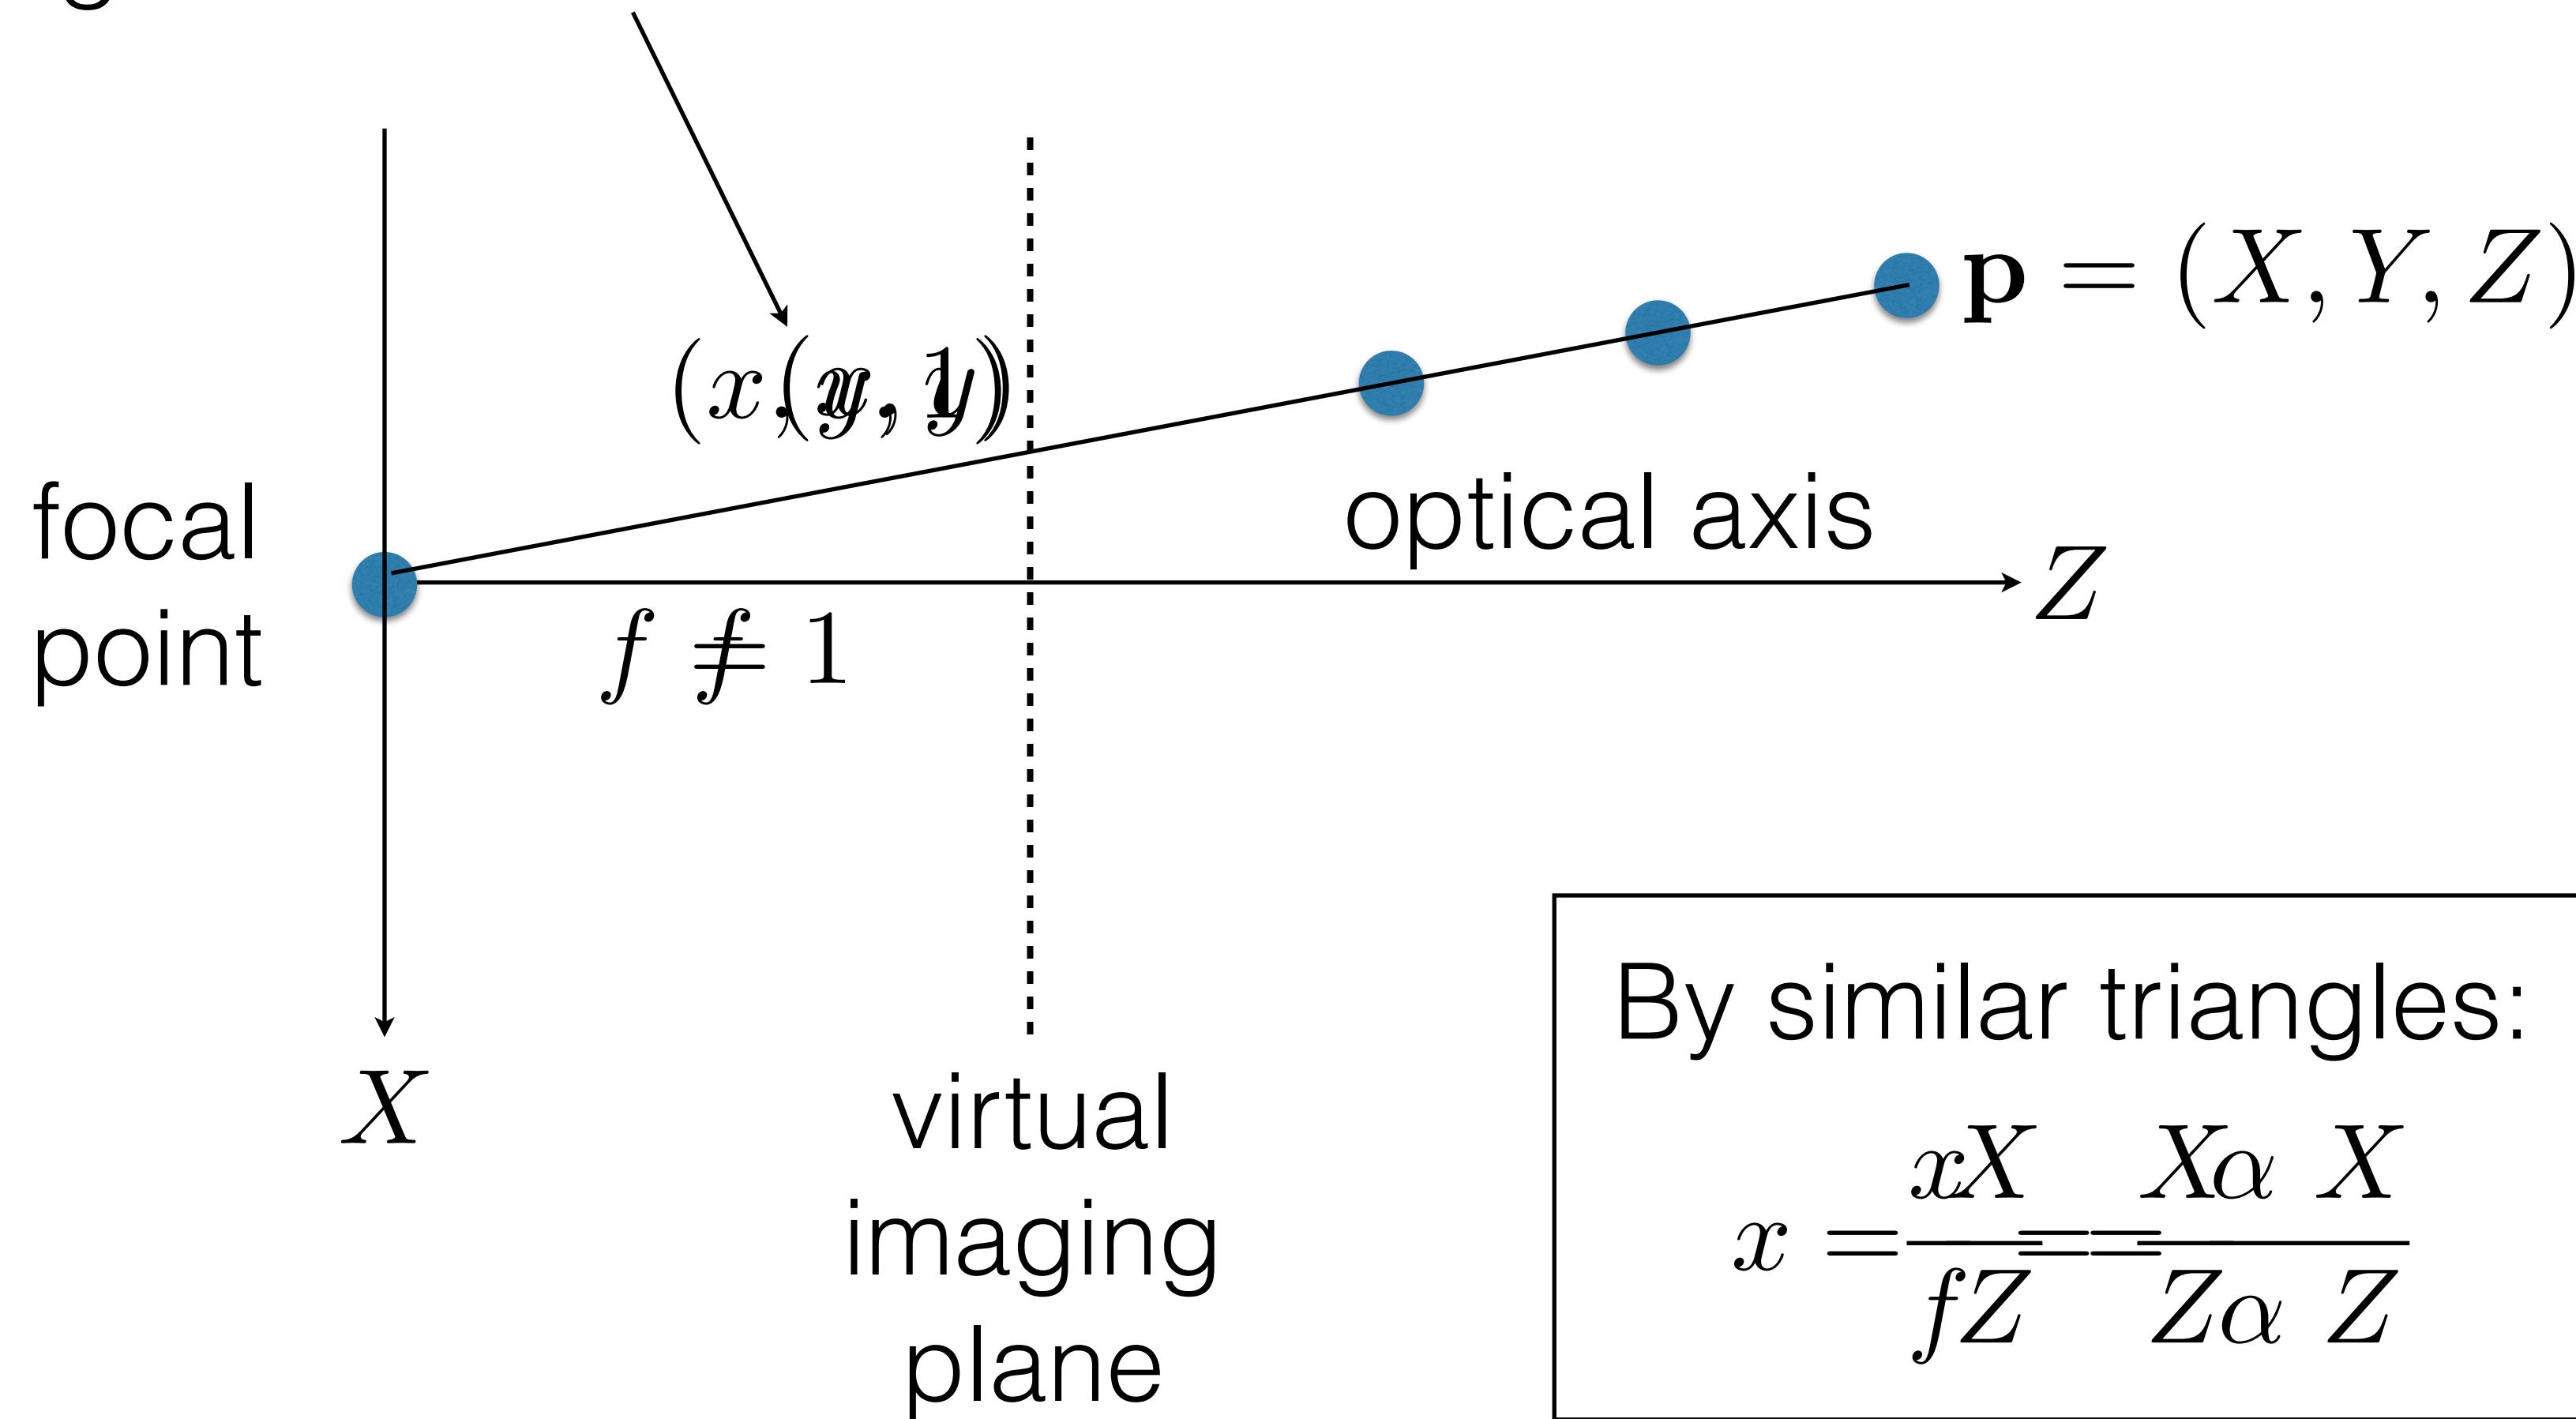
$$\frac{x}{f} = \frac{X}{Z} \quad \frac{y}{f} = \frac{Y}{Z}$$

$$(x, y) = \left( \frac{fX}{Z}, \frac{fY}{Z} \right)$$

Note: this is the projected coordinate in real-world units.  
To get actual pixel location, have to scale by pixel density  
and apply offset to image origin (more on this later...)

# Projection

homogeneous coordinate!

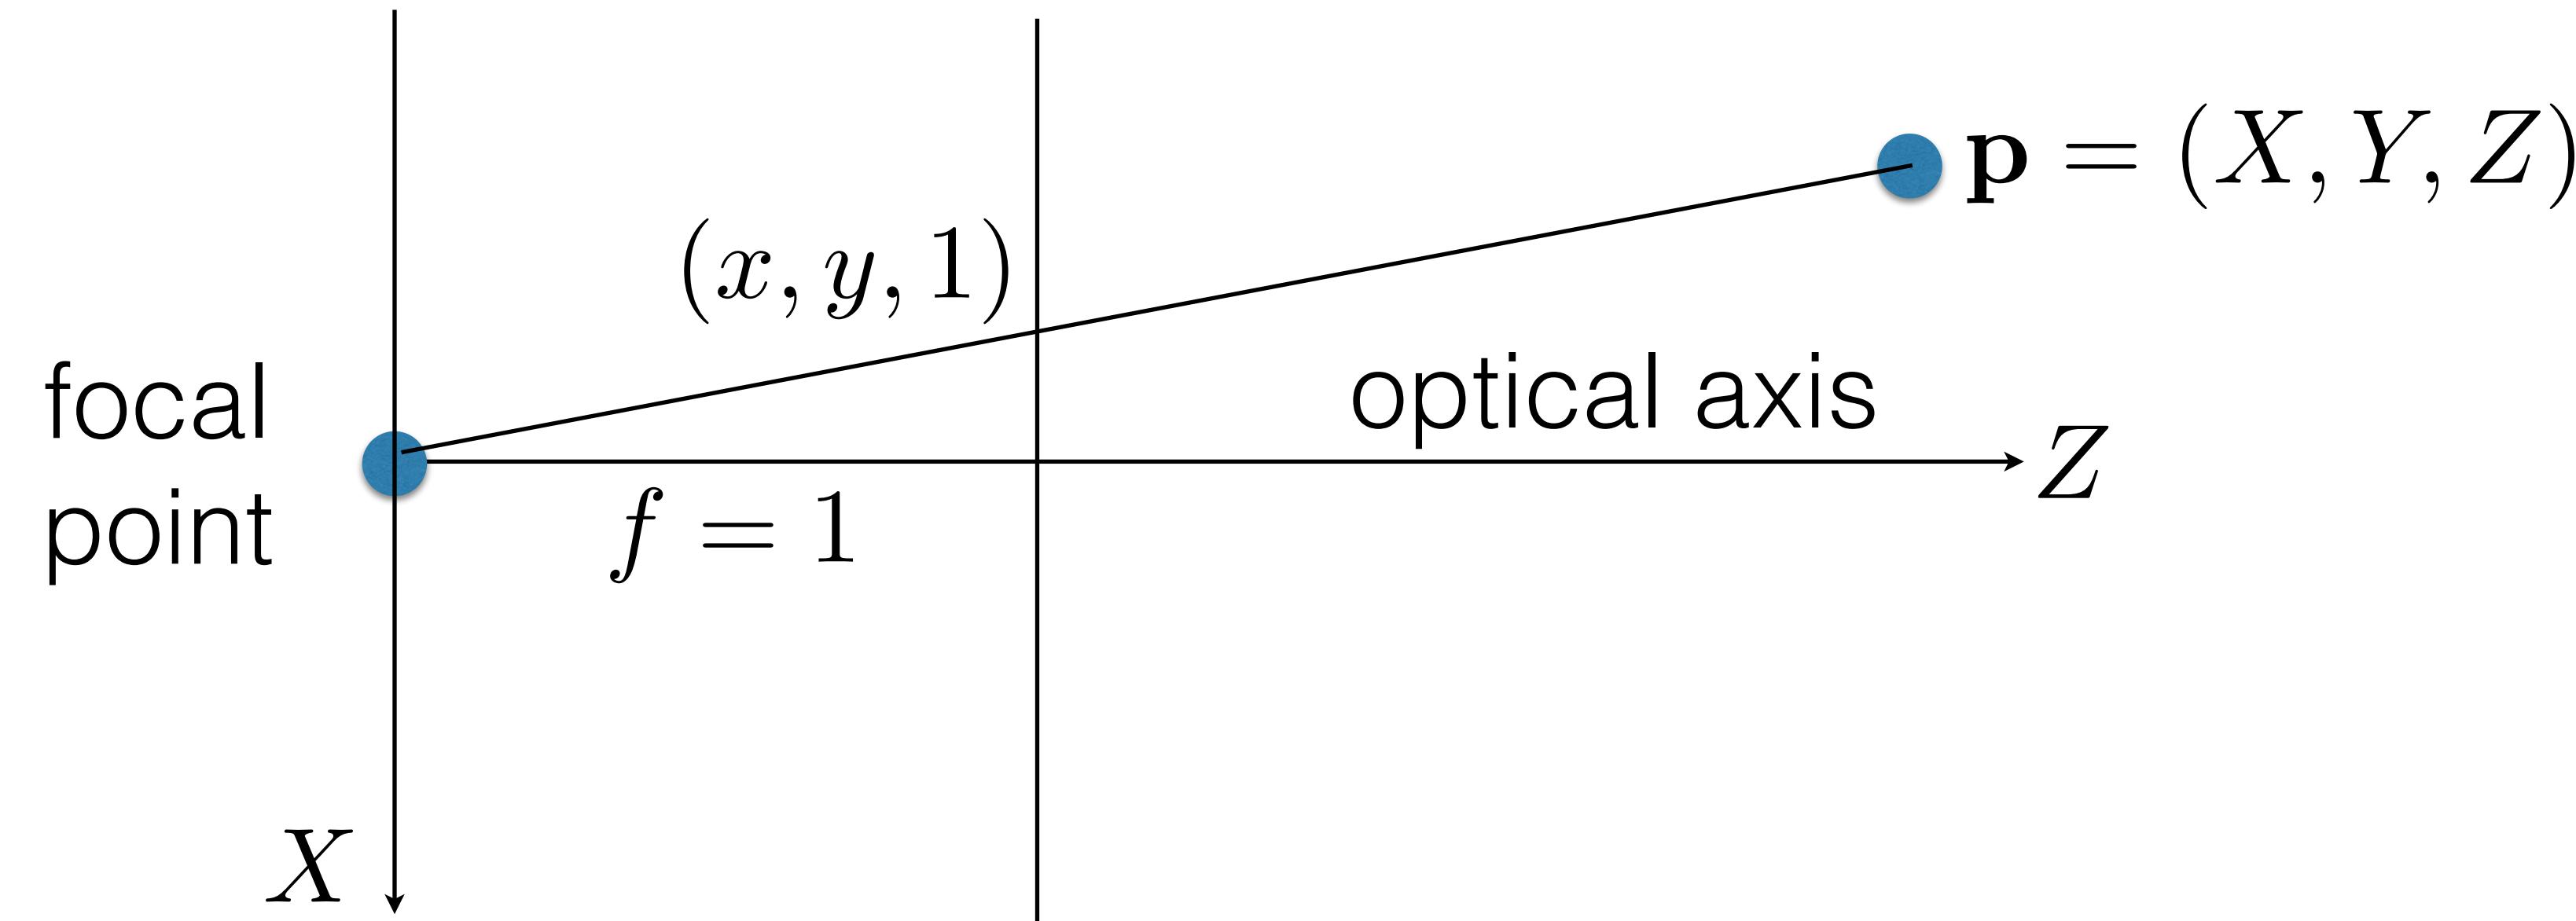


# Homogenous Coordinates

- Homogeneous coordinates are used to represent all 3D points along the ray that falls on the same 2D projection:

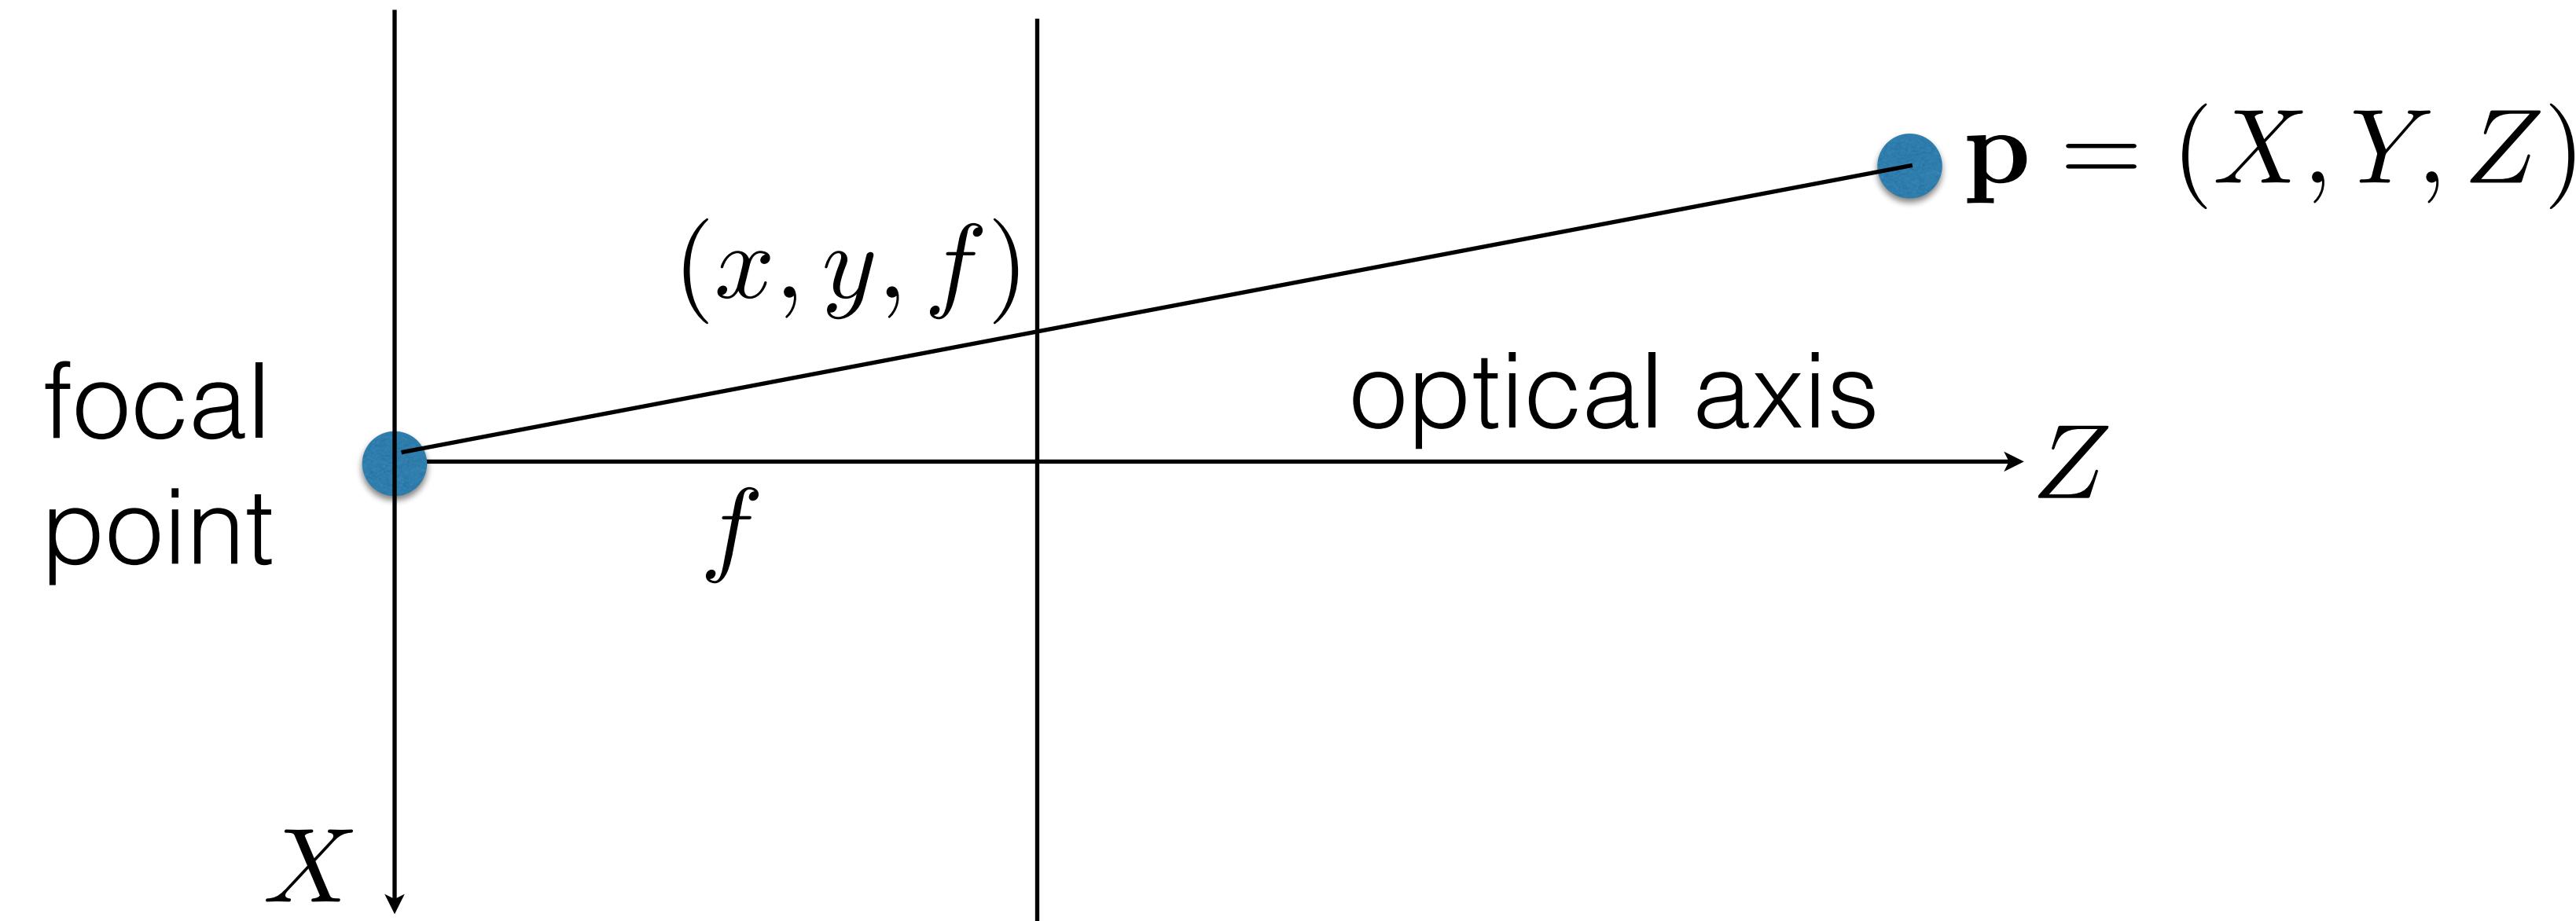
$$\begin{bmatrix} x \\ y \\ 1 \end{bmatrix} \sim \begin{bmatrix} \alpha x \\ \alpha y \\ \alpha \end{bmatrix}$$

# Perspective Projection



$$\begin{bmatrix} x \\ y \\ 1 \\ 1 \end{bmatrix} = \begin{bmatrix} X/Z \\ Y/Z \\ 1 \\ 1 \end{bmatrix} \sim \begin{bmatrix} X \\ Y \\ Z \\ Z \end{bmatrix} = \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 1 & 0 \end{bmatrix} \begin{bmatrix} X \\ Y \\ Z \\ 1 \end{bmatrix}$$

# Perspective Projection



$$\begin{bmatrix} x \\ y \\ f \\ 1 \end{bmatrix} = \begin{bmatrix} fX/Z \\ fY/Z \\ f \\ 1 \end{bmatrix} \sim \begin{bmatrix} X \\ Y \\ Z \\ Z/f \end{bmatrix} = \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 1/f & 0 \end{bmatrix} \begin{bmatrix} X \\ Y \\ Z \\ 1 \end{bmatrix}$$

# Alternative Form

One way (some implementation advantages):

$$\begin{bmatrix} x \\ y \\ \hline f \\ 1 \end{bmatrix} = \begin{bmatrix} fX/Z \\ fY/Z \\ \hline f \\ 1 \end{bmatrix} \sim \begin{bmatrix} X \\ Y \\ Z \\ Z/f \end{bmatrix} = \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 1/f & 0 \end{bmatrix} \begin{bmatrix} X \\ Y \\ Z \\ 1 \end{bmatrix}$$

More commonly used in graphics  
→

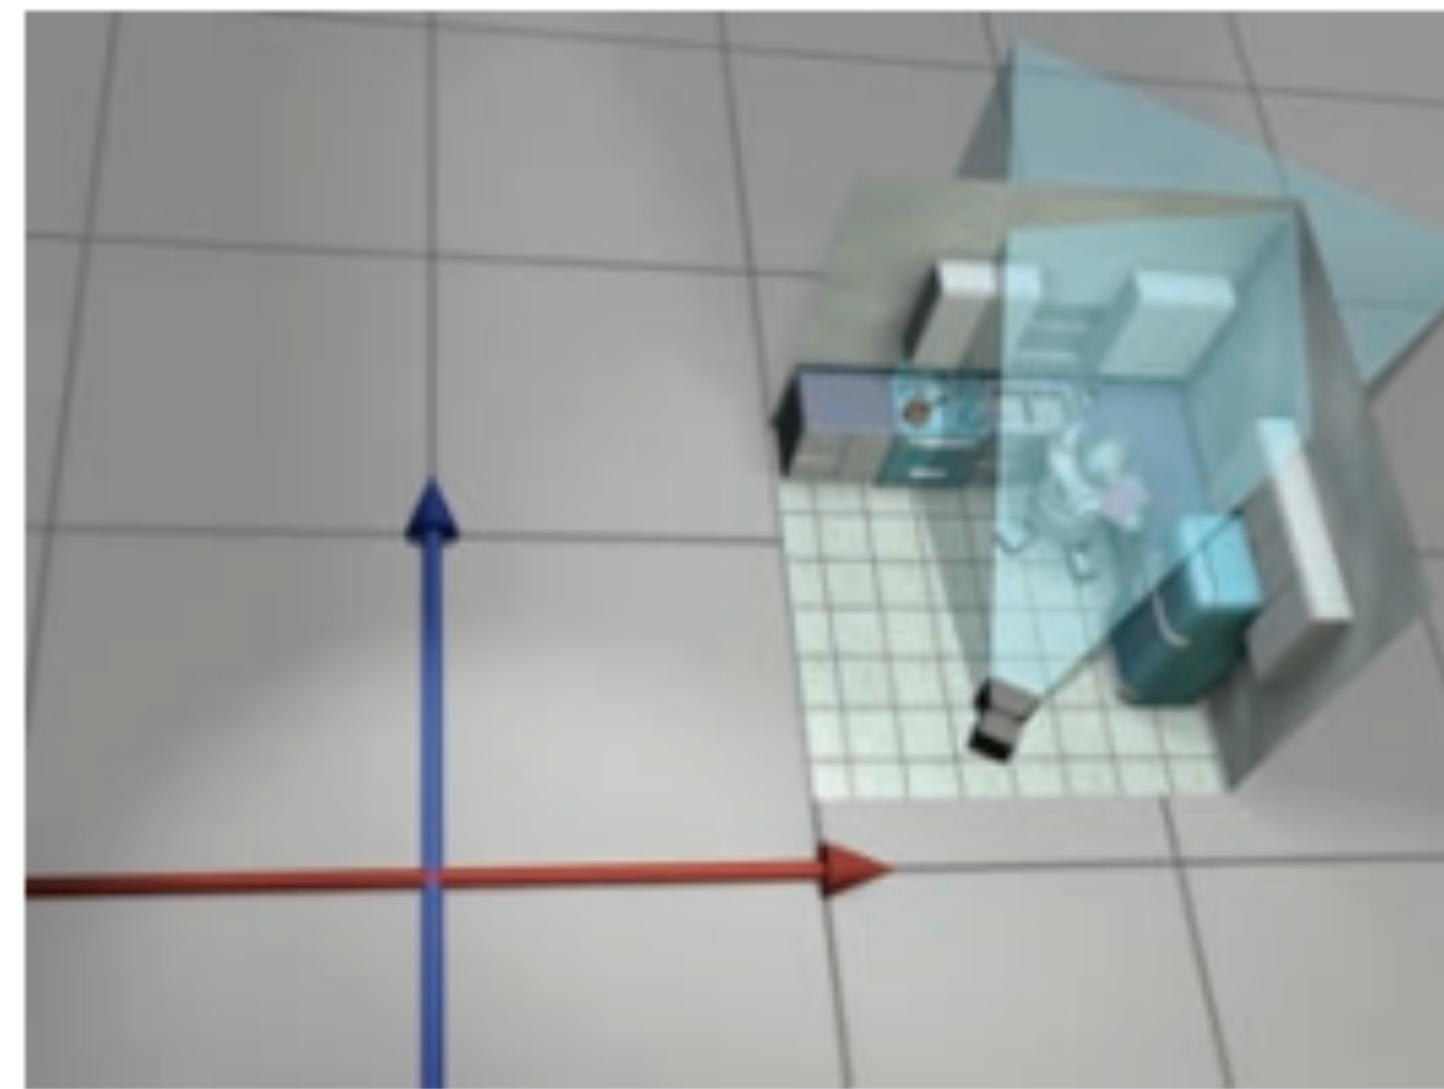
Another way (some conceptual advantages):

$$\begin{bmatrix} x \\ y \\ 1 \end{bmatrix} = \begin{bmatrix} fX/Z \\ fY/Z \\ \hline 1 \end{bmatrix} \sim \begin{bmatrix} X \\ Y \\ Z/f \end{bmatrix} = \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1/f & 0 \end{bmatrix} \begin{bmatrix} X \\ Y \\ Z \\ 1 \end{bmatrix}$$

More commonly used in vision  
→

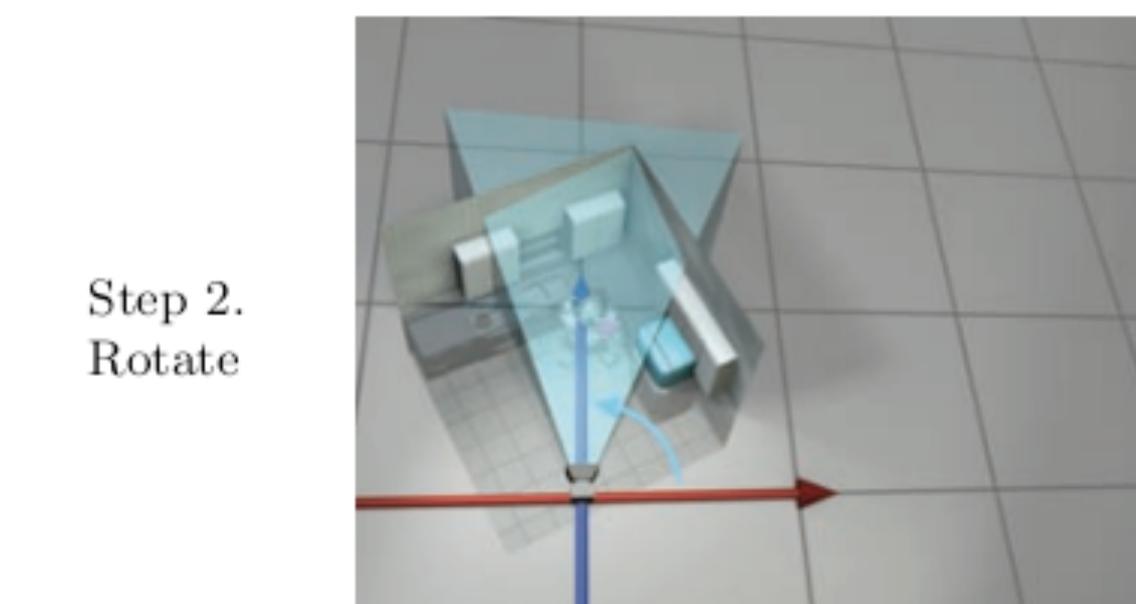
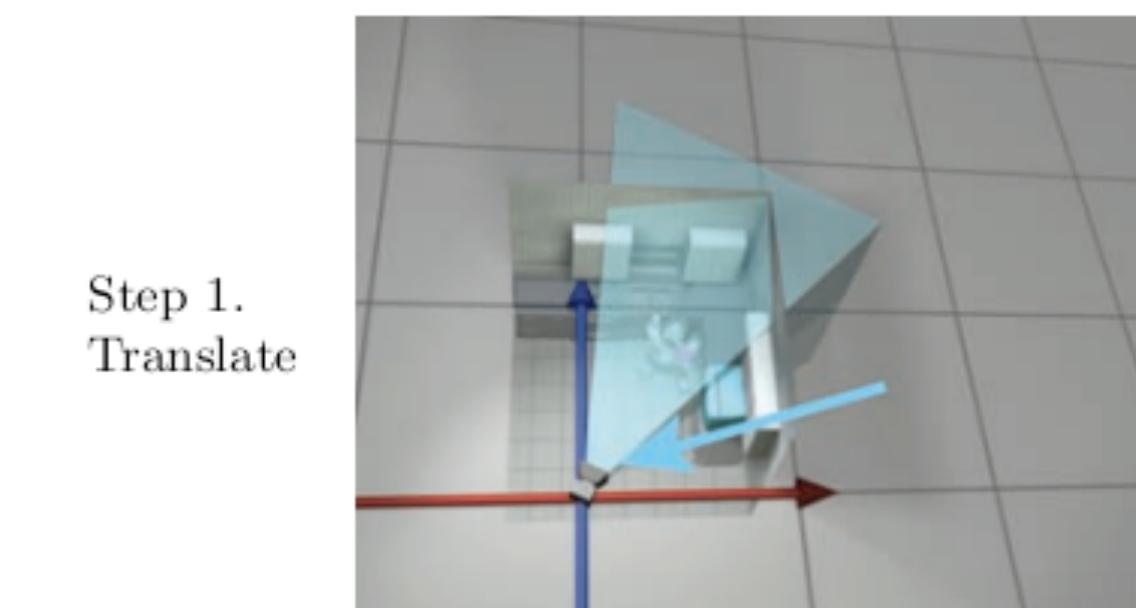
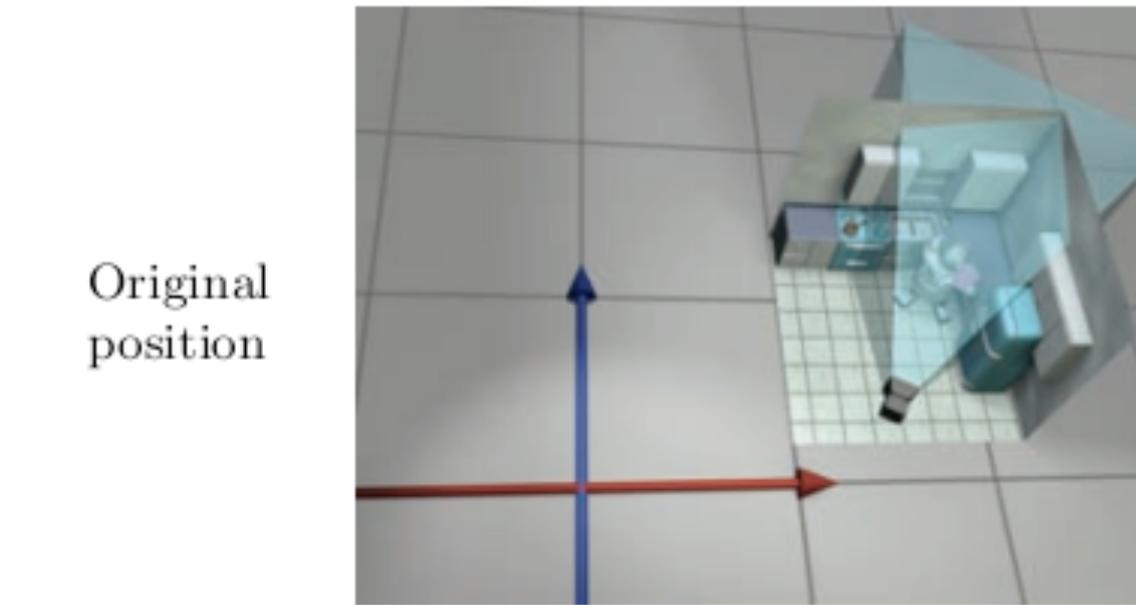
# 3D Spatial Reference Frames

- When working with multiple cameras/views in a 3D environment, we often use multiple 3D coordinate systems (frames of reference)
  - A global “world” coordinate system
  - Each camera’s camera-centric coordinate system
- In graphics, we often use an arbitrarily chosen world coordinate system and place virtual objects in that space
- In vision, we don’t often have a global reference frame, so we often use one of the cameras as the global reference frame



# World to Camera

- Two steps:
  - **Translate** everything to be relative to the camera position
  - **Rotate** into the camera's viewing orientation



# World to Camera

- Two steps:
  - **Translate** everything to be relative to the camera position
  - **Rotate** into the camera's viewing orientation

$$\begin{bmatrix} 1 & 0 & 0 & -c_x \\ 0 & 1 & 0 & -c_y \\ 0 & 0 & 1 & -c_z \\ 0 & 0 & 0 & 1 \end{bmatrix}$$

$$\begin{bmatrix} e_{11} & e_{12} & e_{13} & 0 \\ e_{21} & e_{22} & e_{23} & 0 \\ e_{31} & e_{32} & e_{33} & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix}$$

Camera's x  
Camera's y  
Camera's z

# 3D Perspective Projection

Let's factor this into a  $3 \times 4$   
“canonical projection matrix” and  
a camera-specific 2D scaling  
by focal length  $f$

World-to-camera transformation

---

$$\begin{bmatrix} x \\ y \\ f \\ 1 \end{bmatrix} \sim \begin{bmatrix} X_c \\ Y_c \\ Z_c \\ Z_c/f \end{bmatrix} = \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 1/f & 0 \end{bmatrix} \begin{bmatrix} e_{11} & e_{12} & e_{13} & 0 \\ e_{21} & e_{22} & e_{23} & 0 \\ e_{31} & e_{32} & e_{33} & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} 1 & 0 & 0 & -c_x \\ 0 & 1 & 0 & -c_y \\ 0 & 0 & 1 & -c_z \\ 0 & 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} X_w \\ Y_w \\ Z_w \\ 1 \end{bmatrix}$$

Normalize      Project      Rotate      Translate

---

# 3D Perspective Projection

World-to-camera transformation

---

$$\begin{bmatrix} x \\ y \\ 1 \end{bmatrix} \sim \begin{bmatrix} f & 0 & 0 \\ 0 & f & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \end{bmatrix} \begin{bmatrix} e_{11} & e_{12} & e_{13} & 0 \\ e_{21} & e_{22} & e_{23} & 0 \\ e_{31} & e_{32} & e_{33} & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} 1 & 0 & 0 & -c_x \\ 0 & 1 & 0 & -c_y \\ 0 & 0 & 1 & -c_z \\ 0 & 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} X_w \\ Y_w \\ Z_w \\ 1 \end{bmatrix}$$

←

Camera      Project      Rotate      Translate

**K**            **P**            **R**            **T**

# Intrinsic Camera Parameters

- What we've built so far works fine for modeling an ideal camera, but real ones aren't so simple:
- The optical axis through the center of the lens may not be aligned with the center of the imaging array
- The imaging sensors may not have a square aspect ratio
- The imaging array may not be mounted perpendicular to the optical axis

# Intrinsic Camera Parameters

Can model these intrinsic parameters this way:

$$\mathbf{K} = \begin{bmatrix} s_x & s_\theta & o_x \\ 0 & s_y & o_y \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} f & 0 & 0 \\ 0 & f & 0 \\ 0 & 0 & 1 \end{bmatrix} = \begin{bmatrix} fs_x & fs & o_x \\ 0 & fs_y & o_y \\ 0 & 0 & 1 \end{bmatrix}$$

$f$  focal length (real-world units)

$s_x$   $x$  scaling (pixels)

$s_y$   $y$  scaling ( pixels)

$s$  skew (pixels)

$o_x$   $x$  translation (pixels)

$o_y$   $y$  translation (pixels)

**Can calibrate the camera to determine these parameters**

# 3D Perspective Projection

World-to-camera transformation

$$\begin{bmatrix} x \\ y \\ 1 \end{bmatrix} \sim \begin{bmatrix} fs_x & fs & o_x \\ 0 & fs_y & o_y \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \end{bmatrix} \begin{bmatrix} e_{11} & e_{12} & e_{13} & 0 \\ e_{21} & e_{22} & e_{23} & 0 \\ e_{31} & e_{32} & e_{33} & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} 1 & 0 & 0 & -c_x \\ 0 & 1 & 0 & -c_y \\ 0 & 0 & 1 & -c_z \\ 0 & 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} X_w \\ Y_w \\ Z_w \\ 1 \end{bmatrix}$$

←

Camera      Project      Rotate      Translate

**K**            **P**            **R**            **T**

# 3D Perspective Projection

**Putting it all together:**

$$\mathbf{p} = \mathbf{K} \mathbf{P} \mathbf{R} \mathbf{T} \mathbf{p}_w$$

**Or sometimes written:**

$$\mathbf{p} = \mathbf{K} \mathbf{P} \begin{bmatrix} \mathbf{R} & \mathbf{T} \\ \mathbf{0} & 1 \end{bmatrix} \mathbf{p}_w$$

**Or just:**

$$\mathbf{p} = \mathbf{M} \mathbf{p}_w$$

# Lens Distortion

- One more geometric transformation:  
*radial distortion* caused by the lens:
  - Straight lines before curved
  - Usually radially symmetric  
(function of distance from the center)
  - More pronounced with larger FOV lenses
- How to fix it:
  - Model the spatial distortion  
(usually by imaging a grid of straight lines)
  - Pre-calibrate the camera/lens  
(usually with other calibration as well)
  - Warp the images to correct for the distortion  
before further processing



# Camera Calibration

- Essential idea: take pictures of things with easy-to-find points/lines in known configurations
- Use these to solve for the unknowns in the pipeline
- Intrinsic calibration:  
Solving for the camera's internal parameters  
( $\mathbf{K}$  matrix, radial distortion)
- Extrinsic calibration:  
Solving for the camera's external parameters  
(position  $\mathbf{T}$  and orientation  $\mathbf{R}$ )

